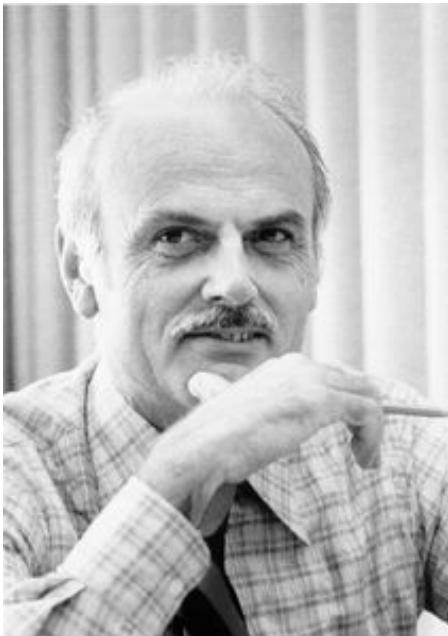


Intro to the Relational Model

Introduction to Database Management

CS348 Spring 2021

Edgar F. Codd (1923-2003)



- Pilot in the Royal Air Force in WW2
- Inventor of the relational model and algebra while at IBM
- Turing Award, 1981

Outline

- Part 1: Relational data model
- Part 2: Relational algebra

Example for Relational data model

User

<i>uid</i>	<i>name</i>	<i>age</i>	<i>pop</i>
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...

relations (or tables)

Group

<i>gid</i>	<i>name</i>
abc	Book Club
gov	Student Government
dps	Dead Putting Society
...	...

Member

<i>uid</i>	<i>gid</i>
142	dps
123	gov
857	abc
857	gov
456	abc
456	gov
...	...

Example for Relational data model

User

uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...

attributes (or columns)

Group

gid	name
abc	Book Club
gov	Student Government
dps	Dead Putting Society
...	...

Member

uid	gid
142	dps
123	gov
857	abc
857	gov
456	abc
456	gov
...	...

Example for Relational data model

User

<i>uid</i>	<i>name</i>	<i>age</i>	<i>pop</i>
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...

domain (or type)

Group

<i>gid</i>	<i>name</i>
abc	Book Club
gov	Student Government
dps	Dead Putting Society
...	...

Member

<i>uid</i>	<i>gid</i>
142	dps
123	gov
857	abc
857	gov
456	abc
456	gov
...	...

Example for Relational data model

User

uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...

tuples (or rows)

Duplicates are not allowed

Ordering of rows doesn't matter
(even though output is
always in some order)

Group

gid	name
abc	Book Club
gov	Student Government
dps	Dead Putting Society
...	...

Member

uid	gid
142	dps
123	gov
857	abc
857	gov
456	abc
456	gov
...	...

Example for Relational data model

User

uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...

Group

gid	name
abc	Book Club
gov	Student Government
dps	Dead Putting Society
...	...

Member

uid	gid
142	dps
123	gov
857	abc
857	gov
456	abc
456	gov
...	...

User: {⟨142, Bart, 10, 0.9⟩, ⟨857, Milhouse, 10, 0.2⟩, ... }

Group: {⟨abc, Book Club⟩, ⟨gov, Student Government⟩, ... }

Member: {⟨142, dps⟩, ⟨123, gov⟩, ... }

Relational data model

- A database is a collection of **relations** (or **tables**)
- Each relation has a set of **attributes** (or **columns**)
- Each attribute has a name and a **domain** (or **type**)
 - The domains are required to be atomic
- Each relation contains a set of **tuples** (or **rows**)
 - Each tuple has a value for each attribute of the relation
 - **Duplicate tuples are not allowed**
 - Two tuples are duplicates if they agree on all attributes

☞ Simplicity is a virtue!

Schema vs. instance

- Schema (metadata)
 - Specifies the logical structure of data
 - Is defined at setup time, rarely changes

User (uid int, name string, age int, pop float)

Group (gid string, name string)

Member (uid int, gid string)

- Instance

- Represents the data content
- Changes rapidly, but always conforms to the schema

User: {⟨142, Bart, 10, 0.9⟩, ⟨857, Milhouse, 10, 0.2⟩, ... }

Group: {⟨abc, Book Club⟩, ⟨gov, Student Government⟩, ... }

Member: {⟨142, dps⟩, ⟨123, gov⟩, ... }

Integrity constraints

- A set of rules that database instances should follow
- Example:
 - *age* cannot be **negative**
 - *uid* should be **unique** in the *User* relation
 - *uid* in *Member* must **refer to** a row in *User*

User (*uid* int, *name* string, *age* int, *pop* float)

Group (*gid* string, *name* string)

Member (*uid* int, *gid* string)

User: {⟨142, Bart, 10, 0.9⟩, ⟨857, Milhouse, 10, 0.2⟩, ... }

Group: {⟨abc, Book Club⟩, ⟨gov, Student Government⟩, ... }

Member: {⟨142, dps⟩, ⟨123, gov⟩, ... }

Integrity constraints

- An instance is only **valid** if it satisfies all the integrity constraints.
- Reasons to use constraints:
 - Ensure data entry/modification respects to database design
 - Protect data from bugs in applications

Types of integrity constraints

- Tuple-level
 - Domain restrictions, attribute comparisons, etc.
 - E.g. *age* cannot be **negative**
- Relation-level
 - **Key constraints** (focus in this lecture)
 - E.g. *uid* should be **unique** in the *User* relation
 - Functional dependencies (Textbook, Ch. 7)
- Database-level
 - Referential integrity – **foreign key** (focus in this lecture)
 - *uid* in *Member* must **refer to** a row in *User* with the same *uid*

Key (Candidate Key)

Def: A set of attributes K for a relation R if

- **Condition 1:** In no instance of R will two different tuples agree on all attributes of K
 - That is, K can serve as a “tuple identifier”
- **Condition 2:** No proper subset of K satisfies the above condition
 - That is, K is **minimal**
- Example: User (uid , $name$, age , pop)
 - uid is a key of User
 - age is not a key (not an identifier)
 - $\{uid, name\}$ is not a key (not minimal), but a **superkey**

Satisfies only
Condition 1

Schema vs. instance

<i>uid</i>	<i>name</i>	<i>age</i>	<i>pop</i>
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3

- Is *name* a key of *User*?
 - Yes? Seems reasonable for this instance
 - No! User names are not unique **in general**
- Key declarations are part of the schema

More examples of keys

- Member (*uid*, *gid*)

- {*uid*, *gid*}

☞ A key can contain multiple attributes

- Address (*street_address*, *city*, *state*, *zip*)

- Key 1: {*street_address*, *city*, *state*}
 - Key 2: {*street_address*, *zip*}

☞ A relation can have multiple keys!

- Primary key: a designated candidate key in the schema declaration

- Underline all its attributes, e.g., Address (*street_address*, *city*, *state*, *zip*)

Member

<i>uid</i>	<i>gid</i>
142	dps
123	gov
857	abc
857	gov
456	abc
456	gov
...	...

Use of keys

- More constraints on data, fewer mistakes
- Look up a row by its key value
 - Many selection conditions are “key = value”
- “Pointers” to other rows (often across tables)

“Pointers” to other rows

- **Foreign key:** primary key of one relation appearing as attribute of another relation

User	<u>uid</u>	<u>name</u>	<u>age</u>	<u>pop</u>
	142	Bart	10	0.9
	123	Milhouse	10	0.2
	857	Lisa	8	0.7
	456	Ralph	8	0.3

Group	
<u>gid</u>	<u>name</u>
abc	Book Club
gov	Student Government
dps	Dead Putting Society
...	...

Member	<u>uid</u>	<u>gid</u>
	142	dps
	123	gov
	857	abc
	857	gov
	456	abc
	456	gov

“Pointers” to other rows

- **Referential integrity:** A tuple with a non-null **value** for a **foreign key** that does **not match the primary key value** of a tuple in the referenced relation is not allowed.

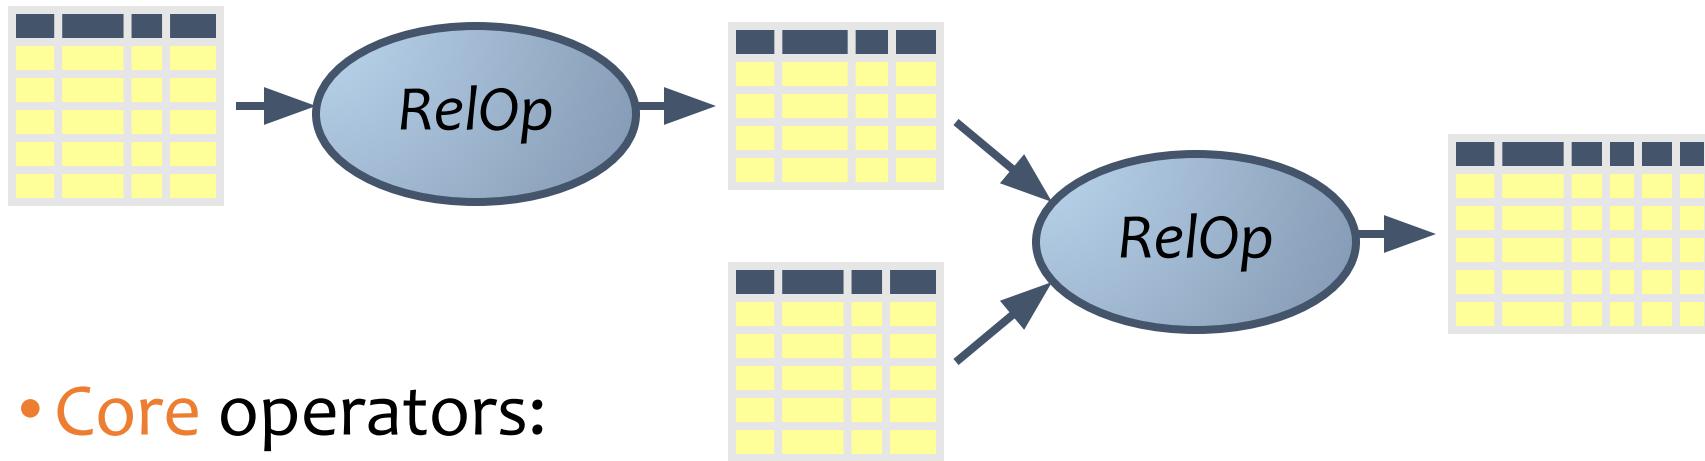
Group		Member	
<u>gid</u>	<u>name</u>	<u>uid</u>	<u>gid</u>
abc	Book Club	142	dps
gov	Student Government	123	gov
dps	Dead Putting Society	857	Xi
		857	gov
		456	abc
		456	gov
	

Outline

- Part 1: Relational data model
 - Data model
 - Database schema
 - Integrity constraints (keys)
 - Languages
 - Relational algebra (focus in this lecture)
 - SQL (next lecture)
 - Relational calculus (textbook, Ch. 27)
- Part 2: Relational algebra

Relational algebra

A language for querying relational data based on “operators”



- **Core** operators:
 - Selection, projection, cross product, union, difference, and renaming
- Additional, **derived** operators:
 - Join, natural join, intersection, etc.
- Compose operators to make complex queries

Core operator 1: Selection

- Example: Users with popularity higher than 0.5

$$\sigma_{pop > 0.5} User$$

uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...



uid	name	age	pop
142	Bart	10	0.9
857	Lisa	8	0.7
...

Core operator 1: Selection

- Input: a table R
- Notation: $\sigma_p R$
 - p is called a **selection condition** (or **predicate**)
- Purpose: filter rows according to some criteria
- Output: same columns as R , but only rows of R that satisfy p

More on selection

- Selection condition can include any column of R , constants, comparison ($=, \leq$, etc.) and Boolean connectives (\wedge : and, \vee : or, \neg : not)
 - Example: users with popularity at least 0.9 and age under 10 or above 12

$$\sigma_{pop \geq 0.9 \wedge (age < 10 \vee age > 12)} User$$

- You must be able to evaluate the condition over **each single row** of the input table!
 - Example: the most popular user

$$\sigma_{pop \geq \text{every } pop \text{ in } User} User$$

**WRON
G!**

Core operator 2: Projection

- Example: IDs and names of all users

$$\pi_{uid, name} \text{ } User$$

uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...



uid	name
142	Bart
123	Milhouse
857	Lisa
456	Ralph
...	...

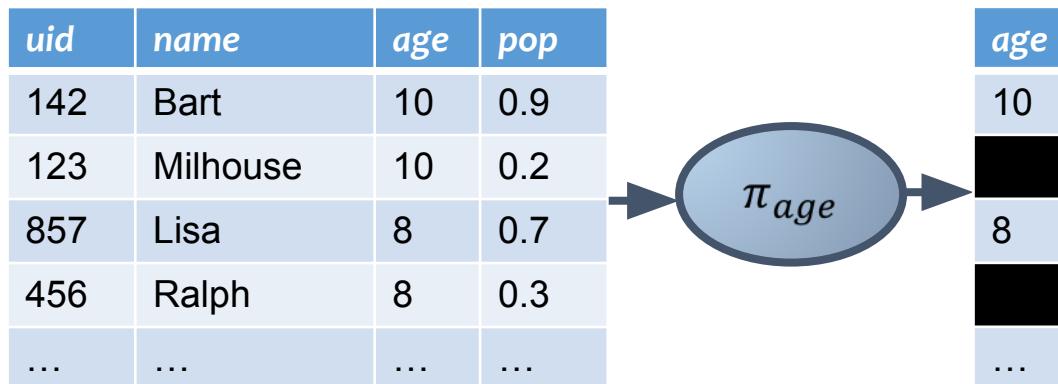
Core operator 2: Projection

- Input: a table R
- Notation: $\pi_L R$
 - L is a list of columns in R
- Purpose: output chosen columns
- Output: “same” rows, but only the columns in L

More on projection

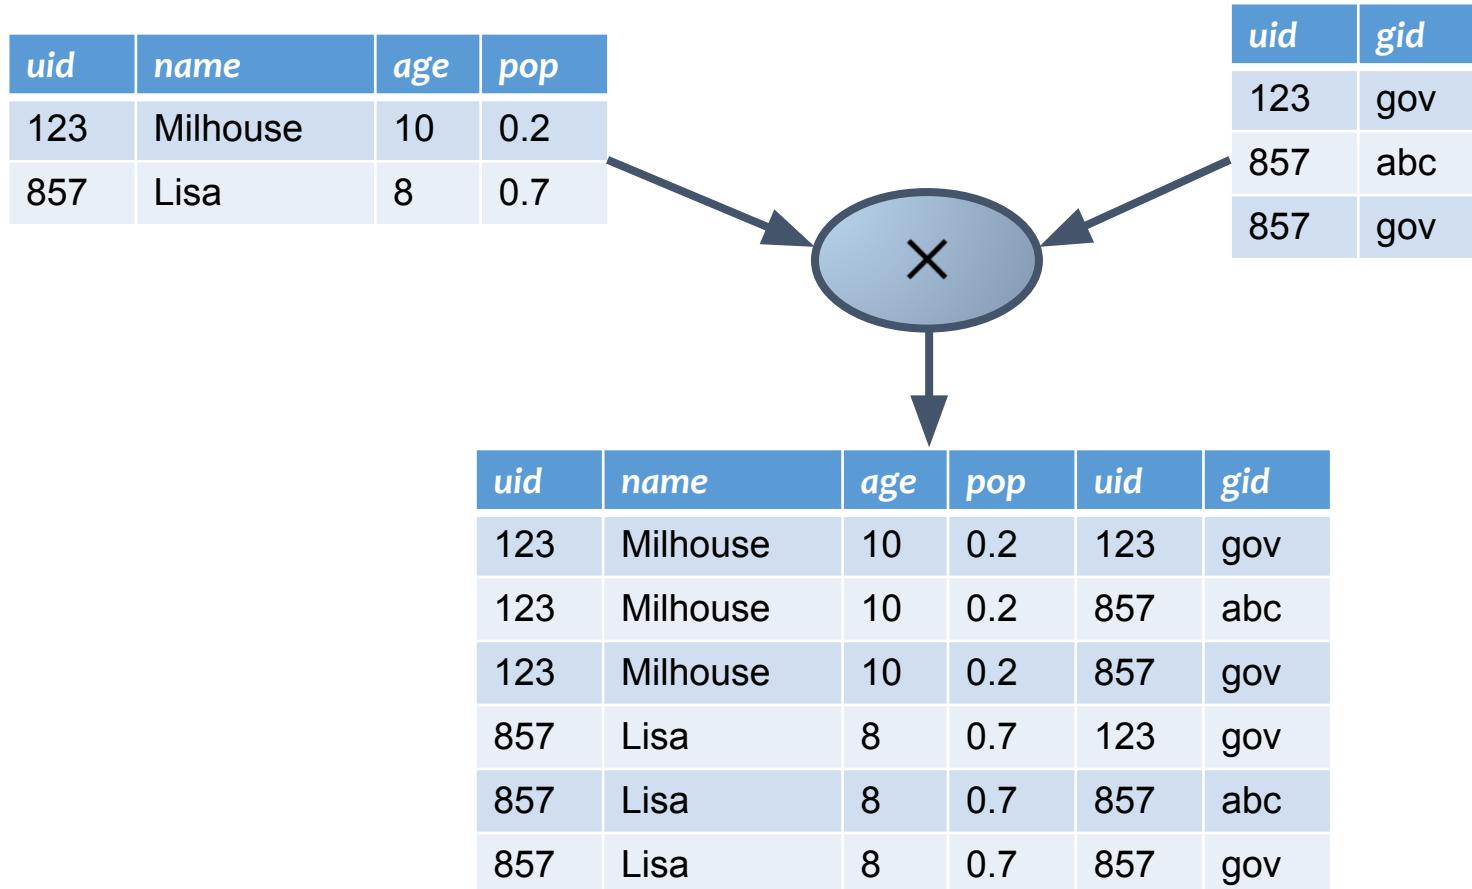
- Duplicate output rows are removed (by definition)
 - Example: user ages

$\pi_{age} User$



Core operator 3: Cross product

- $User \times Member$



Core operator 3: Cross product

- Input: two tables R and S
- Natation: $R \times S$
- Purpose: pairs rows from two tables
- Output: for each row r in R and each s in S , output a row rs (concatenation of r and s)

A note a column ordering

- Ordering of columns is unimportant as far as contents are concerned

<i>uid</i>	<i>name</i>	<i>age</i>	<i>pop</i>	<i>uid</i>	<i>gid</i>
123	Milhouse	10	0.2	123	gov
123	Milhouse	10	0.2	857	abc
123	Milhouse	10	0.2	857	gov
857	Lisa	8	0.7	123	gov
857	Lisa	8	0.7	857	abc
857	Lisa	8	0.7	857	gov
...

=

<i>uid</i>	<i>gid</i>	<i>uid</i>	<i>name</i>	<i>age</i>	<i>pop</i>
123	gov	123	Milhouse	10	0.2
857	abc	123	Milhouse	10	0.2
857	gov	123	Milhouse	10	0.2
123	gov	857	Lisa	8	0.7
857	abc	857	Lisa	8	0.7
857	gov	857	Lisa	8	0.7
...

- So cross product is **commutative**, i.e., for any R and S , $R \times S = S \times R$ (up to the ordering of columns)

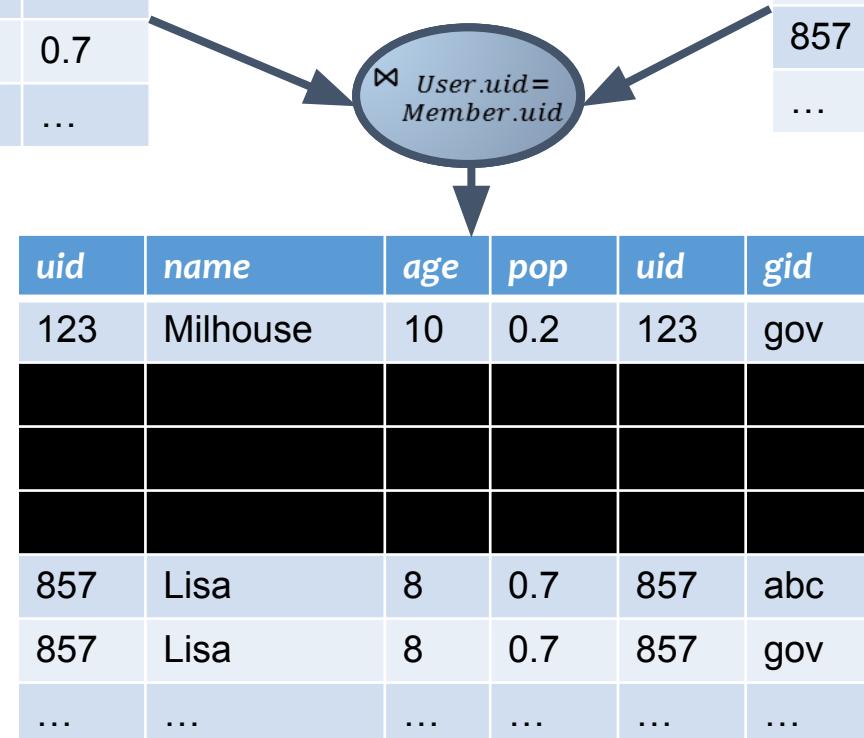
Derived operator 1: Join

- Info about users, plus IDs of their groups

$User \bowtie_{User.uid=Member.uid} Member$

uid	name	age	pop
123	Milhouse	10	0.2
857	Lisa	8	0.7
...

uid	gid
123	gov
857	abc
857	gov
...	...



Prefix a column reference with table name and “.” to disambiguate identically named columns from different tables

Derived operator 1: Join

- Input: two tables R and S
- Notation: $R \bowtie_p S$
 - p is called a **join condition** (or **predicate**)
- Purpose: relate rows from two tables according to some criteria
- Output: for each row r in R and each row s in S , output a row rs if r and s satisfy p
- Shorthand for $\sigma_p(R \times S)$
- (A.k.a. “theta-join”)

Derived operator 2: Natural join

$User \bowtie Member$

$$= \pi_{uid, name, age, pop, gid} (User \bowtie_{User.uid = Member.uid} Member)$$

uid	name	age	pop
123	Milhouse	10	0.2
857	Lisa	8	0.7
...

uid	gid
123	gov
857	abc
857	gov
...	...

uid	name	age	pop		gid
123	Milhouse	10	0.2		gov
857	Lisa	8	0.7		abc
857	Lisa	8	0.7		gov
...

Derived operator 2: Natural join

- Input: two tables R and S
- Notation: $R \bowtie S$
- Purpose: relate rows from two tables, and
 - Enforce equality between identically named columns
 - Eliminate one copy of identically named columns
- Shorthand for $\pi_L(R \bowtie_p S)$, where
 - p equates each pair of columns common to R and S
 - L is the union of column names from R and S (with duplicate columns removed)

Core operator 4: Union

- Input: two tables R and S
- Notation: $R \cup S$
 - R and S must have identical schema
- Output:
 - Has the same schema as R and S
 - Contains all rows in R and all rows in S (with duplicate rows removed)

<i>uid</i>	<i>gid</i>
123	gov
857	abc

\cup

<i>uid</i>	<i>gid</i>
123	gov
901	edf

=

<i>uid</i>	<i>gid</i>
123	gov
857	abc
901	edf

Core operator 5: Difference

- Input: two tables R and S
- Notation: $R - S$
 - R and S must have identical schema
- Output:
 - Has the same schema as R and S
 - Contains all rows in R that are not in S

$$\begin{array}{|c|c|} \hline uid & gid \\ \hline 123 & gov \\ \hline 857 & abc \\ \hline \end{array} \quad - \quad \begin{array}{|c|c|} \hline uid & gid \\ \hline 123 & gov \\ \hline 901 & edf \\ \hline \end{array} \quad = \quad \begin{array}{|c|c|} \hline uid & gid \\ \hline 857 & abc \\ \hline \end{array}$$

Derived operator 3: Intersection

- Input: two tables R and S
- Notation: $R \cap S$
 - R and S must have identical schema
- Output:
 - Has the same schema as R and S
 - Contains all rows that are in both R and S
- Shorthand for $R - (R - S)$
- Also equivalent to $S - (S - R)$
- And to $R \bowtie S$

Core operator 6: Renaming

- Input: a table R
- Notation: $\rho_S R$, $\rho_{(A_1 \rightarrow A'_1, \dots)} R$, or $\rho_{S(A_1 \rightarrow A'_1, \dots)} R$
- Purpose: “rename” a table and/or its columns
- Output: a table with the same rows as R , but called differently

Member

<i>uid</i>	<i>gid</i>
123	gov
857	abc

$\rho_{M1(uid \rightarrow uid_1, gid \rightarrow gid_1)} Member$

$M1$

<i>uid₁</i>	<i>gid₁</i>
123	gov
857	abc

Core operator 6: Renaming

- As with all other relational operators, it doesn't modify the database
 - Think of the renamed table as a copy of the original
- Used to
 - Create identical column names for natural joins
 - Example: $R(rid, \dots), S(sid, \dots)$
 - $R \bowtie_{rid=sid} S$ can be written as $(\rho_{(rid \rightarrow id)} R) \bowtie (\rho_{sid \rightarrow id} S)$
 - Avoid confusion caused by identical column names

Core operator 6: Renaming

- IDs of users who belong to at least two groups

uid	gid
100	gov
100	abc
200	gov

Member $\bowtie_?$ Member

uid	gid
100	gov
100	abc
200	gov

uid	gid	uid	gid
100	gov	100	gov
100	gov	100	abc
100	gov	200	gov
100	abc	100	gov
100	abc	100	abc
100	abc	200	gov
200	gov	100	gov
200	gov	100	abc
200	gov	200	gov

Condition 1: same uid

Condition 2: different gids

Renaming example

- IDs of users who belong to at least two groups
 $Member \bowtie_? Member$

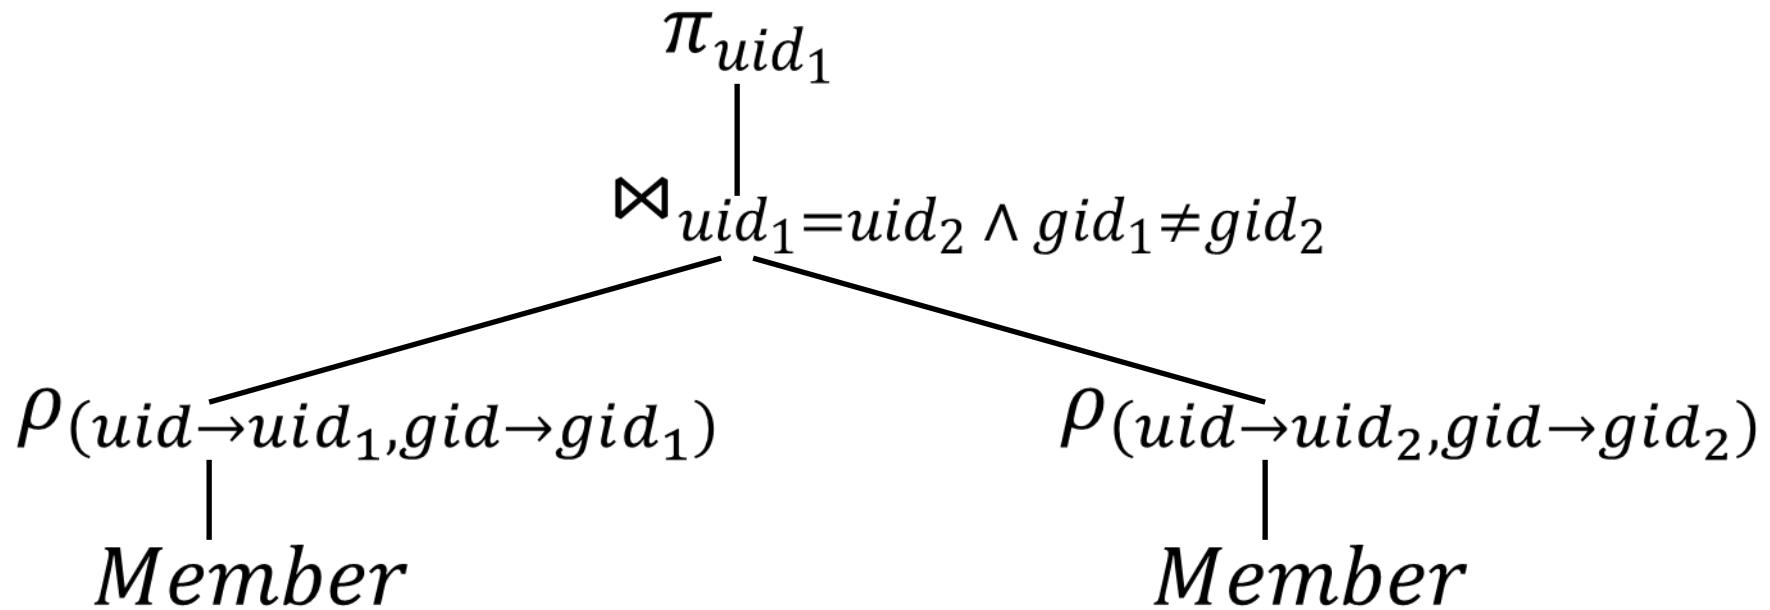
$$\pi_{uid} (Member \bowtie_{\substack{Member.uid=Member.uid \wedge Member \\ Member.gid \neq Member.gid}} Member)$$

WRON

G!

$$\pi_{uid_1} \left(\begin{array}{c} \rho_{(uid \rightarrow uid_1, gid \rightarrow gid_1)} Member \\ \bowtie_{uid_1 = uid_2 \wedge gid_1 \neq gid_2} \\ \rho_{(uid \rightarrow uid_2, gid \rightarrow gid_2)} Member \end{array} \right)$$

Expression tree notation



Take-home Exercises

- Exercise 1: IDs of groups who have **at least 2 users?**
- Exercise 2: IDs of users who belong to **at least three groups?**

Summary of operators

Core Operators

1. Selection: $\sigma_p R$
2. Projection: $\pi_L R$
3. Cross product: $R \times S$
4. Union: $R \cup S$
5. Difference: $R - S$
6. Renaming: $\rho_{S(A_1 \rightarrow A'_1, A_2 \rightarrow A'_2, \dots)} R$
Does not really add “processing” power

Note: Only use these operators for assignments & quiz

Derived Operators

1. Join: $R \bowtie_p S$
2. Natural join: $R \bowtie S$
3. Intersection: $R \cap S$

More example

User (uid int, name string, age int, pop float)
Group (gid string, name string)
Member (uid int, gid string)

- Names of users in Lisa's groups

If you have a different solution and you are uncertain about it, you can post your solution on Piazza with the lecture number and slide number; or discuss it during office hours.

More example

```
User (uid int, name string, age int, pop float)  
Group (gid string, name string)  
Member (uid int, gid string)
```

- Names of users in Lisa's groups

Writing a query bottom-up:

Their names

π_{name}

Users in
Lisa's groups

π_{uid}

User

Lisa's groups

π_{gid}

Member

Who's Lisa?

$\sigma_{name="Lisa"}$

User

Member

More example

User (uid int, name string, age int, pop float)
Group (gid string, name string)
Member (uid int, gid string)

- IDs of groups that Lisa doesn't belong to

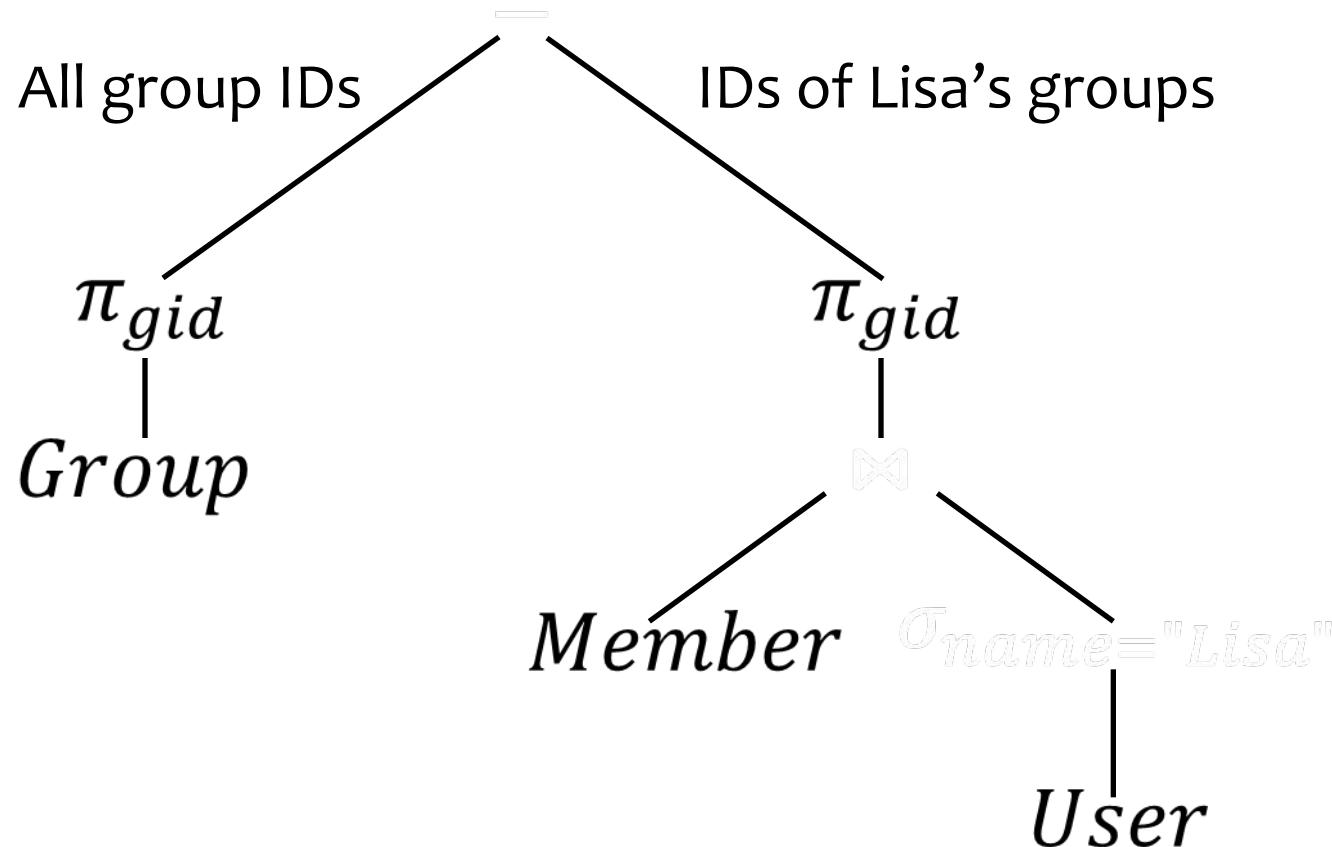
Writing a query top-down:

More example

User (uid int, name string, age int, pop float)
Group (gid string, name string)
Member (uid int, gid string)

- IDs of groups that Lisa doesn't belong to

Writing a query top-down:



A trickier example

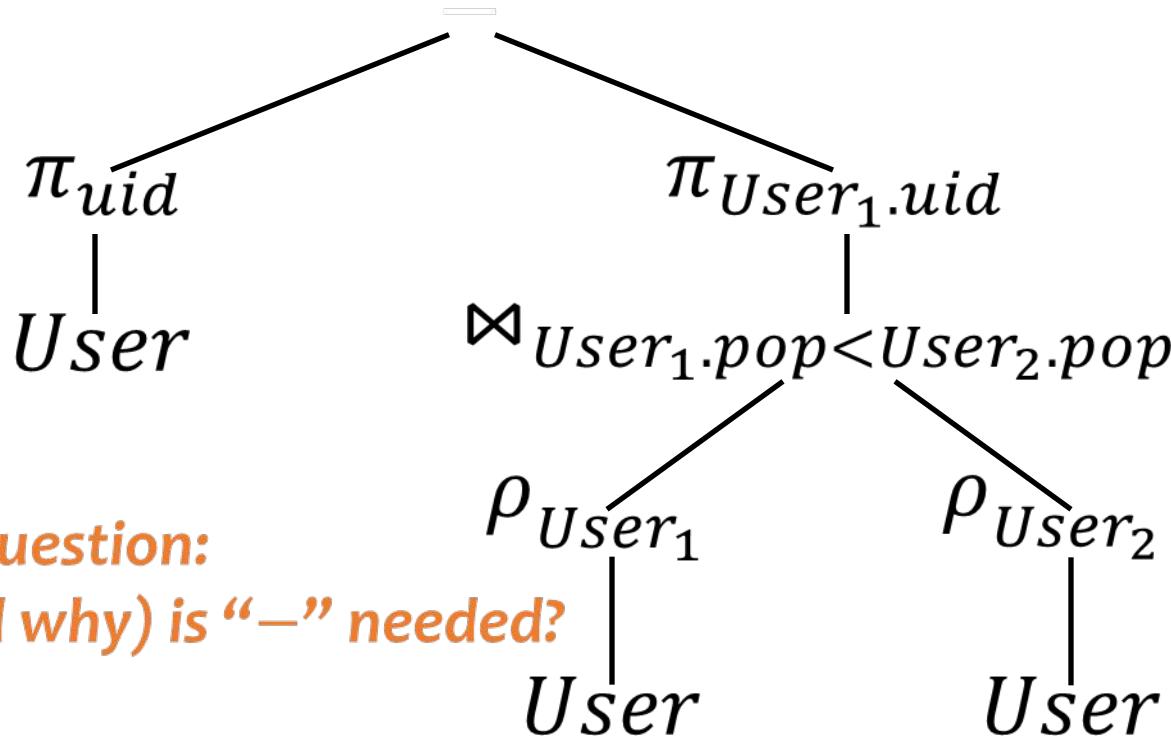
User (uid int, name string, age int, pop float)
Group (gid string, name string)
Member (uid int, gid string)

- Who are the most popular?
 - Who do NOT have the highest pop rating?
 - Whose pop is lower than somebody else's?

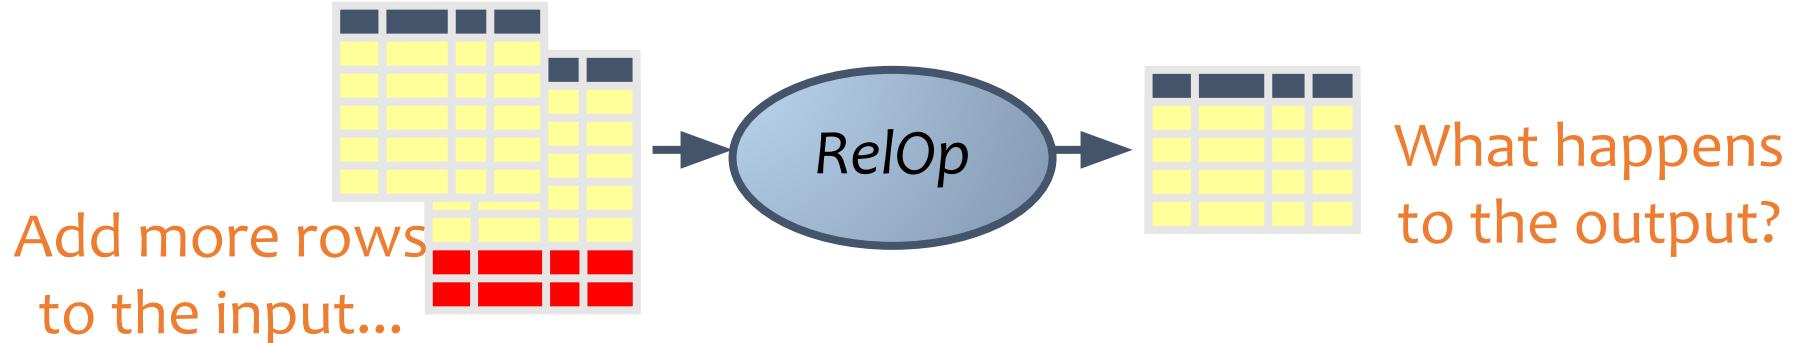
A trickier example

User (uid int, name string, age int, pop float)
Group (gid string, name string)
Member (uid int, gid string)

- Who are the most popular?
 - Who do NOT have the highest pop rating?
 - Whose pop is lower than somebody else's?



Non-monotone operators

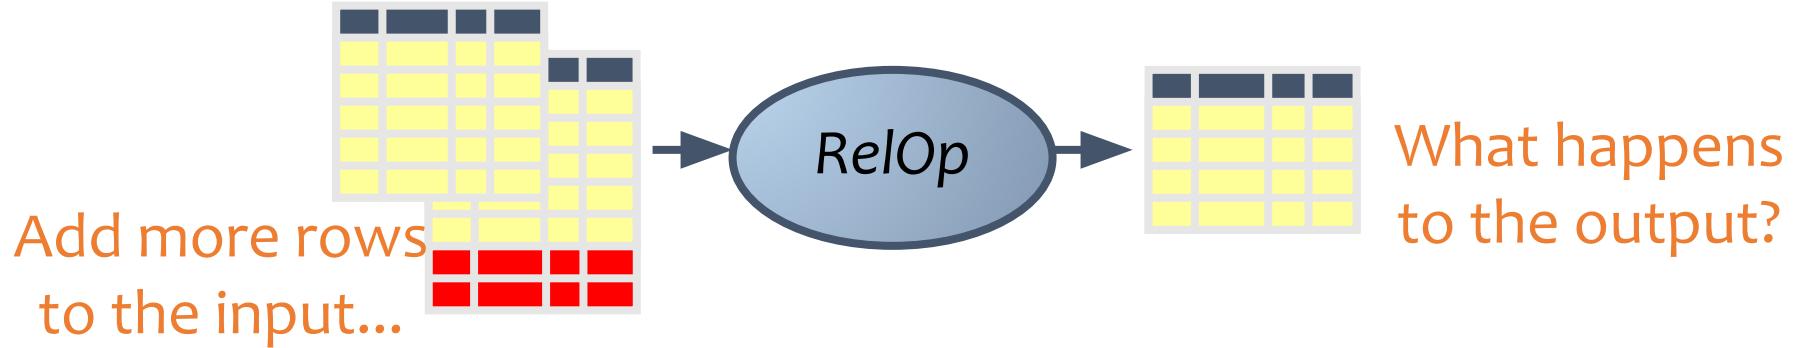


- If some **old output rows** may become **invalid**, and need to be removed → the operator is **non-monotone**
- Example: difference operator $R - S$

R	S	$= R - S$																				
<table border="1"><thead><tr><th>uid</th><th>gid</th></tr></thead><tbody><tr><td>123</td><td>gov</td></tr><tr><td>857</td><td>abc</td></tr></tbody></table>	uid	gid	123	gov	857	abc	<table border="1"><thead><tr><th>uid</th><th>gid</th></tr></thead><tbody><tr><td>123</td><td>gov</td></tr><tr><td>901</td><td>edf</td></tr><tr><td>857</td><td>abc</td></tr></tbody></table>	uid	gid	123	gov	901	edf	857	abc	<table border="1"><thead><tr><th>uid</th><th>gid</th></tr></thead><tbody><tr><td>901</td><td>edf</td></tr><tr><td>857</td><td>abc</td></tr></tbody></table>	uid	gid	901	edf	857	abc
uid	gid																					
123	gov																					
857	abc																					
uid	gid																					
123	gov																					
901	edf																					
857	abc																					
uid	gid																					
901	edf																					
857	abc																					

This old row becomes invalid because the new row added to S

Non-monotone operators



- If some **old output rows** may become **invalid**, and need to **be removed** the operator is **non-monotone**
- Otherwise (**old output rows** always remain “**correct**”) the operator is **monotone**

uid	gid
123	gov
857	abc
189	abc

–

uid	gid
123	gov
901	edf

=

uid	gid
857	abc
189	abc

This old row is always valid no matter what rows are added to R

R

S

Classification of relational operators

- Selection: $\sigma_p R$ Monotone
- Projection: $\pi_L R$ Monotone
- Cross product: $R \times S$ Monotone
- Join: $R \bowtie_p S$ Monotone
- Natural join: $R \bowtie S$ Monotone
- Union: $R \cup S$ Monotone
- Difference: $R - S$ Monotone w.r.t. R ; non-monotone w.r.t S
- Intersection: $R \cap S$ Monotone

Why is “–” needed for “highest”?

- Composition of monotone operators produces a **monotone query**
 - Old output rows remain “correct” when more rows are added to the input
- Is the “highest” query monotone? (slide 50)
 - No!
 - Current highest pop is 0.9
 - Add another row with *pop* 0.91
 - Old answer is invalidated

☞ So it must use difference!

Why do we need core operator X ?

- Difference
 - The only **non-monotone** operator
- Projection
 - The only operator that **removes columns**
- Cross product
 - The only operator that **adds columns**
- Union
 - ?
- Selection
 - ?

Extensions to relational algebra

- Duplicate handling (“bag algebra”)
- Grouping and aggregation
- “Extension” (or “extended projection”) to allow new column values to be computed

- ☞ All these will come up when we talk about SQL
- ☞ But for now we will stick to standard relational algebra without these extensions

Why is r.a. a good query language?

- Simple
 - A small set of core operators
 - Semantics are easy to grasp
- Declarative?
 - Yes, compared with older languages like CODASYL
 - Though operators do look somewhat “procedural”
- Complete?
 - With respect to what?

Relational calculus

User (uid int, name string, age int, pop float)
Group (gid string, name string)
Member (uid int, gid string)

- First-order logic
- Example: Who are the most popular?
 - $\{u.\text{uid} \mid u \in \text{User} \wedge \neg(\exists u' \in \text{User}: u.\text{pop} < u'.\text{pop})\}$, or
 - $\{u.\text{uid} \mid u \in \text{User} \wedge (\forall u' \in \text{User}: u.\text{pop} \geq u'.\text{pop})\}$

Relational calculus

- Relational algebra = “**safe**” relational calculus
 - Every query expressible as a safe relational calculus query is also expressible as a relational algebra query
 - And vice versa
- Example of an “**unsafe**” relational calculus query
 - $\{u.name \mid \neg(u \in User)\}$ → users not in the database
 - **Cannot evaluate it just by looking at the database**
- A query is **safe** if, for all database instances conforming to the schema, the query result can be computed using **only constants appearing in the database instance** or in the query itself.

Turing machine

How does relational algebra compare with a Turing machine?

- A conceptual device that can execute any computer algorithm
- Approximates what general-purpose programming languages can do
 - E.g., Python, Java, C++, ...



Alan Turing (1912-1954)

Limits of relational algebra

- Relational algebra has **no recursion**
 - Example: given relation *Friend(uid₁, uid₂)*, who can Bart reach in his social network with any number of hops?
 - Writing this query in r.a. is impossible!
 - So r.a. is not as powerful as general-purpose languages
- But why not?
 - Optimization becomes **undecidable**
 - ☞ Simplicity is empowering
 - Besides, you can always implement it at the application level, and recursion is added to SQL nevertheless!

Summary

- Part 1: Relational data model
 - Data model
 - Database schema
 - Integrity constraints (**keys**)
 - Languages (relational algebra, relational calculus, SQL)
- Part 2: Relational algebra – basic language
 - Core operators & derived operators
(**how to write a query**)
 - V.s. relational calculus
 - V.s. general programming language
- What's next?
 - SQL – query language used in practice (4 lectures)