

SQL: Part I

Introduction to Database Management

CS348 Spring 2021

We know a bit more about you

CSXXX games Working Economics
co-op listening procrastinating CS/Math Skating
go workout development Cool
cross-stitch part-time watching coop plants
Biking Drawing NBA fan golf family
cycling ML AI Cooking Bake more Read
disc disc Networks COURSES applying
play CS youtube Courses Stuff search comics
likes web walks house jobs productive
school cat movies
Networks music work Painting tennis
taking netflix sing Going Learning
big Guitar outside Neuroscience
reading journaling soccer Board
Hiking chess database electronic
Currently piano

Announcements (Tue. May 18)

- Project details (note+video) are released on Learn
 - Milestone 0 by May 24 (Mon), 11pm
 - Form **groups of 4-5** on Markus
 - Submit **report.pdf** and **code.zip** to Markus
 - Mainly for testing of the groups, not graded
- Assignment #1 is released on Learn
 - Part 1: general questions and r.a.
 - Submit a **a1part1.pdf** to Markus
 - Part 2: writing SQL on DB2 on school servers (try soon)
 - Submit **1.sql, 2.sql, ..., 6.sql** to Markus
 - Due by May 31 (Mon), 11pm

SQL

- SQL: **Structured Query Language**
 - Pronounced “S-Q-L” or “sequel”
 - The standard query language supported by most DBMS
- A brief history
 - IBM System R
 - ANSI SQL86
 - ANSI SQL89
 - ANSI SQL92 (SQL2)
 - ANSI SQL99 (SQL3)
 - ANSI SQL 2003 (added OLAP, XML, etc.)
 - ANSI SQL 2006 (added more XML)
 - ANSI SQL 2008, ...

SQL

- Data-definition language (DDL): define/modify schemas, delete relations
- Data-manipulation language (DML): query information, and insert/delete/modify tuples
- Integrity constraints: specify constraints that the data stored in the database must satisfy
- Intermediate/Advanced topics: (next week)
 - E.g., triggers, views, indexes, programming, recursive queries

this
week

DDL

User (uid int, name string, age int, pop float)
Group (gid string, name string)
Member (uid int, gid string)

- **CREATE TABLE** *table_name*
(..., *column_name column_type*, ...);

```
CREATE TABLE User(uid DECIMAL(3,0), name VARCHAR(30), age DECIMAL  
(2,0), pop DECIMAL(3,2));  
CREATE TABLE Group (gid CHAR(10), name VARCHAR(100));  
CREATE TABLE Member (uid DECIMAL (3,0), gid CHAR(10));
```

- **DROP TABLE** *table_name*;

```
DROP TABLE User;  
DROP TABLE Group;  
DROP TABLE Member;
```

How does it
work with
MySQL?

-- everything from -- to the end of line is ignored.
-- SQL is insensitive to white space.
-- SQL is insensitive to case (e.g., ...CREATE... is
-- equivalent to ...create...).

Basic queries for DML: SFW statement

- **SELECT** A_1, A_2, \dots, A_n
FROM R_1, R_2, \dots, R_m
WHERE *condition*;
- Also called an SPJ (select-project-join) query
- Corresponds to (but not really equivalent to)
relational algebra query:

$$\pi_{A_1, A_2, \dots, A_n}(\sigma_{\text{condition}}(R_1 \times R_2 \times \dots \times R_m))$$

Examples

User (uid int, name string, age int, pop float)
Group (gid string, name string)
Member (uid int, gid string)

- List all rows in the User table

```
SELECT * FROM User;
```

- * is a short hand for “all columns”

- List name of users under 18 (selection, projection)

```
SELECT name FROM User where age <18;
```

- When was Lisa born?

```
SELECT 2021-age FROM User where name = 'Lisa';
```

- SELECT list can contain expressions
 - Can also use built-in functions such as SUBSTR, ABS, etc.
 - String literals (case sensitive) are enclosed in single quotes

Example: join

User (uid int, name string, age int, pop float)
Group (gid string, name string)
Member (uid int, gid string)

- List ID's and names of groups with a user whose name contains “Simpson”

```
SELECT Group.gid, Group.name  
      FROM User, Member, Group  
     WHERE User.uid = Member.uid  
           AND Member.gid = Group.gid  
           AND ...;;
```

Example: join

User (uid int, name string, age int, pop float)
Group (gid string, name string)
Member (uid int, gid string)

- List ID's and names of groups with a user whose name **contains “Simpson”**

```
SELECT Group.gid, Group.name  
      FROM User, Member, Group  
     WHERE User.uid = Member.uid  
           AND Member.gid = Group.gid  
           AND User.name LIKE '%Simpson%';
```

- **LIKE** matches a string against a pattern
 - **%** matches any sequence of zero or more characters
- Okay to omit *table_name* in *table_name.column_name* if *column_name* is unique

Example: rename

User (uid int, name string, age int, pop float)
Group (gid string, name string)
Member (uid int, gid string)

- ID's of all pairs of users that belong to one group
 - Relational algebra query:

$$\pi_{m_1.uid, m_2.uid}$$
$$(\rho_{m_1} Member \bowtie_{m_1.gid = m_2.gid \wedge m_1.uid > m_2.uid} \rho_{m_2} Member$$

- SQL (not exactly):

```
SELECT m1.uid AS uid1, m2.uid AS uid2  
      FROM Member AS m1, Member AS m2  
     WHERE m1.gid = m2.gid  
       AND m1.uid > m2.uid;
```

- AS keyword is completely optional

A more complicated example

- Names of all groups that Lisa and Ralph are both in

Tip: Write the FROM clause first, then WHERE, and then SELECT

```
User (uid int, name string, age int, pop float)  
Group (gid string, name string)  
Member (uid int, gid string)
```

A more complicated example

- Names of all **groups that Lisa** and Ralph are both in

```
SELECT g.name  
    FROM User u1, ..., Member m1, ...  
   WHERE u1.name = 'Lisa' AND ...  
         AND u1.uid = m1.uid AND ...  
         AND ...;
```

User (uid int, name string, age int, pop float)
Group (gid string, name string)
Member (uid int, gid string)

A more complicated example

- Names of all **groups that** Lisa and **Ralph** are both in

```
SELECT g.name  
    FROM User u1, User u2, Member m1, Member m2, ...  
   WHERE u1.name = 'Lisa' AND u2.name = 'Ralph'  
         AND u1.uid = m1.uid AND u2.uid=m2.uid  
         AND ...;
```

User (uid int, name string, age int, pop float)
Group (gid string, name string)
Member (uid int, gid string)

A more complicated example

- Names of all groups that Lisa and Ralph are both in

```
SELECT g.name  
  FROM User u1, User u2, Member m1, Member m2, Group g  
 WHERE u1.name = 'Lisa' AND u2.name = 'Ralph'  
       AND u1.uid = m1.uid AND u2.uid=m2.uid  
       AND m1.gid = g.gid AND m2.gid = g.gid;
```

User (uid int, name string, age int, pop float)
Group (gid string, name string)
Member (uid int, gid string)

Why SFW statements?

- Many queries can be written using only selection, projection, and cross product (or join)
- These queries can be written in a canonical form which is captured by SFW:

$$\pi_L \left(\sigma_p (R_1 \times \dots \times R_m) \right)$$

- Example: $\pi_{R.A,S.B} (R \bowtie_{p_1} S) \bowtie_{p_2} (\pi_{T.C} \sigma_{p_3} T)$
 $= \pi_{R.A,S.B,T.C} \sigma_{p_1 \wedge p_2 \wedge p_3} (R \times S \times T)$

Set versus bag

User

uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...

```
SELECT age  
FROM User;
```

$\pi_{age} User$

age
10
8
...

Set

- No duplicates
- Relational model and algebra use set semantics

Bag

age
10
10
8
8
...

- Duplicates allowed

- Number of duplicates is significant

- SQL uses bag semantics by default

A case for bag semantics

- Efficiency
 - Saves time of eliminating duplicates

- Which one is more useful?

$\pi_{age} User$

```
SELECT age  
FROM User;
```

- The first query just returns all possible user ages
- The second query returns the user age distribution
- Besides, SQL provides the option of set semantics with **DISTINCT** keyword

Forcing set semantics

- ID's of all pairs of users that belong to one group

```
SELECT m1.uid AS uid1, m2.uid AS uid2  
      FROM Member AS m1, Member AS m2  
     WHERE m1.gid = m2.gid  
       AND m1.uid > m2.uid;
```

- Say Lisa and Ralph are in both the book club and the student government, they id pairs will appear twice

- Remove duplicate (uid1, uid2) pairs from the output

```
SELECT DISTINCT m1.uid AS uid1, m2.uid AS uid2  
      FROM Member AS m1, Member AS m2  
     WHERE m1.gid = m2.gid;  
       AND m1.uid > m2.uid;
```

Semantics of SFW

- $\text{SELECT } [\text{DISTINCT}] \ E_1, E_2, \dots, E_n$
 $\text{FROM } R_1, R_2, \dots, R_m$
 $\text{WHERE } condition;$
- For each t_1 in R_1 :
 - For each t_2 in R_2 :
 - For each t_m in R_m :
 - If $condition$ is true over t_1, t_2, \dots, t_m :
 - Compute and output E_1, E_2, \dots, E_n as a row
- If DISTINCT is present
 - Eliminate duplicate rows in output
- t_1, t_2, \dots, t_m are often called tuple variables

SQL set and bag operations

- Set: UNION, EXCEPT, INTERSECT
 - Exactly like set \cup , $-$, and \cap in relational algebra
 - Duplicates in input tables, if any, are first eliminated
 - Duplicates in result are also eliminated (for UNION)

Bag1	Bag2	(SELECT * FROM Bag1) UNION (SELECT * FROM Bag2);	(SELECT * FROM Bag1) EXCEPT (SELECT * FROM Bag2);	(SELECT * FROM Bag1) INTERSECT (SELECT * FROM Bag2);
fruit	fruit	fruit	fruit	fruit
apple	orange	apple		apple
apple	orange	orange	apple	
orange	orange		orange	orange

SQL set and bag operations

- Set: UNION, EXCEPT, INTERSECT
 - Exactly like set \cup , $-$, and \cap in relational algebra
- Bag: UNION ALL, EXCEPT ALL, INTERSECT ALL
 - Think of each row as having an implicit count (the number of times it appears in the table)



SQL set and bag operations

- Set: UNION, EXCEPT, INTERSECT
 - Exactly like set \cup , $-$, and \cap in relational algebra
- Bag: UNION ALL, EXCEPT ALL, INTERSECT ALL
 - Think of each row as having an implicit count (the number of times it appears in the table)

Bag1	Bag2	(SELECT * FROM Bag1) EXCEPT ALL (SELECT * FROM Bag2)	proper-subtract the two counts
fruit	fruit	fruit	
apple	apple	apple	
apple	orange		
orange	orange		
apple: 2 orange:1	apple: 1 orange:2	apple: 1 orange:0	

SQL set and bag operations

- Set: UNION, EXCEPT, INTERSECT
 - Exactly like set \cup , $-$, and \cap in relational algebra
- Bag: UNION ALL, EXCEPT ALL, INTERSECT ALL
 - Think of each row as having an implicit count (the number of times it appears in the table)

Bag1	Bag2	(SELECT * FROM Bag1) INTERSECT ALL (SELECT * FROM E <i>fruit</i>)	take the minimum of the two counts
<i>fruit</i>	<i>fruit</i>		
apple	apple	apple	
apple	orange	orange	
orange	orange		
apple: 2 orange:1	apple: 1 orange:2		
		apple: 1 orange:1	

Set versus bag operations

Poke (uid1, uid2, timestamp)

- uid1 poked uid2 at timestamp

Question: How do these two queries differ?

```
(SELECT uid1 FROM  
Poke)  
EXCEPT  
(SELECT uid2 FROM
```

```
(SELECT uid1 FROM  
Poke)  
EXCEPT ALL  
(SELECT uid2 FROM
```

Set versus bag operations

Poke (uid1, uid2, timestamp)

- uid1 poked uid2 at timestamp

Question: How do these two queries differ?

```
(SELECT uid1 FROM  
Poke)  
EXCEPT  
(SELECT uid2 FROM
```

```
(SELECT uid1 FROM  
Poke)  
EXCEPT ALL  
(SELECT uid2 FROM
```

Users who poked others but never got poked by others

Users who poked others more than others poked them

SQL features covered so far

- SELECT-FROM-WHERE statements
- Set and bag operations

☞ Next: how to nest SQL queries

Table subqueries

- Use query result as a table
 - In set and bag operations, FROM clauses, etc.
- Example: names of users who poked others more than others poked them

```
SELECT DISTINCT name
  FROM User,
       (SELECT uid1 FROM Poke)
 EXCEPT ALL
       (SELECT uid2 FROM Poke) AS T
 WHERE User.uid = T.uid;
```

Scalar subqueries

- A query that returns a **single row** can be used as **a value in WHERE, SELECT, etc.**
- Example: users at the same age as Bart

```
SELECT *
  FROM User,
       WHERE age = (SELECT age
                      FROM User
                     WHERE name = 'Bart');
```

- When can this query go wrong?
 - Return more than 1 row
 - Return no rows

IN subqueries

- $x \text{ IN } (\text{subquery})$ checks if x is in the result of *subquery*
- Example: users at the same age as (some) Bart

```
SELECT *
FROM User,
WHERE age IN (SELECT age
                FROM User
                WHERE name = 'Bart');
```

EXISTS subqueries

- EXISTS (*subquery*) checks if the result of *subquery* is non-empty
- Example: users at the same age as (some) Bart

```
SELECT *  
FROM User AS u,  
WHERE EXISTS (SELECT * FROM User  
WHERE name = 'Bart'  
AND age = u.age);
```

- This happens to be a correlated subquery—a subquery that references tuple variables in surrounding queries

Another example

User (uid int, name string, age int, pop float)
Group (gid string, name string)
Member (uid int, gid string)

- Users who join at least two groups

```
SELECT * FROM User u
WHERE EXISTS
  (SELECT * FROM Member m
   WHERE uid = u.uid
   AND EXISTS
     (SELECT * FROM Member
      WHERE uid = u.uid
      AND gid <> m.gid));
```

Use
table_name.
column_name
notation and AS
(renaming) to avoid
confusion

- How to find which table a column belongs to?
 - Start with the immediately surrounding query
 - If not found, look in the one surrounding that; repeat if necessary

Quantified subqueries

- Universal quantification (for all):
 - ... WHERE $x \ op \text{ ALL}(\text{subquery})$...
 - True iff for all t in the result of *subquery*, $x \ op \ t$

```
SELECT *
FROM User
WHERE pop >= ALL(SELECT pop FROM User);
```

- Existential quantification (exists):
 - ... WHERE $x \ op \text{ ANY}(\text{subquery})$...
 - True iff there exists *some* t in *subquery* result s.t. $x \ op \ t$

```
SELECT *
FROM User
WHERE NOT
    (pop < ANY(SELECT pop FROM User));
```

More ways to get the most popular

- Which users are the most popular?

```
Q1. SELECT *  
    FROM User  
    WHERE pop >= ALL(SELECT pop FROM User);
```

```
Q2. SELECT *  
    FROM User  
    WHERE NOT  
        (pop < ANY(SELECT pop FROM User));
```

```
Q3. SELECT *  
    FROM User AS u  
    WHERE NOT [EXISTS or IN?]  
        (SELECT * FROM User  
        WHERE pop > u.pop);
```

EXISTS or IN?

```
Q4. SELECT * FROM User  
    WHERE uid NOT [EXISTS or IN?]  
        (SELECT u1.uid  
        FROM User AS u1, User AS u2  
        WHERE u1.pop < u2.pop);
```

SQL features covered so far

- SELECT-FROM-WHERE statements
- Set and bag operations
- Subqueries
 - Subqueries allow queries to be written in more declarative ways (recall the “most popular” query)
 - But in many cases, they don’t add expressive power

☞ Next: aggregation and grouping

Aggregates

- Standard SQL aggregate functions: **COUNT, SUM, AVG, MIN, MAX**
- Example: number of users under 18, and their average popularity
 - COUNT(*) counts the number of rows

```
SELECT COUNT(*), AVG(pop)
  FROM User
 WHERE age <18;
```

Aggregates with DISTINCT

- Example: How many users are in some group?

```
SELECT COUNT(*)  
FROM (SELECT DISTINCT uid FROM Member);
```

Is equivalent
to

```
SELECT COUNT(DISTINCT uid)  
FROM Member;
```

Grouping

- SELECT ... FROM ... WHERE ...
GROUP BY *list_of_columns*;
- Example: compute average popularity for each age group

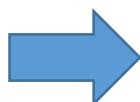
```
SELECT age, AVG(pop)  
FROM User  
GROUP BY age;
```

Example of computing GROUP BY

```
SELECT age, AVG(pop) FROM User GROUP BY age;
```

uid	name	age	pop
142	Bart	10	0.9
857	Lisa	8	0.7
123	Milhouse	10	0.2
456	Ralph	8	0.3

Compute GROUP BY: group rows according to the values of GROUP BY columns



uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3

Compute SELECT for each group



age	avg_pop
10	0.55
8	0.50

Semantics of GROUP BY

- **SELECT ... FROM ... WHERE ... GROUP BY ...;**

1. Compute FROM (\times)
2. Compute WHERE (σ)
3. Compute GROUP BY: group rows according to the values of GROUP BY columns
4. Compute SELECT for each group (π)
 - For aggregation functions with DISTINCT inputs, first eliminate duplicates within the group

☞ Number of groups =
number of rows in the final output

Aggregates with no GROUP BY

- An aggregate query with no GROUP BY clause = all rows go into one group

```
SELECT AVG(pop) FROM User;
```

Group all rows into
one group

Aggregate over
the whole group

uid	name	age	pop
142	Bart	10	0.9
857	Lisa	8	0.7
123	Milhouse	10	0.2
456	Ralph	8	0.3



uid	name	age	pop
142	Bart	10	0.9
857	Lisa	8	0.7
123	Milhouse	10	0.2
456	Ralph	8	0.3

avg_pop
0.525

Restriction on SELECT

- If a query uses aggregation/group by, then every column referenced in SELECT must be either
 - Aggregated, or
 - A GROUP BY column

Why?

- ☞ This restriction ensures that any SELECT expression produces only one value for each group

```
SELECT uid, age FROM User GROUP BY age;
```

WRON
G!

```
SELECT uid, MAX(pop) FROM User;
```

WRON
G!

HAVING

- Used to filter groups based on the group properties (e.g., aggregate values, GROUP BY column values)
- $\text{SELECT} \dots \text{FROM} \dots \text{WHERE} \dots \text{GROUP BY} \dots$
HAVING *condition*;
 1. Compute FROM (\times)
 2. Compute WHERE (σ)
 3. Compute GROUP BY: group rows according to the values of GROUP BY columns
 4. Compute HAVING (another σ over the groups)
 5. Compute SELECT (π) for each group that passes HAVING

HAVING examples

- List the average popularity for each age group with more than a hundred users

```
SELECT age, AVG(pop)
FROM User
GROUP BY age
HAVING COUNT(*)>100;
```

- Can be written using WHERE and table subqueries

```
SELECT T.age, T.apop
FROM (SELECT age, AVG(pop) AS apop, COUNT(*) AS gsize
      FROM User GROUP BY age) AS T
WHERE T.gsize>100;
```

HAVING examples

- Find average popularity for each age group over 10

```
SELECT age, AVG(pop)
FROM User
GROUP BY age
HAVING age >10;
```

- Can be written using WHERE without table subqueries

```
SELECT age, AVG(pop)
FROM User
WHERE age >10
GROUP BY age;
```

SQL features covered so far

- SELECT-FROM-WHERE statements
- Set and bag operations
- Subqueries
- Aggregation and grouping
 - More expressive power than relational algebra

☞ Next: ordering output rows

ORDER BY

- SELECT [DISTINCT] ...
FROM ... WHERE ... GROUP BY ... HAVING ...
ORDER BY *output_column* [ASC | DESC], ...;
- ASC = ascending, DESC = descending
- Semantics: After SELECT list has been computed
and optional duplicate elimination has been carried
out, **sort the output according to ORDER BY
specification**

ORDER BY example

- List all users, sort them by popularity (descending) and name (ascending)

```
SELECT uid, name, age, pop  
FROM User  
ORDER BY pop DESC, name;
```

- ASC is the default option
- Strictly speaking, only output columns can appear in ORDER BY clause (although some DBMS support more)
- Can use sequence numbers instead of names to refer to output columns: ORDER BY 4 DESC, 2;

SQL features covered so far

- Query
 - SELECT-FROM-WHERE statements
 - Set/bag (DISTINCT, UNION/EXCEPT/INTERSECT (ALL))
 - Subqueries (table, scalar, IN, EXISTS, ALL, ANY)
 - Aggregation and grouping (GROUP BY, HAVING)
 - Ordering (ORDER)
 - Outerjoins (and Nulls)
- Modification
 - INSERT/DELETE/UPDATE
- Constraints

Lecture 4