

# SQL: Triggers, Views, Indexes

Introduction to Database Management

CS348 Spring 2021

# Announcements (Tue., May 25)

- Milestone 0 - Project groups are formed by tonight!
  - Questions are welcome on Piazza
  - Watch out for updates
- Assignment #1 due next Mon (May 31, 11pm)
  - Servers:
    - ubuntu2004-002.student.cs.uwaterloo.ca
    - ubuntu2004-004.student.cs.uwaterloo.ca
  - Setup db2 environment:
    - \$ source ~cs348/public/db2profile

# SQL

- Basic SQL (queries, modifications, and constraints)
  - Intermediate SQL
    - Triggers
    - Views
    - Indexes
  - Advanced SQL
    - Programming
    - Recursive queries
- 
- this  
week

# Still remember “referential integrity”?

Example: *Member.uid* references *User.uid*

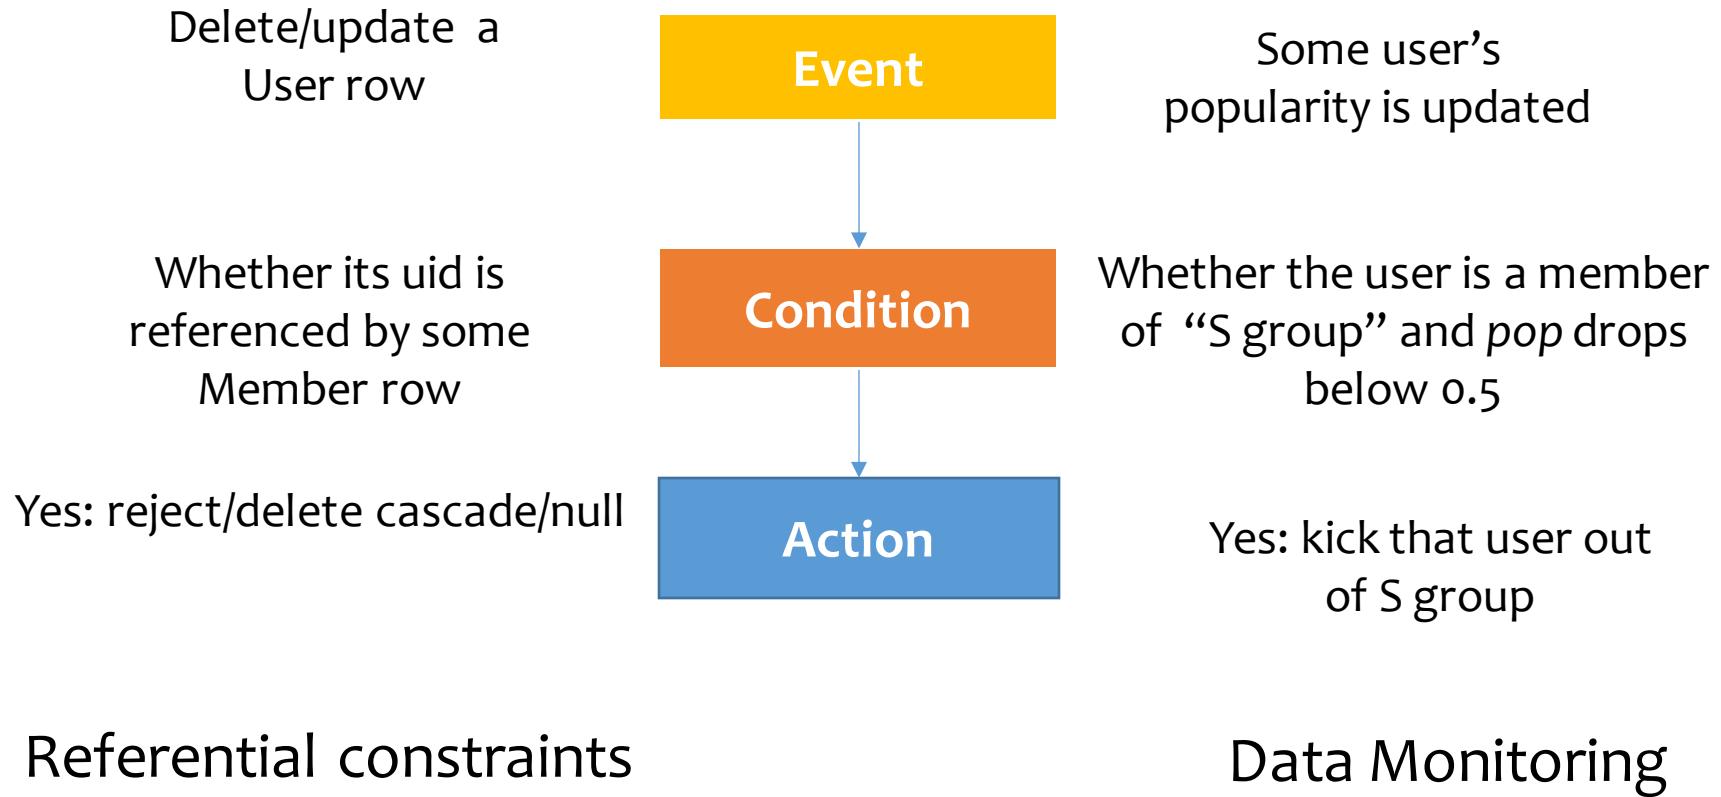
- Delete or update a *User* row whose *uid* is referenced by some *Member* row
  - Multiple Options (in SQL)

User			Member	
uid	name	...	uid	gid
142	Bart	...	142	dps
123	Milhouse	...	123	gov
857	Lisa	...	857	abc
456	Ralph	...	857	gov
789	Nelson	...	456	abc
...	...	...	456	gov
			...	....

```
CREATE TABLE Member
(uid DECIMAL(3,0) NOT NULL
REFERENCES User(uid)
ON DELETE CASCADE,
....);
```

**Option 1: Reject**  
**Option 2: Cascade**  
(ripple changes to all  
referring rows)

# Can we generalize it?



# Triggers

- A **trigger** is an event-condition-action (ECA) rule
  - When **event** occurs, test **condition**; if condition is satisfied, execute **action**

```
CREATE TRIGGER PickySGroup
AFTER UPDATE OF pop ON User
REFERENCING NEW ROW AS newUser
FOR EACH ROW
    WHEN (newUser.pop < 0.5)
        AND (newUser.uid IN (SELECT uid
                             FROM Member
                             WHERE gid = 'sgroup'))
        DELETE FROM Member
        WHERE uid = newUser.uid AND gid = 'sgroup';
```

The diagram illustrates the structure of the trigger definition. It highlights the 'Event' (the trigger type and target table), 'Transition variable' (the new row being processed), 'Condition' (the logic that triggers the action), and 'Action' (the delete operation). The code uses standard SQL syntax with some specific annotations.

# Trigger option 1 – possible events

- Possible events include:

- **INSERT ON table; DELETE ON table; UPDATE [OF column] ON table**

```
CREATE TRIGGER PickySGroup
AFTER UPDATE OF pop ON User
REFERENCING NEW ROW AS newUser
FOR EACH ROW
    WHEN (newUser.pop < 0.5)
        AND (newUser.uid IN (SELECT uid
                             FROM Member
                             WHERE gid = 'sgroup'))
        DELETE FROM Member
        WHERE uid = newUser.uid AND gid = 'sgroup';
```

**Event**

**Condition**

**Action**

# Trigger option 2 – timing

- Timing—action can be executed:
  - AFTER or BEFORE the triggering event
  - INSTEAD OF the triggering event on views (more later)

```
CREATE TRIGGER NoFountainOfYouth  
BEFORE UPDATE OF age ON User  
REFERENCING OLD ROW AS o, NEW ROW AS n  
FOR EACH ROW  
    WHEN (n.age < o.age)  
        SET n.age = o.age;
```

The diagram illustrates the structure of the trigger definition. It uses orange brackets and labels to identify specific parts of the code:

- Event:** Points to the line "BEFORE UPDATE OF age ON User".
- Condition:** Points to the line "WHEN (n.age < o.age)".
- Action:** Points to the line "SET n.age = o.age;".

# Trigger option 3 – granularity

- Granularity—trigger can be activated:
  - FOR EACH ROW modified

```
CREATE TRIGGER PickySGroup
AFTER UPDATE OF pop ON User
REFERENCING NEW ROW AS newUser
FOR EACH ROW
WHEN (newUser.pop < 0.5)
      AND (newUser.uid IN (SELECT uid
                           FROM Member
                           WHERE gid = 'sgroup'))
DELETE FROM Member
WHERE uid = newUser.uid AND gid = 'sgroup';
```

Event

Condition

Action

# Trigger option 3 – granularity

- Granularity—trigger can be activated:
  - FOR EACH ROW modified
  - FOR EACH STATEMENT that performs modification

```
CREATE TRIGGER PickySGroup2
AFTER UPDATE OF pop ON User
REFERENCING NEW TABLE AS newUsers
FOR EACH STATEMENT
    DELETE FROM Member
        WHERE gid = 'sgroup'
        AND uid IN (SELECT uid
                     FROM newUsers
                     WHERE pop < 0.5);
```

Event

Transition table:  
contains all the  
affected rows

Condition  
& Action

# Trigger option 3 – granularity

- Granularity—trigger can be activated:
  - **FOR EACH ROW** modified
  - **FOR EACH STATEMENT** that performs modification

```
CREATE TRIGGER PickySGroup2
AFTER UPDATE OF pop ON User
REFERENCING NEW TABLE AS newUsers
FOR EACH STATEMENT
```

```
    DELETE FROM Member
        WHERE gid = 'sgroup'
        AND uid IN (SELECT uid
```

```
            FROM newUsers
            WHERE pop < 0.5);
```

Transition table:  
contains all the  
affected rows

Only can be used  
with **AFTER**  
triggers

# Transition variables/tables

- **OLD ROW**: the modified row before the triggering event
- **NEW ROW**: the modified row after the triggering event
- **OLD TABLE**: a hypothetical read-only table containing all rows to be modified before the triggering event
- **NEW TABLE**: a hypothetical table containing all modified rows after the triggering event

Event	Row	Statement
Delete	old r; old t	old t
Insert	new r; new t	new t
Update	old/new r; old/new t	old/new t

AFTER Trigger

Event	Row	Statement
Update	old/new r	-
Insert	new r	-
Delete	old r	-

BEFORE Trigger

# Statement- vs. row-level triggers

- Simple row-level triggers are easier to implement
  - Statement-level triggers: require significant amount of state to be maintained in OLD TABLE and NEW TABLE
- Exercise 1: However, can you think of a case when a row-level trigger may be less efficient?
- Exercise 2: Certain triggers are only possible at statement level. Can you think of an example?

# System issues

- Recursive firing of triggers
  - Action of one trigger causes another trigger to fire
  - Can get into an infinite loop
- Interaction with constraints (tricky to get right!)
  - When to check if a triggering event violates constraints?
    - After a BEFORE trigger
    - Before an AFTER trigger
    - (based on db2, other DBMS may differ)
- Be best avoided when alternatives exist

# SQL features covered so far

- Basic SQL
- Intermediate SQL
  - Triggers
  - Views

# Views

- A **view** is like a “virtual” table
  - Defined by a query, which describes **how to compute the view contents on the fly**
  - Stored by DBMS instead of view contents
  - Can be used in queries just like a regular table

```
CREATE VIEW SGroup AS  
    SELECT * FROM User  
    WHERE uid IN (SELECT uid  
                  FROM Member  
                  WHERE gid = 'sgroup');
```

Base  
tables

```
SELECT AVG(pop) FROM SGroup;
```

```
SELECT MIN(pop) FROM SGroup;
```

```
SELECT ... FROM SGroup;
```

```
SELECT AVG(pop)  
FROM (SELECT * FROM User  
      WHERE uid IN  
          (SELECT uid FROM Member  
          WHERE gid = 'jes'))  
AS SGroup;
```

```
DROP VIEW SGroup;
```



# Why use views?

- To hide complexity from users
- To hide data from users
- Logical data independence
- To provide a uniform interface for different implementations or sources

# Modifying views

- Does it even make sense, since views are virtual?
- It does make sense if we want users to really see views as tables
- Goal: modify the base tables such that the modification would appear to have been accomplished on the view

# A simple case

```
CREATE VIEW UserPop AS  
    SELECT uid, pop FROM User;
```

```
DELETE FROM UserPop WHERE uid = 123;
```

translates to:

```
DELETE FROM User WHERE uid = 123;
```

# An impossible case

```
CREATE VIEW PopularUser AS  
    SELECT uid, pop FROM User  
    WHERE pop >= 0.8;
```

```
INSERT INTO PopularUser VALUES(987, 0.3);
```

- No matter what we do on *User*, the inserted row will not be in *PopularUser*

# A case with too many possibilities

```
CREATE VIEW AveragePop(pop) AS  
    SELECT AVG(pop) FROM User;
```

Renamed  
column

```
UPDATE AveragePop SET pop = 0.5;
```

- Set everybody's *pop* to 0.5?
- Adjust everybody's *pop* by the same amount?
- Just lower one user's *pop*?

# SQL92 updateable views

- More or less just single-table selection queries
  - No join
  - No aggregation
  - No subqueries
- Arguably somewhat restrictive
- Still might get it wrong in some cases
  - See the slide titled “An impossible case” (slide 20)
  - Adding **WITH CHECK OPTION** to the end of the view definition will make DBMS reject such modifications

# INSTEAD OF triggers for views

```
CREATE VIEW AveragePop(pop) AS  
    SELECT AVG(pop) FROM User;
```

```
CREATE TRIGGER AdjustAveragePop  
INSTEAD OF UPDATE ON AveragePop  
REFERENCING OLD ROW AS o,  
    NEW ROW AS n  
FOR EACH ROW  
    UPDATE User  
    SET pop = pop + (n.pop-o.pop);
```

- What does this trigger do?

```
UPDATE AveragePop SET pop = 0.5;
```

# INSTEAD OF triggers for views

```
CREATE VIEW AveragePop(pop) AS  
    SELECT AVG(pop) FROM User;
```

```
CREATE TRIGGER AdjustAveragePop  
INSTEAD OF UPDATE ON AveragePop  
REFERENCING OLD ROW AS o,  
        NEW ROW AS n  
FOR EACH ROW  
    UPDATE User  
    SET pop = pop + (n.pop-o.pop);
```

User

...	pop	...
	0.4	+0.1
	0.4	+0.1
	0.6	+0.1
	0.3	+0.1

- What does this trigger do?

```
UPDATE AveragePop SET pop = 0.5;
```

# SQL features covered so far

- Basic SQL
- Intermediate SQL
  - Triggers
  - Views
  - Indexes

# Motivating examples of using indexes

```
SELECT * FROM User WHERE name = 'Bart';
```

- Can we go “directly” to rows with `name='Bart'` instead of scanning the entire table?  
→ index on `User.name`

```
SELECT * FROM User, Member  
WHERE User.uid = Member.uid AND Member.gid = 'sgroup';
```

- Can we find relevant `Member` rows “directly”?  
→ index on `Member.gid` or `(gid, uid)`
- For each relevant `Member` row, can we “directly” look up `User` rows with matching `uid`  
→ index on `User.uid`

# Indexes

- An **index** is an auxiliary persistent data structure
  - Search tree (e.g., B<sup>+</sup>-tree), lookup table (e.g., hash table), etc.
  - ☞ More on indexes later in this course!
- **CREATE [UNIQUE] INDEX *indexname* ON *tablename*(*columnname*<sub>1</sub>, ..., *columnname*<sub>*n*</sub>);**
  - With UNIQUE, the DBMS will also enforce that  $\{\text{columnname}_1, \dots, \text{columnname}_n\}$  is a key of *tablename*
- **DROP INDEX *indexname*;**
- Typically, the DBMS will automatically create indexes for PRIMARY KEY and UNIQUE constraint declarations

# Indexes

- An index on  $R.A$  can speed up accesses of the form
  - $R.A = value$
  - $R.A > value$  (sometimes; depending on the index type)
- An index on  $(R.A_1, \dots, R.A_n)$  can speed up
  - $R.A_1 = value_1 \wedge \dots \wedge R.A_n = value_n$
  - $(R.A_1, \dots, R.A_n) > (value_1, \dots, value_n)$  (again depends)

## Questions (lecture 12):

- ☞ Ordering of index columns is important—is an index on  $(R.A, R.B)$  equivalent to one on  $(R.B, R.A)$ ?
- ☞ How about an index on  $R.A$  plus another on  $R.B$ ?
- ☞ More indexes = better performance?

# SQL features covered so far

## Basic & Intermediate SQL

- Query
- Modification
- Constraints
- Triggers
- Views
- Indexes

☞ Next: Recursion, programming