

SQL: Recursion, Programming

Introduction to Database Management

CS348 Spring 2021


SQL

- Basic SQL (queries, modifications, and constraints)
- Intermediate SQL
 - Triggers
 - Views
 - Indexes
- Advanced SQL
 - Recursion
 - Programming

A motivating example

Parent (parent, child)

parent	child
Homer	Bart
Homer	Lisa
Marge	Bart
Marge	Lisa
Abe	Homer
Ape	Abe



- Example: find Bart's ancestors
- “Ancestor” has a recursive definition
 - X is Y 's ancestor if
 - X is Y 's parent, or
 - X is Z 's ancestor and Z is Y 's ancestor

Recursion in SQL

- SQL2 had no recursion
 - You can find Bart's parents, grandparents, great grandparents, etc.

```
SELECT p1.parent AS grandparent  
FROM Parent p1, Parent p2  
WHERE p1.child = p2.parent  
AND p2.child = 'Bart';
```

- But you cannot find all his ancestors with a single query
- SQL3 introduces recursion
 - **WITH** clause
 - Implemented in PostgreSQL (**common table expressions**)

Ancestor query in SQL3

WITH RECURSIVE

Ancestor(anc, desc) AS

((SELECT parent, child FROM Parent))

base case

UNION

(SELECT a1.anc, a2.desc

a1.anc (X) \rightarrow a1.desc(Z)
a2.anc (Z) \rightarrow a2.desc (Y)

FROM Ancestor a1, Ancestor a2

WHERE a1.desc = a2.anc))

Define
a relation
recursively

SELECT anc

recursion step

FROM Ancestor

Query using the relation
defined in WITH clause

WHERE desc = 'Bart';

Finding ancestors


WITH RECURSIVE

```
Ancestor(anc, desc) AS base case
  ((SELECT parent, child FROM Parent)
UNION
  (SELECT a1.anc, a2.desc
   FROM Ancestor a1, Ancestor a2
   WHERE a1.desc = a2.anc))
   ....;
```

recursive step

anc	desc
Homer	Bart
Homer	Lisa
Marge	Bart
Marge	Lisa
Abe	Homer
Ape	Abe

parent	child
Homer	Bart
Homer	Lisa
Marge	Bart
Marge	Lisa
Abe	Homer
Ape	Abe



anc	desc
Homer	Bart
Homer	Lisa
Marge	Bart
Marge	Lisa
Abe	Homer
Ape	Abe
Abe	Bart
Abe	Lisa
Ape	Homer
Ape	Bart
Ape	Lisa

Fixed point of a function

- If $f: D \rightarrow D$ is a function from a type D to itself, a **fixed point** of f is a value x such that $f(x) = x$
 - Example: what is the fixed point of $f(x) = x/2$?
 - Ans: 0, as $f(0)=0$ With seed 1: $1, 1/2, 1/4, 1/8, 1/16, \dots \rightarrow 0$
- To compute a fixed point of f
 - Start with a “seed”: $x \leftarrow x_0$
 - Compute $f(x)$
 - If $f(x) = x$, stop; x is fixed point of f
 - Otherwise, $x \leftarrow f(x)$; repeat

Fixed point of a query

- A query q is just a function that maps an input table to an output table, so a **fixed point** of q is a table T such that $q(T) = T$
- To compute fixed point of q
 - Start with an empty table: $T \leftarrow \emptyset$
 - Evaluate q over T
 - If the result is identical to T , stop; T is a fixed point
 - Otherwise, let T be the new result; repeat

Non-linear v.s. linear recursion

- Non-linear

```
WITH RECURSIVE Ancestor(anc, desc) AS
((SELECT parent, child FROM Parent)
UNION
(SELECT a1.anc, a2.desc
FROM Ancestor a1, Ancestor a2
WHERE a1.desc = a2.anc)) .....;
```

- Linear: a recursive definition can make only one reference to itself

```
WITH RECURSIVE Ancestor2(anc, desc) AS
((SELECT parent, child FROM Parent)
UNION
(SELECT anc, child
FROM Ancestor, Parent
WHERE desc = parent))
```

Linear vs. non-linear recursion

- Linear recursion is easier to implement
 - For linear recursion, just keep joining newly generated Ancestor rows with *Parent*
 - For non-linear recursion, need to join newly generated Ancestor rows with **all existing Ancestor rows**
- Non-linear recursion may take fewer steps to converge, but perform more work
 - Example: Given $a \rightarrow b \rightarrow c \rightarrow d \rightarrow e$, i.e., a is parent of b, b is parent of c, ..., d is parent of e.
 - The **linear recursion** takes 4 steps to find (a, e) is an ancestor-descendant pair (slide 9, Ancestor2)
 - Question: How about **non-linear recursion**? (slide 9, Ancestor)

Mutual recursion example

- Table *Natural* (*n*) contains 1, 2, ..., 100
- Which numbers are even/odd?
 - An even number plus 1 is an odd number
 - An odd number plus 1 is an even number
 - 1 is an odd number

```
WITH RECURSIVE Even(n) AS
  (SELECT n FROM Natural
   WHERE n = ANY(SELECT n+1 FROM Odd)),
RECURSIVE Odd(n) AS
  (SELECT n FROM Natural WHERE n = 1)
UNION
  (SELECT n FROM Natural
   WHERE n = ANY(SELECT n+1 FROM Even))
```

Computing mutual recursion

- $\text{Even} = \emptyset, \text{Odd} = \emptyset$
- $\text{Even} = \emptyset, \text{Odd} = \{1\}$
- $\text{Even} = \{2\}, \text{Odd} = \{1\}$
- $\text{Even} = \{2\}, \text{Odd} = \{1, 3\}$
- $\text{Even} = \{2, 4\}, \text{Odd} = \{1, 3\}$
- $\text{Even} = \{2, 4\}, \text{Odd} = \{1, 3, 5\}$
- ...

```
WITH RECURSIVE Even(n) AS
  (SELECT n FROM Natural
   WHERE n = ANY(SELECT n+1 FROM Odd)),
RECURSIVE Odd(n) AS
  ((SELECT n FROM Natural WHERE n = 1)
   UNION
   (SELECT n FROM Natural
    WHERE n = ANY(SELECT n+1 FROM Even
```

Semantics of WITH

- WITH RECURSIVE R_1 AS Q_1 ,
 ...,
 RECURSIVE R_n AS Q_n
 Q ;
 - Q and Q_1, \dots, Q_n may refer to R_1, \dots, R_n
- Semantics
 1. $R_1 \leftarrow \emptyset, \dots, R_n \leftarrow \emptyset$
 2. Evaluate Q_1, \dots, Q_n using the current contents of R_1, \dots, R_n :
 $R_1^{new} \leftarrow Q_1, \dots, R_n^{new} \leftarrow Q_n$
 3. If $R_i^{new} \neq R_i$ for some i
 - 3.1. $R_1 \leftarrow R_1^{new}, \dots, R_n \leftarrow R_n^{new}$
 - 3.2. Go to 2.
 4. Compute Q using the current contents of R_1, \dots, R_n and output the result

Starting with non-empty set

WITH RECURSIVE

```
Ancestor(anc, desc) AS
  base case
  ((SELECT parent, child FROM Parent)
UNION
  (SELECT a1.anc, a2.desc
  recursive
  FROM Ancestor a1, Ancestor a2
  WHERE a1.desc = a2.anc))
  ....;
```

anc	desc
Bogus	Bogus
Homer	Bart
Homer	Lisa
Marge	Bart
Marge	Lisa
Abe	Homer
Ape	Abe
Bogus	Bogus

parent	child
Homer	Bart
Homer	Lisa
Marge	Bart
Marge	Lisa
Abe	Homer
Ape	Abe

anc	desc
Homer	Bart
Homer	Lisa
Marge	Bart
Marge	Lisa
Abe	Homer
Ape	Abe

anc	desc
Homer	Bart
Homer	Lisa
Marge	Bart
Marge	Lisa
Abe	Homer
Ape	Abe
Abe	Bart
Abe	Lisa
Ape	Homer
Ape	Bart
Ape	Lisa
Bogus	Bogus

Fixed points are not unique

WITH RECURSIVE

Ancestor(anc, desc) AS

((SELECT parent, child FROM Parent)

UNION

(SELECT a1.anc, a2.desc

FROM Ancestor a1, Ancestor a2

WHERE a1.desc = a2.anc))

.....;

parent	child
Homer	Bart
Homer	Lisa
Marge	Bart
Marge	Lisa
Abe	Homer
Ape	Abe
Abe	Bart
Abe	Lisa
Ape	Homer
Ape	Bart
Ape	Lisa
Bogus	Bogus

anc	desc
Homer	Bart
Homer	Lisa
Marge	Bart
Marge	Lisa
Abe	Homer
Ape	Abe
Abe	Bart
Abe	Lisa
Ape	Homer
Ape	Bart
Ape	Lisa
Bogus	Bogus

Lecture 2

Note how the bogus tuple reinforces itself!

- If q is monotone, then starting from \emptyset produces the unique minimal fixed point
 - All these fixed points must contain this fixed point
→ the unique minimal fixed point is the “natural” answer

Mixing negation with recursion


- If q is non-monotone
 - The fixed-point iteration may never converge
 - There could be multiple minimal fixed points
- Example: popular users ($\text{pop} \geq 0.8$) join either SGroup or PGroup
 - Those not in SGroup should be in PGroup
 - Those not in GGroup should be in SGroup

```
WITH RECURSIVE PGroup(uid) AS
  (SELECT uid FROM User WHERE pop >= 0.8
   AND uid NOT IN (SELECT uid FROM SGroup)),
RECURSIVE SGroup(uid) AS
  (SELECT uid FROM User WHERE pop >= 0.8
   AND uid NOT IN (SELECT uid FROM PGroup))
```

Fixed-point iter may not converge

```
WITH RECURSIVE PGroup(uid) AS
  (SELECT uid FROM User WHERE pop >= 0.8
   AND uid NOT IN (SELECT uid FROM SGroup)),
RECURSIVE SGroup(uid) AS
  (SELECT uid FROM User WHERE pop >= 0.8
   AND uid NOT IN (SELECT uid FROM PGroup))
```


uid	name	age	pop
142	Bart	10	0.9
121	Allison	8	0.85



Multiple minimal fixed points



```
WITH RECURSIVE PGroup(uid) AS
  (SELECT uid FROM User WHERE pop >= 0.8
   AND uid NOT IN (SELECT uid FROM SGroup)),
RECURSIVE SGroup(uid) AS
  (SELECT uid FROM User WHERE pop >= 0.8
   AND uid NOT IN (SELECT uid FROM PGroup))
```

uid	name	age	pop
142	Bart	10	0.9
121	Allison	8	0.85



Legal mix of negation and recursion


- Construct a **dependency graph**
 - One node for each table defined in WITH
 - A directed edge $R \rightarrow S$ if R is defined in terms of S
 - Label the directed edge “ $-$ ” if the query defining R is not monotone with respect to S
- Legal SQL3 recursion: no cycle with a “ $-$ ” edge
 - Called **stratified negation**
- Bad mix: a cycle with at least one edge labeled “ $-$ ”



Stratified negation example


- Find pairs of persons with no common ancestors

```
WITH RECURSIVE Ancestor(anc, desc) AS
((SELECT parent, child FROM Parent) UNION
 (SELECT a1.anc, a2.desc
  FROM Ancestor a1, Ancestor a2
  WHERE a1.desc = a2.anc)),
RECURSIVE Person(person) AS
((SELECT parent FROM Parent) UNION
 (SELECT child FROM Parent)),
RECURSIVE NoCommonAnc(person1, person2) AS
((SELECT p1.person, p2.person
  FROM Person p1, Person p2
  WHERE p1.person <> p2.person)
EXCEPT
(SELECT a1.desc, a2.desc
  FROM Ancestor a1, Ancestor a2
  WHERE a1.anc = a2.anc))
SELECT * FROM NoCommonAnc;
```



Evaluating stratified negation

- The **stratum** of a node R is the maximum number of “–” edges on any path from R
 - Ancestor: stratum 0
 - Person: stratum 0
 - NoCommonAnc: stratum 1
 - Evaluation strategy
 - Compute tables lowest-stratum first
 - For each stratum, use fixed-point iteration on all nodes in that stratum
 - Stratum 0: Ancestor and Person
 - Stratum 1: NoCommonAnc
- ☞ Intuitively, there is **no negation within each stratum**



SQL features covered so far

- Basic SQL (queries, modifications, and constraints)
- Intermediate SQL(triggers, views, indexes)
- Recursion
 - SQL3 WITH recursive queries
 - Solution to a recursive query (with no negation)
 - Mixing negation and recursion is tricky
- Programming

Motivation

- Pros and cons of SQL
 - Very high-level, possible to optimize
 - Not intended for general-purpose computation
- Solutions
 - Augment SQL with constructs from general-purpose programming languages
 - E.g.: SQL/PSM
 - Use SQL together with general-purpose programming languages: many possibilities
 - Through an API, e.g., Python psycopg2
 - Embedded SQL, e.g., in C
 - Automatic object-relational mapping, e.g.: Python SQLAlchemy
 - Extending programming languages with SQL-like constructs, e.g.: LINQ

An “impedance mismatch”

- SQL operates on **a set of records at a time**
- Typical low-level general-purpose programming languages operate on **one record at a time**

☞ Solution: **cursor**

- **Open** (a result table), **Get next**, **Close**
- ☞ Found in virtually every database language/API
 - With slightly different syntaxes

Augmenting SQL: SQL/PSM

- PSM = Persistent Stored Modules
- `CREATE PROCEDURE proc_name(param_decls)
local_decls
proc_body;`
- `CREATE FUNCTION func_name(param_decls)
RETURNS return_type
local_decls
func_body;`
- `CALL proc_name(params);`
- Inside procedure body:
`SET variable = CALL func_name(params);`

SQL/PSM example

```
CREATE FUNCTION SetMaxPop(IN newMaxPop FLOAT)
RETURNS INT
-- Enforce newMaxPop; return # rows modified.
BEGIN
DECLARE rowsUpdated INT DEFAULT 0;
DECLARE thisPop FLOAT;
    -- A cursor to range over all users:
DECLARE userCursor CURSOR FOR
    SELECT pop FROM User
FOR UPDATE;
    -- Set a flag upon “not found” exception:
DECLARE noMoreRows INT DEFAULT 0;
DECLARE CONTINUE HANDLER FOR NOT FOUND
    SET noMoreRows = 1;
        ... (see next slide) ...
RETURN rowsUpdated;
END
```

Declare
local
variables

SQL/PSM example continued

```
-- Fetch the first result row:  
OPEN userCursor;  
FETCH FROM userCursor INTO thisPop;  
  
-- Loop over all result rows:  
WHILE noMoreRows <> 1 DO  
    IF thisPop > newMaxPop THEN  
        -- Enforce newMaxPop:  
        UPDATE User SET pop = newMaxPop  
        WHERE CURRENT OF userCursor;  
        -- Update count:  
        SET rowsUpdated = rowsUpdated + 1;  
    END IF;  
    -- Fetch the next result row:  
    FETCH FROM userCursor INTO thisPop;  
END WHILE;  
CLOSE userCursor;
```

Function
body

Other SQL/PSM features

- Assignment using scalar query results
 - SELECT INTO
 - Other loop constructs
 - FOR, REPEAT UNTIL, LOOP
 - Flow control
 - GOTO
 - Exceptions
 - SIGNAL, RESIGNAL
- ...
- For more PostgreSQL-specific information, look for “PL/pgSQL” in PostgreSQL documentation
 - <https://www.postgresql.org/docs/9.6/plpgsql.html>

Working with SQL through an API

- E.g.: Python psycopg2, JDBC, ODBC (C/C++/VB)
 - All based on the SQL/CLI (Call-Level Interface) standard
- The application program sends SQL commands to the DBMS at runtime
- Responses/results are converted to objects in the application program

Example API: Python psycopg2

```
import psycopg2
conn = psycopg2.connect(dbname='beers')
cur = conn.cursor()
# list all drinkers:
cur.execute('SELECT * FROM Drinker')
for drinker, address in cur:
    print(drinker + ' lives at ' + address)
# print menu for bars whose name contains "a":
cur.execute('SELECT * FROM Serves WHERE bar LIKE %s', ('%a%',))
for bar, beer, price in cur:
    print('{0} serves {1} at ${2:.2f}'.format(bar, beer, price))
cur.close()
conn.close()
```

You can iterate over `cur` one tuple at a time

Placeholder for query parameter

Tuple of parameter values, one for each `%s`

More psycopg2 examples

```
# “commit” each change immediately—need to set this option just once at  
the start of the session
```

```
conn.set_session(autocommit=True)
```

```
# ...
```

```
bar = input('Enter the bar to update: ').strip()
```

```
beer = input('Enter the beer to update: ').strip()
```

```
price = float(input('Enter the new price: '))
```

```
try:
```

```
    cur.execute("""  
        UPDATE Serves  
        SET price = %s  
        WHERE bar = %s AND beer = %s""", (price, bar, beer))
```

```
if cur.rowcount != 1:
```

```
    print('{} row(s) updated: correct bar/beer?\n'.format(cur.rowcount))
```

```
except Exception as e:
```

```
    print(e)
```

Perform parsing,
semantic analysis,
optimization,
compilation, and finally
execution

More psycopg2 examples

```
....  
while true:  
# Input bar, beer, price...
```

```
    cur.execute(  
        "UPDATE Serves  
        SET price = %s  
        WHERE bar = %s AND beer = %s'", (price, bar, beer))
```

```
....  
# Check result...
```

Perform passing,
semantic analysis,
optimization,
compilation, and finally
execution

Execute many times
Can we reduce this overhead?

Prepared statements: example

```
cur.execute("")          # Prepare once (in SQL).  Prepare only once  
PREPARE update_price AS    # Name the prepared plan,  
UPDATE Serves  
SET price = $1            # and note the $1, $2, ... notation for  
WHERE bar = $2 AND beer = $3") # parameter placeholders.
```


while true:

```
# Input bar, beer, price...
```

```
cur.execute('  
EXECUTE update_price(%s, %s, %s)',\ # Execute many times.  
(price, bar, beer))....
```

```
# Check result...
```

“Exploits of a mom”



- The school probably had something like:

```
cur.execute("SELECT * FROM Students " + \  
           "WHERE (name = '" + name + "')")
```

where **name** is a string input by user

- Called an **SQL injection attack**

Guarding against SQL injection

- Escape certain characters in a user input string, to ensure that it remains a single string
 - E.g., ', which would terminate a string in SQL, must be replaced by " (two single quotes in a row) within the input string
- Luckily, most API's provide ways to “sanitize” input automatically (if you use them properly)
 - E.g., pass parameter values in psycopg2 through %s's

Augmenting SQL vs. API

- Pros of augmenting SQL:
 - More processing features for DBMS
 - More application logic can be pushed closer to data
- Cons of augmenting SQL:
 - SQL is already too big
 - Complicate optimization and make it impossible to guarantee safety

A brief look at other approaches

- “Embed” SQL in a general-purpose programming language
 - E.g.: embedded SQL
- Support database features through an object-oriented programming language
 - E.g., object-relational mappers (ORM) like Python SQLAlchemy
- Extend a general-purpose programming language with SQL-like constructs
 - E.g.: LINQ (Language Integrated Query for .NET)

Embedding SQL in a language

Example in C

```
EXEC SQL BEGIN DECLARE SECTION;  
int thisUid; float thisPop;  
EXEC SQL END DECLARE SECTION;  
EXEC SQL DECLARE ABCMember CURSOR FOR  
    SELECT uid, pop FROM User  
    WHERE uid IN (SELECT uid FROM Member WHERE gid = 'abc')  
    FOR UPDATE;  
EXEC SQL OPEN ABCMember;  
EXEC SQL WHENEVER NOT FOUND DO break;  
while (1) {  
    EXEC SQL FETCH ABCMember INTO :thisUid, :thisPop;  
    printf("uid %d: current pop is %f\n", thisUid, thisPop);  
    printf("Enter new popularity: ");  
    scanf("%f", &thisPop);  
    EXEC SQL UPDATE User SET pop = :thisPop  
        WHERE CURRENT OF ABCMember;  
}  
EXEC SQL CLOSE ABCMember;
```

} Declare variables to be “shared”
between the application and DBMS

Embedded SQL v.s. API

- Pros of embedded SQL:
 - Be processed by a preprocessor prior to compilation → may catch SQL-related errors at preprocessing time
 - API: SQL statements are interpreted at runtime
- Cons of embedded SQL:
 - New host language code → complicate debugging

A brief look at other approaches

- “Embed” SQL in a general-purpose programming language
 - E.g.: embedded SQL
- Support database features through an object-oriented programming language
 - E.g., object-relational mappers (ORM) like Python SQLAlchemy
- Extend a general-purpose programming language with SQL-like constructs
 - E.g.: LINQ (Language Integrated Query for .NET)

Object-relational mapping

- Example: Python SQLAlchemy

```
class User(Base):
    __tablename__ = 'users'
    id = Column(Integer, primary_key=True)
    name = Column(String)
    password = Column(String)
```

```
class Address(Base):
    __tablename__ = 'addresses'
    id = Column(Integer, primary_key=True)
    email_address = Column(String, nullable=False)
    user_id = Column(Integer, ForeignKey('users.id'))
```

```
Address.user = relationship("User", back_populates="addresses")
User.addresses = relationship("Address", order_by=Address.id, back_populates="user")
```

```
jack = User(name='jack', password='gjffdd')
jack.addresses = [Address(email_address='jack@google.com'),
                  Address(email_address='j25@yahoo.com')]
session.add(jack)
session.commit()
```

```
session.query(User).join(Address).filter(Address.email_address=='jack@google.com').all()
```

- Automatic data mapping and query translation
- But syntax may vary for different host languages
- Very convenient for simple structures/queries, but quickly get complicated and less intuitive for more complex situations

A brief look at other approaches

- “Embed” SQL in a general-purpose programming language
 - E.g.: embedded SQL
- Support database features through an object-oriented programming language
 - By automatically storing objects in tables and translating methods to SQL
 - E.g., object-relational mappers (ORM) like Python SQLAlchemy
- Extend a general-purpose programming language with SQL-like constructs
 - E.g.: LINQ (Language Integrated Query for .NET)

Deeper language integration

- Example: LINQ (Language Integrated Query) for Microsoft .NET languages (e.g., C#)

```
int someValue = 5;
var results = from c in someCollection
              let x = someValue * 2
              where c.SomeProperty < x
              select new {c.SomeProperty, c.OtherProperty};
foreach (var result in results) {
    Console.WriteLine(result);
}
```

- Again, automatic data mapping and query translation
- Much cleaner syntax, but it still may vary for different host languages

Summary

- Basic SQL (queries, modifications, and constraints)
- Intermediate SQL (triggers, views, indexes)
- Recursion
- Programming
 - Augment SQL, e.g., SQL/PSM
 - Through an API, e.g., Python psycopg2, JDBC
 - Embedded SQL, e.g., in C
 - Automatic object-relational mapping, e.g.: Python SQLAlchemy
 - Extending programming languages with SQL-like constructs, e.g.: LINQ