


Relational Database Design: E/R-Relational Translation

Introduction to Database Management

CS348 Spring 2021

E/R Model


- E/R Concepts
- E/R Schema Design
- Next: Translating E/R to relational schema



- *Building(name, year)*
- *Room(building_name, room_number, capacity)*
- *Seat(building_name, room_number, seat_number, left_or_right)*

Translating entity sets

- An entity set translates directly to a table
 - Attributes → columns
 - Key attributes → key columns




User (uid, name)

Group (gid, name)


Translating weak entity sets

- Remember the “borrowed” key attributes
- Watch out for attribute name conflicts




Translating relationship sets

- A relationship set translates to a table
 - Keys of connected entity sets → columns
 - Attributes of the relationship set (if any) → columns
 - Multiplicity of the relationship set determines the key of the table



Member (uid, gid, fromDate)


More examples



Parent (parent_uid, child_uid)


Translating double diamonds?

- No need to translate because the relationship is implicit in the weak entity set's translation



Translating subclasses & ISA: approach 1

- Entity-in-all-superclasses approach (“E/R style”)
 - An entity is represented in the table for each subclass to which it belongs
 - A table includes only the attributes directly attached to the corresponding entity set, plus the inherited key




$\langle 142, \text{Bart} \rangle \in \text{User}(\underline{\text{uid}}, \underline{\text{name}})$
 $\langle 456, \text{Ralph} \rangle \in \text{User}(\underline{\text{uid}}, \underline{\text{name}})$
 $\langle 456, \odot \rangle \in \text{PaidUser}(\underline{\text{uid}}, \underline{\text{avatar}})$

Group (gid, name)
Member (uid, gid, from_date)


Translating subclasses & ISA: approach 2

- Entity-in-most-specific-class approach (“OO style”)
 - An entity is only represented in one table (the most specific entity set to which the entity belongs)
 - A table includes the attributes attached to the corresponding entity set, plus all inherited attributes



Translating subclasses & ISA: approach 3


- All-entities-in-one-table approach (“NULL style”)
 - One relation for the root entity set, with all attributes found in the network of subclasses
 - (plus a “type” attribute when needed)
 - Use a special NULL value in columns that are not relevant for a particular entity




Comparison of three approaches

- Entity-in-all-superclasses
 - User (uid, name), PaidUser (uid, avatar)
 - Pro: All users are found in one table
 - Con: Attributes of paid users are scattered in different tables
- Entity-in-most-specific-class
 - User (uid, name), PaidUser (uid, name, avatar)
 - Pro: All attributes of paid users are found in one table
 - Con: Users are scattered in different tables
- All-entities-in-one-table
 - User (uid, [type,]name, avatar)
 - Pro: Everything is in one table
 - Con: Lots of NULL's; complicated if class hierarchy is complex

A complete example



A complete example



Simplifications and refinements

*Train (number, engineer), LocalTrain (number), ExpressTrain (number)
Station (name, address), LocalStation (name), ExpressStation (name)
LocalTrainStop (local_train_number, station_name, time)
ExpressTrainStop (express_train_number, express_station_name, time)*

- Eliminate *LocalTrain* table
 - Redundant: can be computed as
$$\pi_{number}(Train) - ExpressTrain$$
 - Slightly harder to check that *local_train_number* is indeed a local train number
- Eliminate *LocalStation* table
 - It can be computed as $\pi_{number}(Station) - ExpressStation$

An alternative design

Train (number, engineer, type)

Station (name, address, type)

TrainStop (train_number, station_name, time)

- Encode the type of train/station as a column rather than creating subclasses
- What about the following constraints?
 - Type must be either “local” or “express”
 - Express trains only stop at express stations
 - ☞ They can be expressed/declared explicitly as database constraints in SQL
 - ☞ Arguably a better design because it is simpler!

Design principles

- KISS
 - Keep It Simple, Stupid
- Avoid redundancy
- Capture essential constraints, but don't introduce unnecessary restrictions
- Use your common sense
 - Warning: mechanical translation procedures given in this lecture are no substitute for your own judgment



More examples

- Representing aggregation
 - Tabular representation of aggregation of R = tabular representation for relationship set R
 - To represent relationship set involving aggregation of R , treat the aggregation like an entity set whose primary key is the primary key of the table for R


Student (StudentNum)

Couse(CourseNum)

Account(UserID)


EnrolledIn(StudentNum,CouseNum)

CouseAccount(UserId, StudentNum, CourseNum, ExpirationDate)



More examples (Exercise)

- ER Diagram





Relational Schema

?


More examples

- ER Diagram



More examples

- ER Diagram



Relational DDL Commands

```
CREATE TABLE Course
(CourseNum INTEGER PRIMARY KEY,
CourseName CHAR(50));
```

```
CREATE TABLE Student
(StudentNum INTEGER PRIMARY KEY,
StudentName CHAR(50),
GPA FLOAT);
```

```
CREATE TABLE Professor
(ProfNum INTEGER PRIMARY KEY,
ProfName CHAR(50));
```

```
CREATE TABLE Section
(CourseNum INTEGER NOT NULL REFERENCES Course(CourseNum),
SectionNum INTEGER NOT NULL,
Term INTEGER NOT NULL,
PRIMARY KEY(CourseNum, SectionNum, Term),
ProfNum INTEGER NOT NULL REFERENCES Professor(ProfNum));
```

```
CREATE TABLE Off-SiteSection
(CourseNum INTEGER NOT NULL,
SectionNum INTEGER NOT NULL,
Term INTEGER NOT NULL,
FOREIGN KEY(CouseNum,SectionNum,Term) REFERENCES
Section(CouseNum,SectionNum,Term),
Location CHAR(50));
```

```
CREATE TABLE EnrolledIn
(CourseNum INTEGER NOT NULL,
SectionNum INTEGER NOT NULL,
Term INTEGER NOT NULL,
StudentNum INTEGER NOT NULL REFERENCES Student(StudentNum),
FOREIGN KEY(CouseNum,SectionNum,Term) REFERENCES
Section(CouseNum,SectionNum,Term),
Primary Key(CouseNum,SectionNum,Term,StudentNum),
Mark INTEGER);
```

Database Design


- Entity-Relationship (E/R) model
- Translating E/R to relational schema
- Next week: Relational design principles

Clarifications

June 03, 2022

(slide 5) Translating relationship sets

- A relationship set translates to a table
 - Keys of connected entity sets → columns
 - Attributes of the relationship set (if any) → columns
 - Multiplicity of the relationship set determines the key of the table



- If we can deduce the general cardinality constraint (0,1) for a component entity set E, then take the primary key attributes for E
- Otherwise, choose primary key attributes of each component entity

(Slide 17) More examples

- Representing aggregation
 - Tabular representation of aggregation of R = tabular representation for relationship set R
 - To represent relationship set involving aggregation of R , treat the aggregation like an entity set whose primary key is the primary key of the table for R


Student (StudentNum)

Couse(CourseNum)

Account(UserID)

EnrolledIn(StudentNum,CouseNum)

CouseAccount(UserId,StudentNum,CourseNum,ExpirationDate)



One-to-one relationships → We can simply take UserId or (StudentNum, CourseNum) as the eky