

CS 348 Lecture 12

Physical Data Organization Finish & Indices

Semih Salihoglu

Feb 17th 2022



UNIVERSITY OF
WATERLOO | **DSg** Data
Systems
Group

Announcements

- Assignment 3 past but still have until tomorrow midnight:
- Reading week: No lectures or OHs next week.

Outline For Today

1. Recap & Finish Physical File Organization Designs
 - Row-oriented Physical Design (Recap)
 - Column-oriented Physical Design
 - Hybrid (PAX) Physical Design
 - Designs for Variable-length Fields
 - Designs for NULLs
2. Database Indices
 - 5 Index Designs in Increasing Level of Robustness
 - Using Indices In Practice

Outline For Today

1. Recap & Finish Physical File Organization Designs
 - Row-oriented Physical Design (Recap)
 - Column-oriented Physical Design
 - Hybrid (PAX) Physical Design
 - Designs for Variable-length Fields
 - Designs for NULLs
2. Database Indices
 - 5 Index Designs in Increasing Level of Robustness
 - Using Indices In Practice

Recall Potential Goals of Physical Data Design

1. Minimize I/O!
 - i. Design files to pack records/columns into few number of pages
 - ii. Keep related data close in storage devices!
 - iii. Cache pages of data that is likely to be used
2. Maximize sequential reads from storage devices!
3. Simplicity/speed of accessing values
 - E.g., don't compress using a complex compression scheme

These are general guidelines and not necessarily independent.

Common Record and Page Layouts

1. Row-oriented
2. Column-oriented
3. Hybrid (PAX) Row & Column-oriented Fixed-length Values
4. Designs for Variable-length Values
5. Designs for NULL Values

Recap: Row-Oriented NSM Design

- Each page contains ~ H/T many rows and all fields of each row

Customer					
cid : int64 (8 bytes)		name : string (60 chars)		isGoldMember : bool (1 byte)	
101	Alice Munro\0\0\0....\0	1	201	Carl Sagan\0\0\0...\0	0
103	Bob \0\0\0....\0	1
<i>(Example Disk Page, e.g., in file customer.mkd)</i>					

- Note: because records are fixed length, can read any field with simple arithmetic (*not a consequence of NSM*)
 - E.g: 5th records name field at offsets: $4*(8+60+1) + 8 = 284$ to 344

Recap: Pros/Cons of NSM

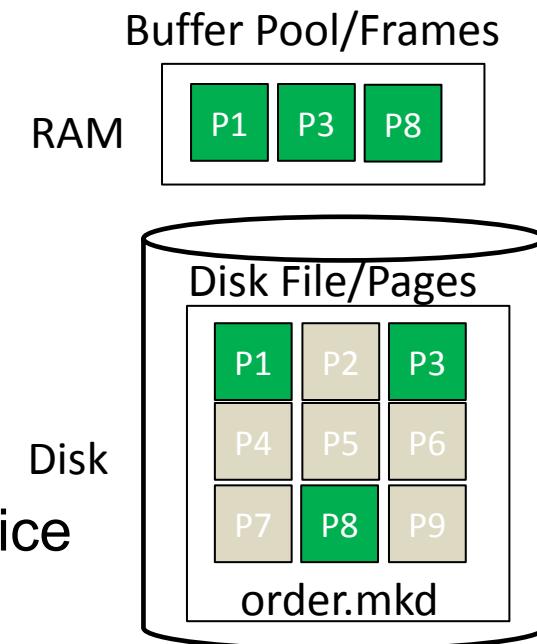
- **Pro:** Good for queries that access multiple or all fields of tuples
- **Con:** If queries access a single field, e.g., counting numGoldMembers, “effective read size” can be small
 - E.g. 1/69 bytes read is useful, in columnar storage you can do 69x less I/O to do the same count.
- **Con:** For some workloads higher # I/O => lower (buffer) cache utilization (next slide)

Recap: Buffer Manager & Ex Op/Storage Interface

- DBMSs have a *buffer manager* (BM) component that keeps a set of file pages in memory to reduce I/O.
- Goal of BM: keep pages that are likely to be accessed in memory
 - Implement a caching alg, e.g., Least-Recently Used (LRU)

```
class DatabaseFileHandle : {  
    string filename; // e.g., order.mkd  
    int tupleLen; // e.g., 69  
    BufferManager bm;  
    void readRecord(Tuple* outTuple, recordID) {  
        (pageIdx, offset)= calculatePageInfo(recordID)  
        Frame* frame = bm.pinPage(pageIdx);  
        memcpy(outTuple, frame->bytes[offset], tupleLen)  
        bm.unpin(frame); } }
```

- Pin: put disk page into buffer pool until further notice
 - possibly evicts an existing page
- Unpin: can now remove the page from buffer pool if needed



Column-Oriented Storage Design

- One file for each field of table (~H/field-size rows worth of data):

101	201	103	330
...
<i>cID.mkd</i>			

Alice Munro\0\0....\0
Carl Sagan\0\0\0....\0
Bob \0\0\0....\0
<i>name.mkd</i>

1	0	1	0	0	0
.....
<i>isGoldMember.mkd</i>					

- **Pros:** Good for queries that access few columns
 - Can reduce I/O significantly
 - Better sequential reads when pages are in memory
 - **Pros:** Easier to apply compression because each page is very “homogenous”
 - **Cons:** Bad for queries that access all columns (1 I/O in row-oriented storage could be done with # cols I/O).
- Select count(*)
from customer
where isGoldMember

PAX Hybrid Storage Design

- ## □ PAX: Partition Attributes Across (Ailamaki et al. VLDB 2001)

- All fields of a row in the same page
 - But pages internally organized by columns



101	201	103	330
Alice Munro				Carl	
1	0	1	0	0	0

- Better sequential reads for queries that read a single column
 - Gives the advantage of having all fields of a row in the same page
 - But still more I/O than pure columnar storage for some queries

Outline For Today

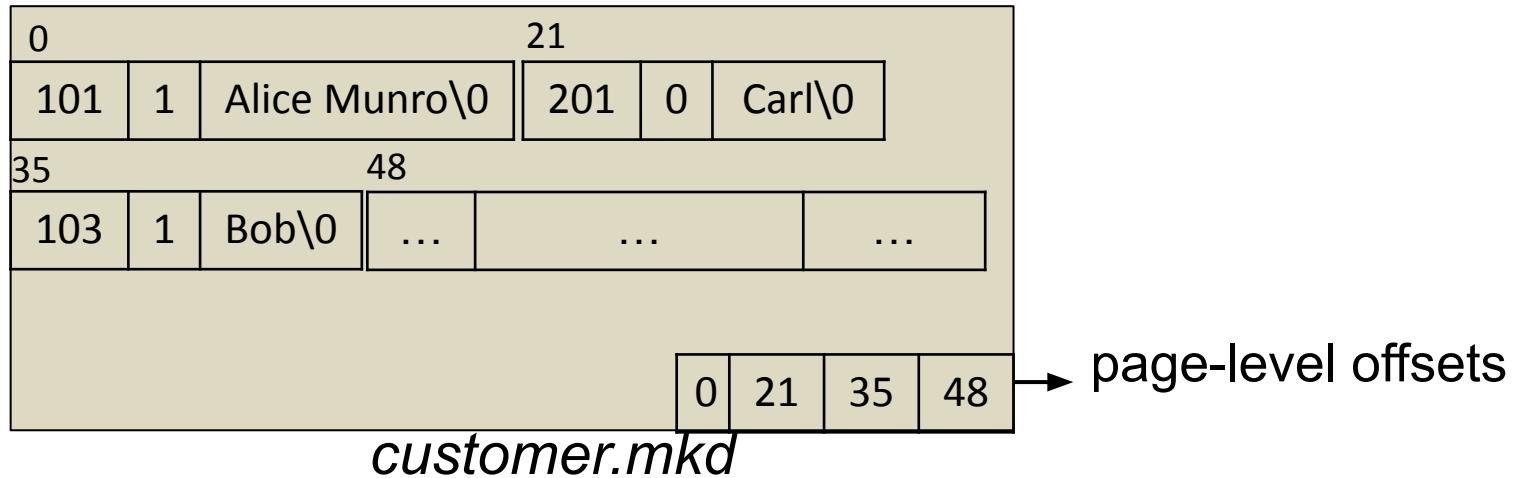
1. High-level DBMS Architecture
2. Storage Manager and Physical File Organization Designs
 - Physical Operators/Storage Interface
 - Fundamental Property of Storage Devices & Potential Goals of Physical Design
 - Row-oriented Physical Design
 - Column-oriented Physical Design
 - Hybrid (PAX) Physical Design
 - **Designs for Variable-length Fields**
 - **Designs for NULLs**

Variable-length Values

- Pages are often more complex than examples in these slides
 - One complexity: variable-length fields
 - Assume row-oriented NSM storage
 - Approaches to encode var-length fields, e.g., var-len strings, in pages
 - 1. Delimiters
 - 2. Field Offsets
 - 3. Pointers to Overflow Pages
- Both:
- (i) put first fixed-len fields and then var-len fields when encoding each row
 - (ii) require storing page-level offset of each row

Delimiter Approach

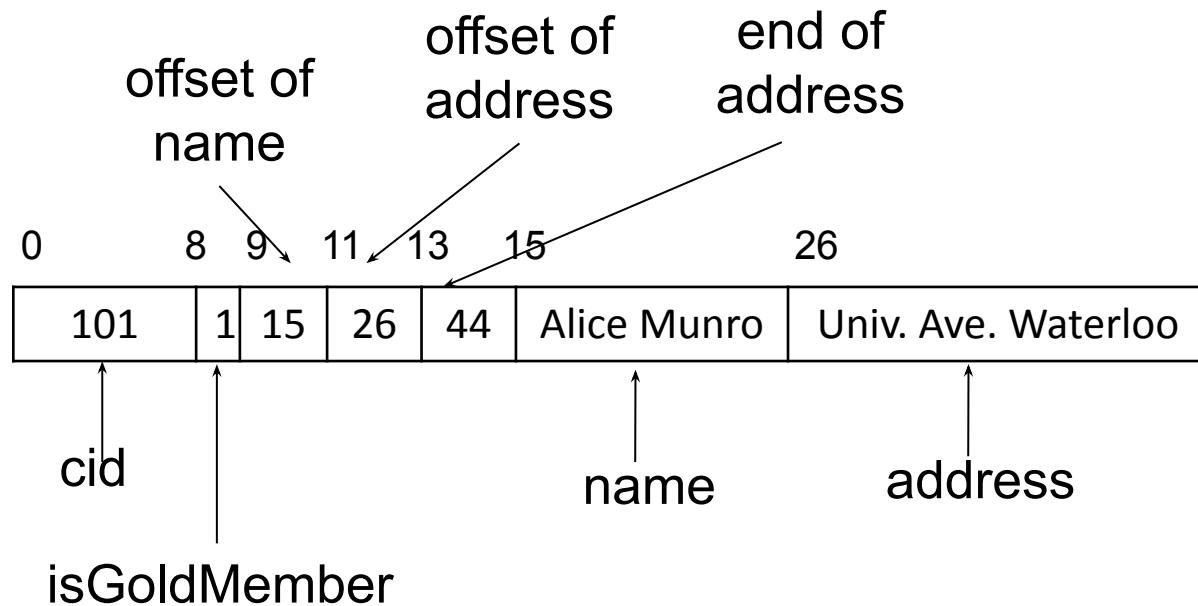
- Example Delimiter: \0 character



- Why are page offsets at the end and fields in the beginning of pages?
 - Can grow them separately without sliding data within pages.
- Advantage of storing fixed-len fields and then var-len fields:
 - Can read fixed-len fields w/ page-level offsets + arithmetic
 - e.g., isGoldMember of rowID 2: $\text{offset}[2] + 8 = 35+8=43$.
 - no need to scan over var-len fields

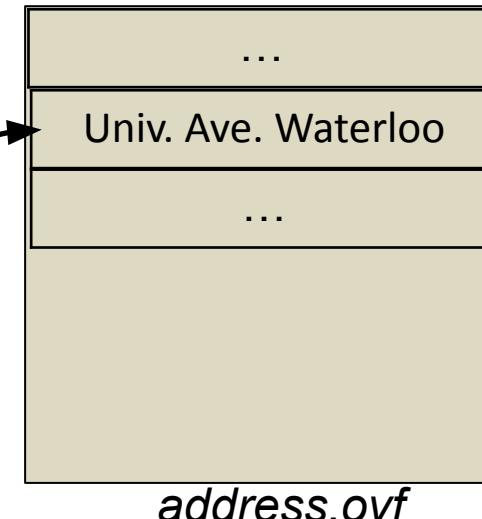
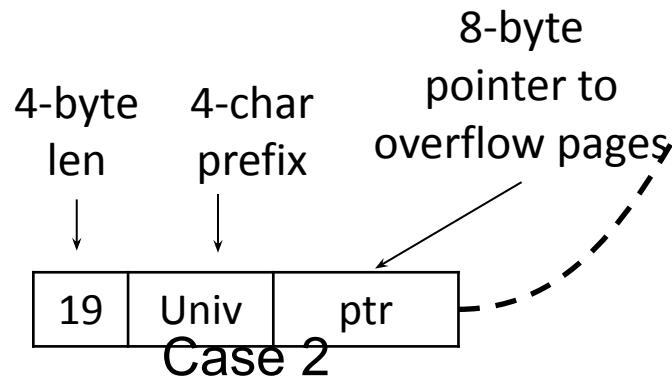
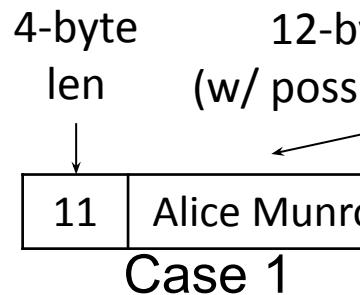
Field Offsets Approach

- Suppose there is a second var-len field called address
- Page structure stays same except each row would be encoded as:



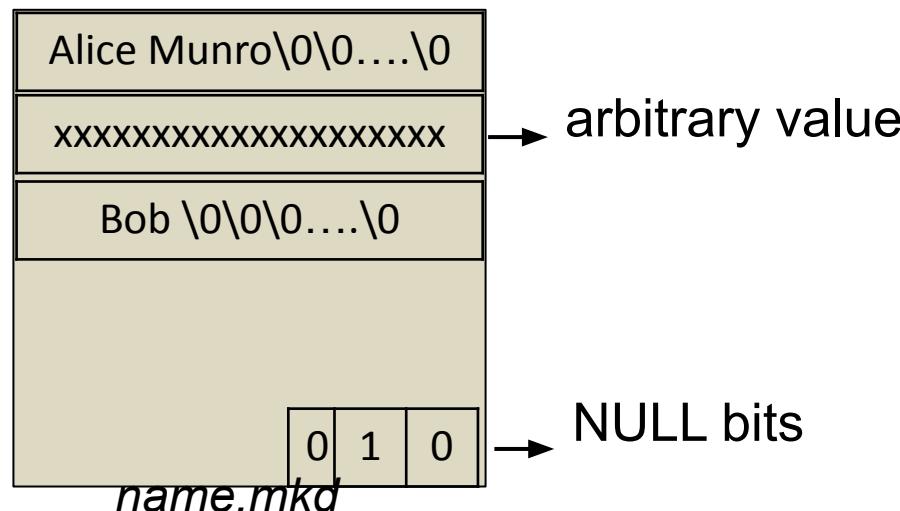
Overflow Pages

- Keep variable length fields, in particular strings, fixed length
- Any overflow points to separate overflow pages
- E.g., Strings in Graphflow (originally by Thomas Neumann)
 - 16 bytes fixed part with 2 possible cases
 - **Pro:** original pages are simpler; **Con:** More I/O for reading large strings



NULL Values

- Several possible approaches. E.g:
 - Store special NULL values for each data type as regular values
 - E.g.: -2^{63} for 8-byte int
 - Store NULL bits in the page
 - Assume columnar storage and (fixed-len) name can be NULL
 - 1 means NULL 0 means NOT NULL
 - Other alternatives may be possible



Summary of Physical Data Design Alternatives

- High-level: row-oriented vs columnar but row-oriented can be hybrid (PAX)
- Variable-length fields and fields that can be NULL require more sophisticated design
- For any physical design we showed, other alternatives are possible
- In practice physical design is principled art!
- Many tradeoff, choices, and complications exist:
 - Deletions of tuples (can leave gaps in pages or shift tuples)
 - Fields larger than page sizes
 - Tuple versions (some systems, e.g., Postgres, support multiple version of tuples so users can keep track of how each record changed)
 - Compression
 - Some crazy ideas put rows from multiple tables into same pages
 - Many other designs

Outline For Today

1. Recap & Finish Physical File Organization Designs
 - Row-oriented Physical Design (Recap)
 - Column-oriented Physical Design
 - Hybrid (PAX) Physical Design
 - Designs for Variable-length Fields
 - Designs for NULLs
2. Database Indices
 - 5 Index Designs in Increasing Level of Robustness
 - Using Indices In Practice

Functionality of Indices (1)

- Indices are the primary mechanism to

- retrieve records quickly*
- search records in sort order*

```
SELECT * FROM Students WHERE ID = 912;
```

```
SELECT * FROM Students WHERE ID > 100;
```

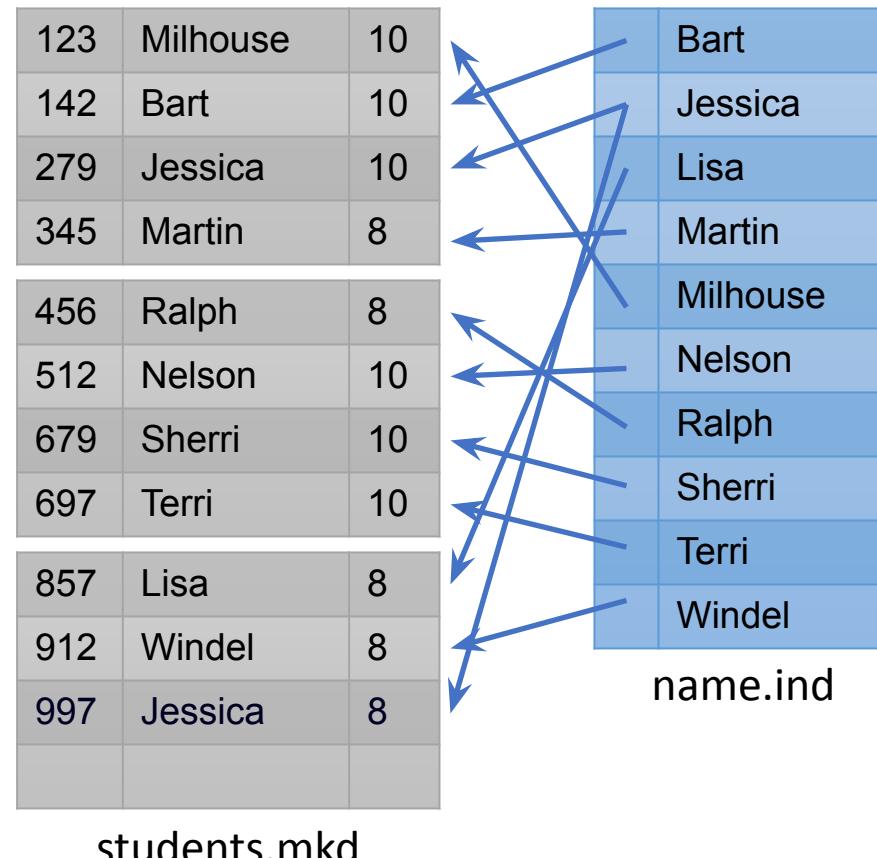
- Default way to find records:: sequential scans

- Read each page and each record
- Can be very slow for large tables
- If a file sorted on some columns:
 - Can now do binary search
 - Key Question: How to do this efficiently?

ID	name	mark
123	Milhouse	10
142	Bart	10
279	Jessica	10
345	Martin	8
456	Ralph	8
512	Nelson	10
679	Sherri	10
697	Terri	10
857	Lisa	8
912	Windel	8
997	Jessica	8
...
...
...

Functionality of Indices (2)

- Indices: are persistent data structures that are stored along with table files that allow fast search.
- An example of a simple index: can be much smaller than the original table.



Naïve Approach for Keeping a Table Sorted

- Sorting: primary technique to find things & data quickly!
- Once a file is sorted, we can do binary search on the pages of the file.
- How to keep a relation file in sorted order (e.g., students.mkd)?
- Assume a sequence of insertions. 2 records/page. Sort on ID
- Simple/naïve approach: Shift-based (index-less) sorting

Naïve Approach for Keeping a Table Sorted

Next Insertion

857	Lisa	8
-----	------	---

PAGES

Naïve Approach for Keeping a Table Sorted

Next Insertion

512	Nelson	10
-----	--------	----

PAGES

857	Lisa	8

pg 1

Naïve Approach for Keeping a Table Sorted

Next Insertion

279	Jessica	10
-----	---------	----

PAGES

512	Nelson	10
857	Lisa	8

pg 1

Naïve Approach for Keeping a Table Sorted

Next Insertion

912	Windel	8
-----	--------	---

PAGES

279	Jessica	10
512	Nelson	10

pg 1

857	Lisa	8

pg 2

Naïve Approach for Keeping a Table Sorted

Next Insertion

345	Martin	8
-----	--------	---

PAGES

279	Jessica	10
512	Nelson	10

pg 1

857	Lisa	8
912	Windel	8

pg 2

Naïve Approach for Keeping a Table Sorted

Next Insertion

697	Terri	10
-----	-------	----

PAGES

279	Jessica	10
345	Martin	8

pg 1

512	Nelson	10
857	Lisa	8

pg 2

912	Windel	8

pg 3

Naïve Approach for Keeping a Table Sorted

Next Insertion

123	Milhouse	10
-----	----------	----

PAGES

279	Jessica	10
345	Martin	8

pg 1

512	Nelson	10
697	Terri	10

pg 2

857	Lisa	8
912	Windel	8

pg 3

Naïve Approach for Keeping a Table Sorted

Next Insertion

PAGES

123	Milhouse	10
279	Jessica	10

pg 1

345	Martin	8
512	Nelson	10

pg 2

697	Terri	10
857	Lisa	8

pg 3

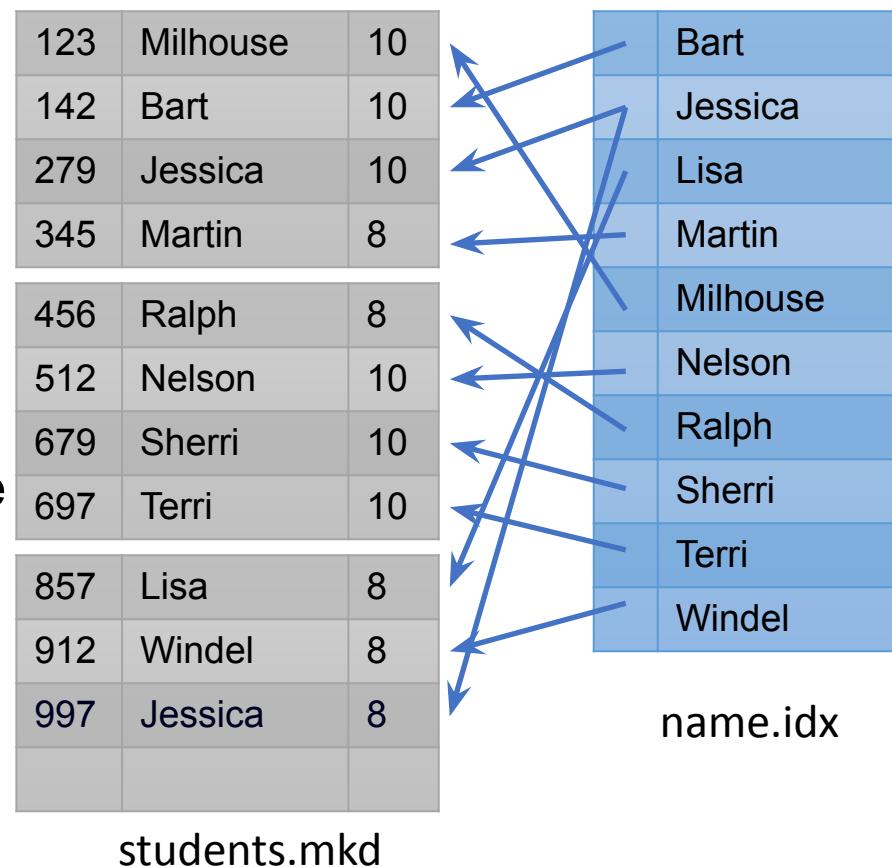
912	Windel	8

pg 4

- **Pro:** Very simple to implement
- **Con:** Each insertion could require up to $2xb$ many I/Os (to read and right pages) if the table has b pages.
 - Will not scale. Not practical.

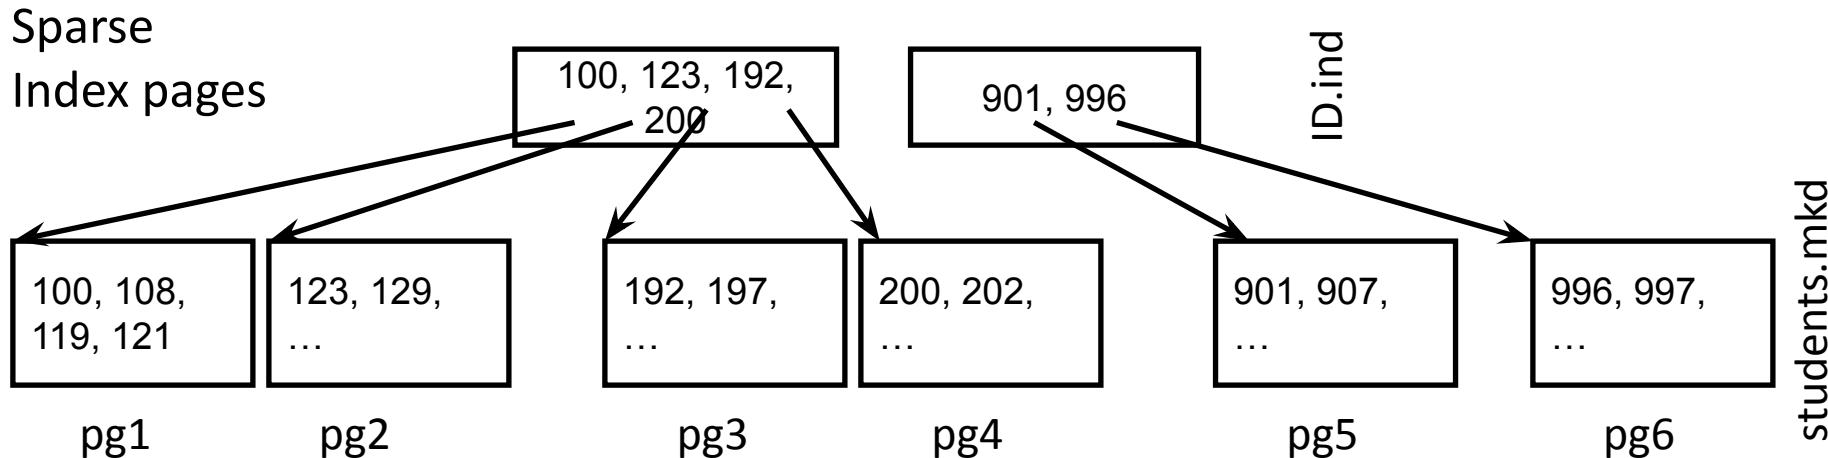
2nd Approach: Single-level Dense Index

- Lookup: find the record in index & follow pointer (page & offset)
- If index file is disk-based:
 - **Con:** Same problem as naïve solution
 - **Pro:** But at a smaller scale b/c the index is smaller (a projection).
- If index is in memory:
 - **Pro:** Optimal I/O. Only store the record in the relation file but no sorting (called unclustered index)
 - **Con:** Index cannot get very large.



3rd Approach: Single-level Sparse Index w/ Overflows

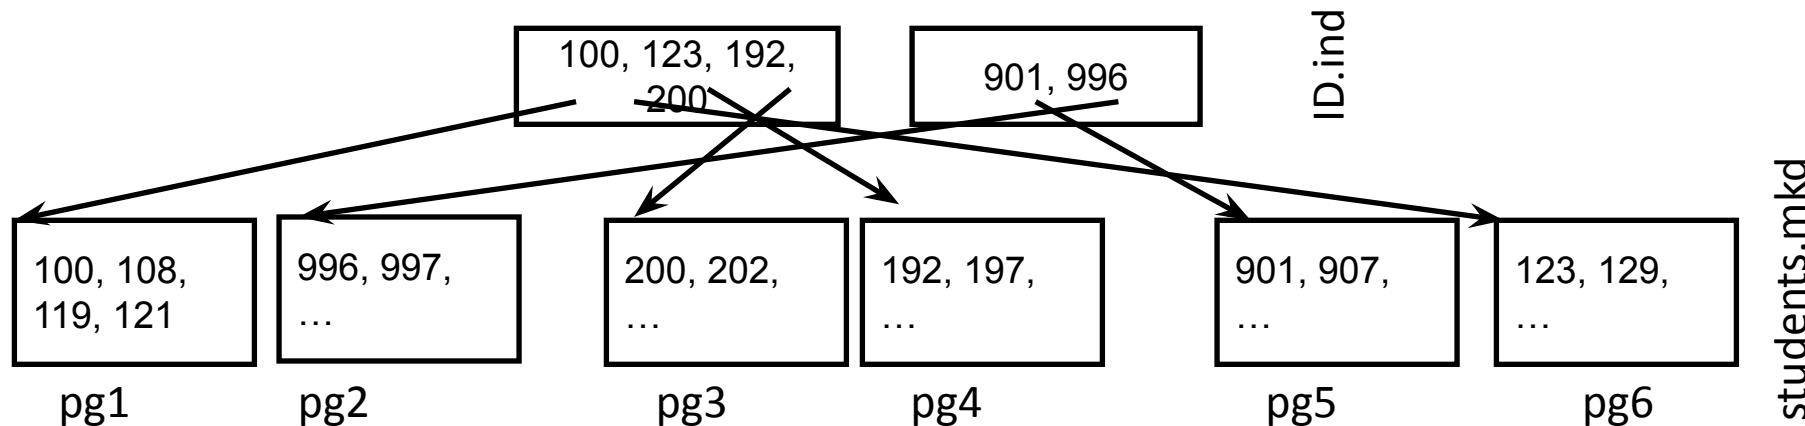
- Suppose the sort column keys have a relatively stable domain and the table is not expected to grow significantly.
- Can do an initial sort upon data ingestion and keep a sparse-index
- Below: Just showing the sort column not entire rows
- Lookup: Find page, follow pointer, scan page (& overflow pages (soon))



- Need the data pages sorted (called clustered index)
- Advantage over dense index: much smaller (can be a few orders of magn.)
- Insertions require chaining & deletions can lead to empty pages (soon)

Note on Clustered Indices

- When a relation file has a clustered index, i.e., when pages are sorted, we do not necessarily need the pages to be sorted.
- If a file has b pages, as long as we can recover the full sort order with b I/Os (ignoring I/Os to read the sparse index), we get similar benefits
- Below design will have similar performance as previous slide if we read pages sequentially through sparse index



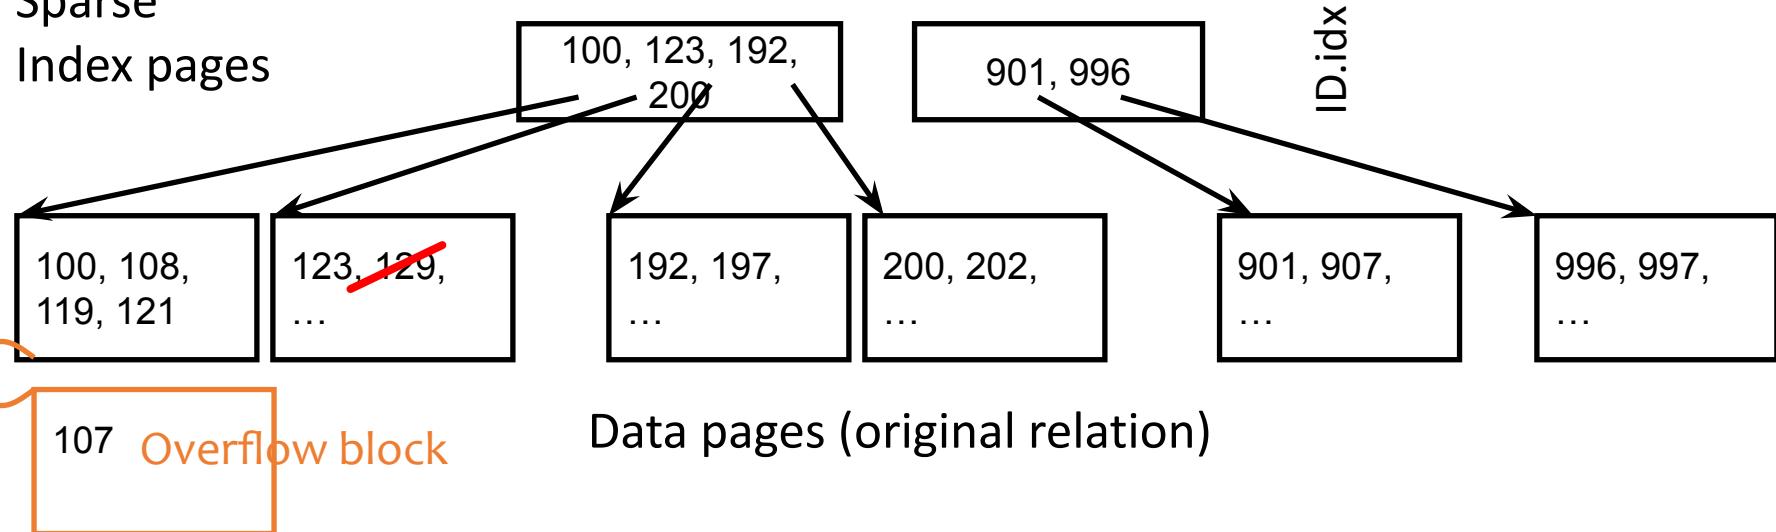
3rd Approach: Single-level Sparse Index w/ Overflows

□ Handing Insertions

Example: insert tuple with key 107

Example: delete tuple with key 129

Sparse
Index pages



`students.mkd`

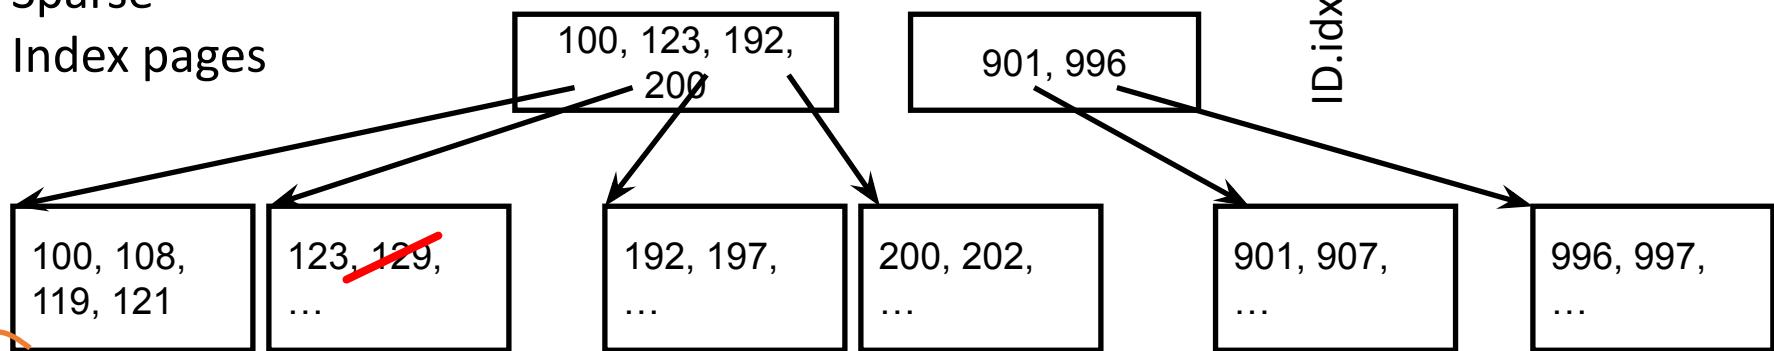
□ Overflow chains and empty data blocks degrade performance

- If there is significant *data distribution skew*: records can go into one long chain, so lookups require scanning all data in worst-case.

3rd Approach: Single-level Sparse Index w/ Overflows

- Pros: Index size smaller than dense index (1 key/ptr in index per page), so can be larger Address w/ multi-level indices Address w/ splitting/merging
- Cons:
 - Can still become very large (GBs) for large tables.
 - Need overflows, which is not robust, if table grows significantly over time (e.g., most pages can become overflows, leading to large scans)
 - Can lead to empty pages (but less of an issue in practice)

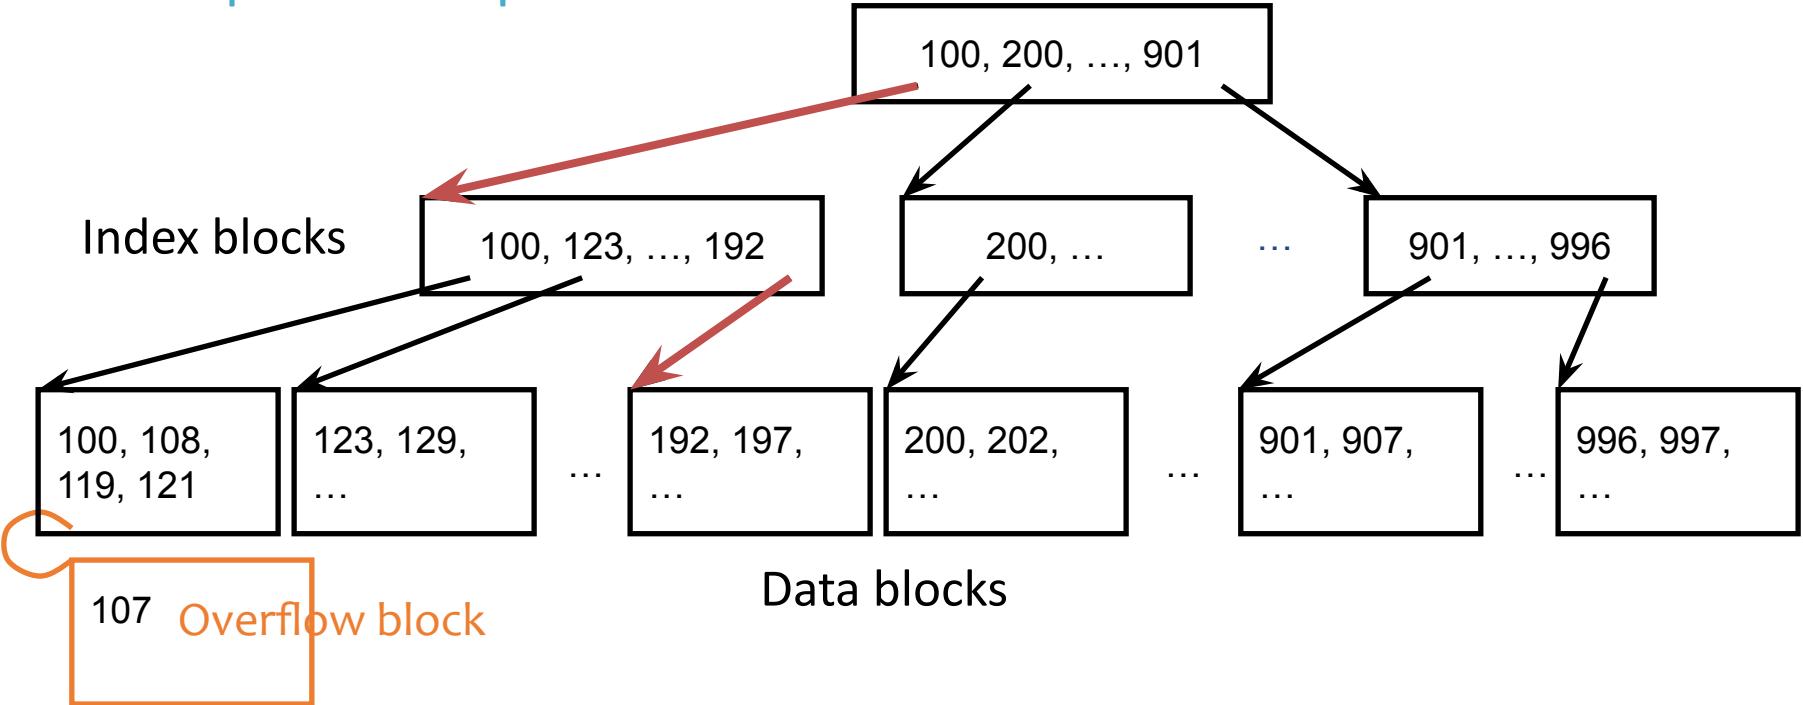
Sparse
Index pages



4th Approach: Multi-level Indices

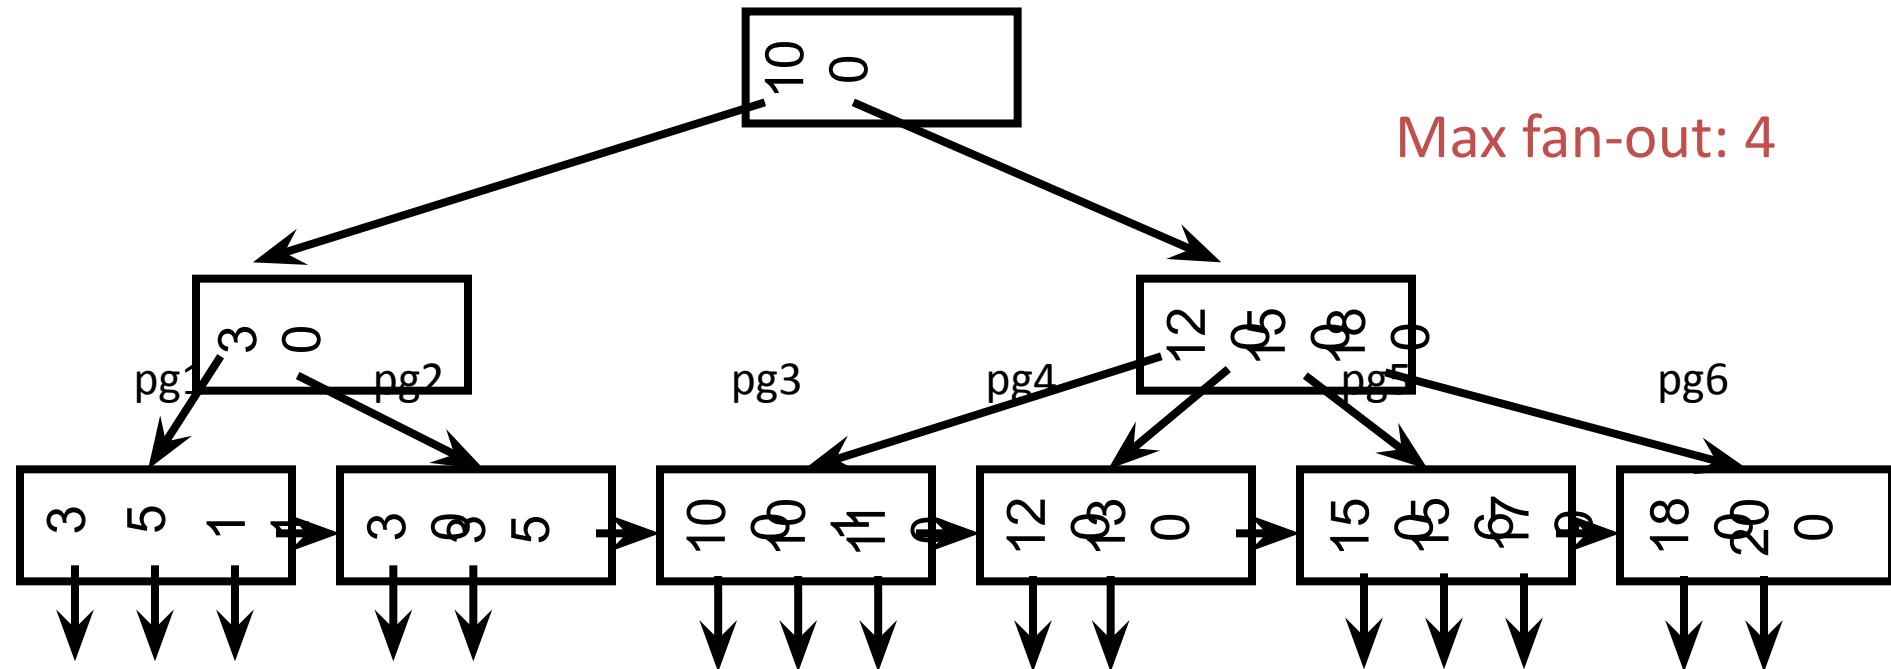
- If an index is too large, can put other layers of sparse indices on the index
- Forms a tree and allows the system to keep higher-level indices in memory
 - Can do depth-1 many I/Os in lookups (ignoring overflows)

Example: look up 197

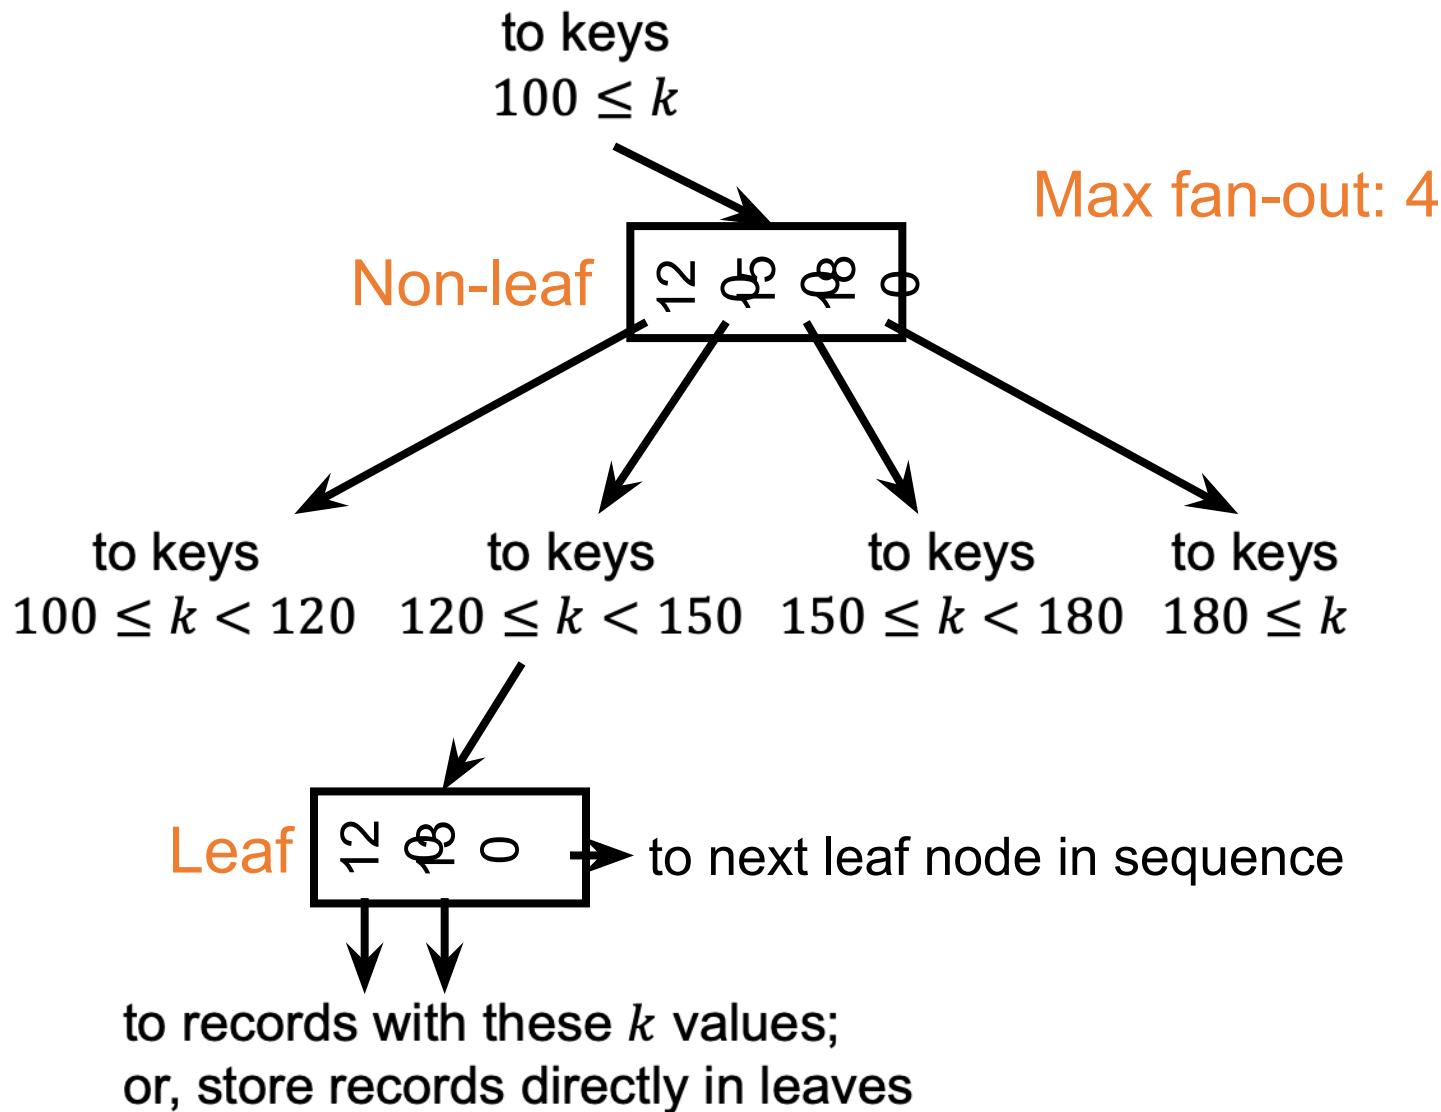


5th Approach: B/B+ Tree Indices

- Multi-level sparse indices on a first level of pages that is either:
 - actual relation pages (if clustered)
 - dense index on the relation pages (works for clustered or unclustered)
 - leaf level consists of *chained pages*
- Forms a k-ary balanced tree
- Instead of overflow pages uses splitting and merging of pages at any layer



Sample B⁺-tree Nodes



B⁺-tree Balancing Properties

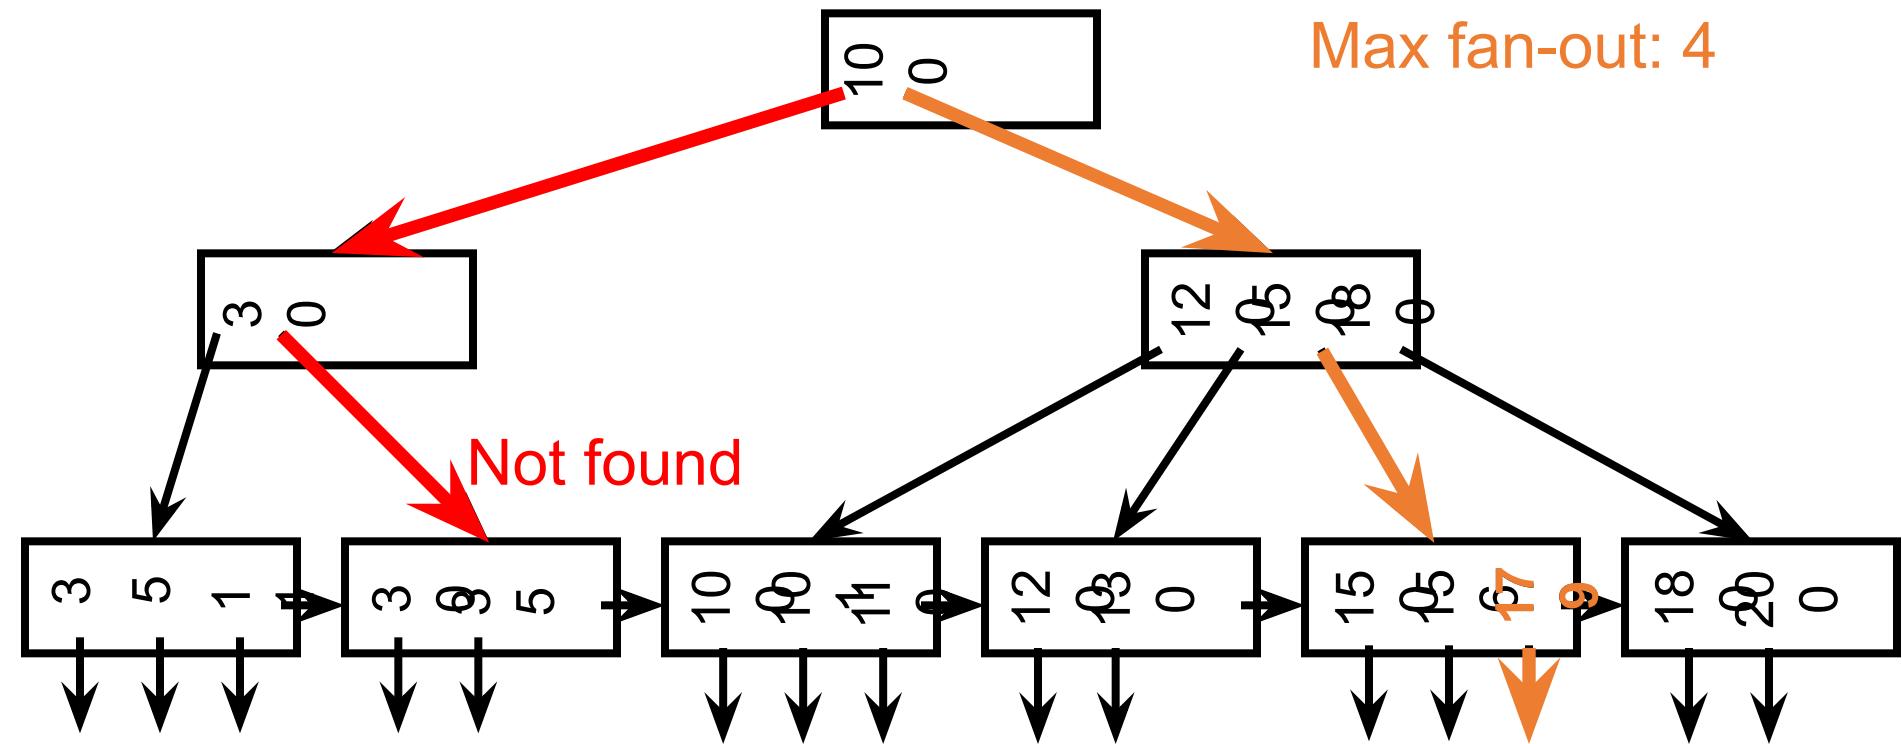
- Height constraint: all leaves at the same lowest level
- Fan-out constraint: all nodes at least half full (except root)

	Max # pointers	Max # keys	Min # active pointers	Min # keys
Non-leaf	f	$f - 1$	$\lceil f/2 \rceil$	$\lceil f/2 \rceil - 1$
Root	f	$f - 1$	2	1
Leaf	f	$f - 1$	$\lfloor f/2 \rfloor$	$\lfloor f/2 \rfloor$

Lookups

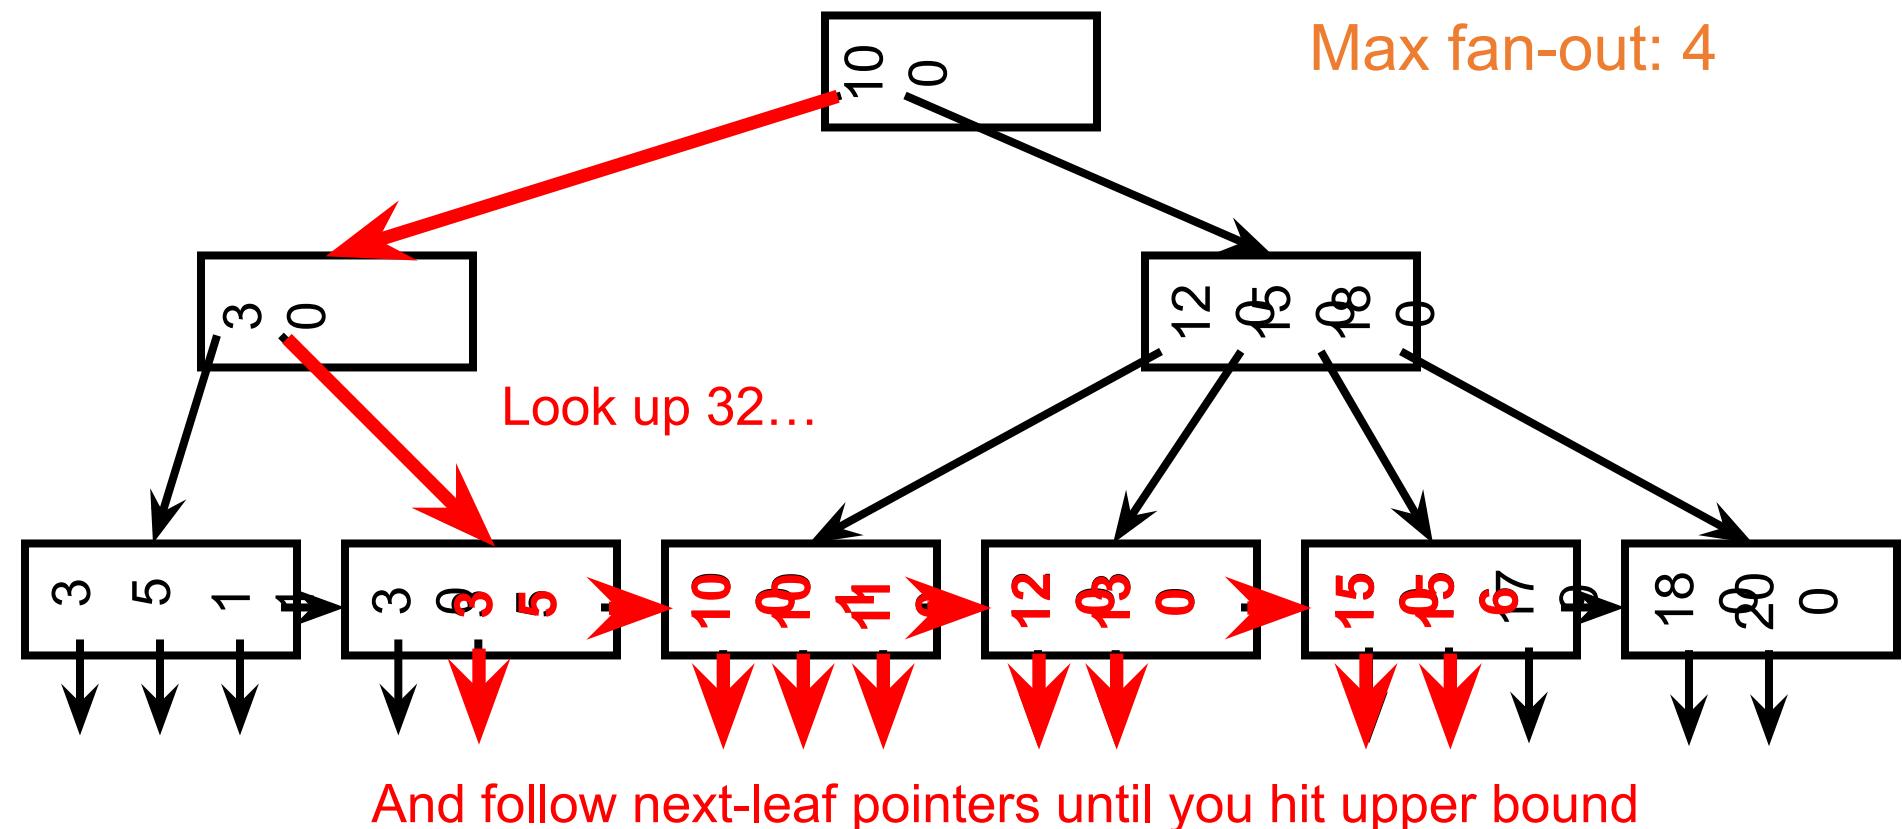
```
SELECT * FROM R WHERE k = 179;
```

```
SELECT * FROM R WHERE k = 32;
```



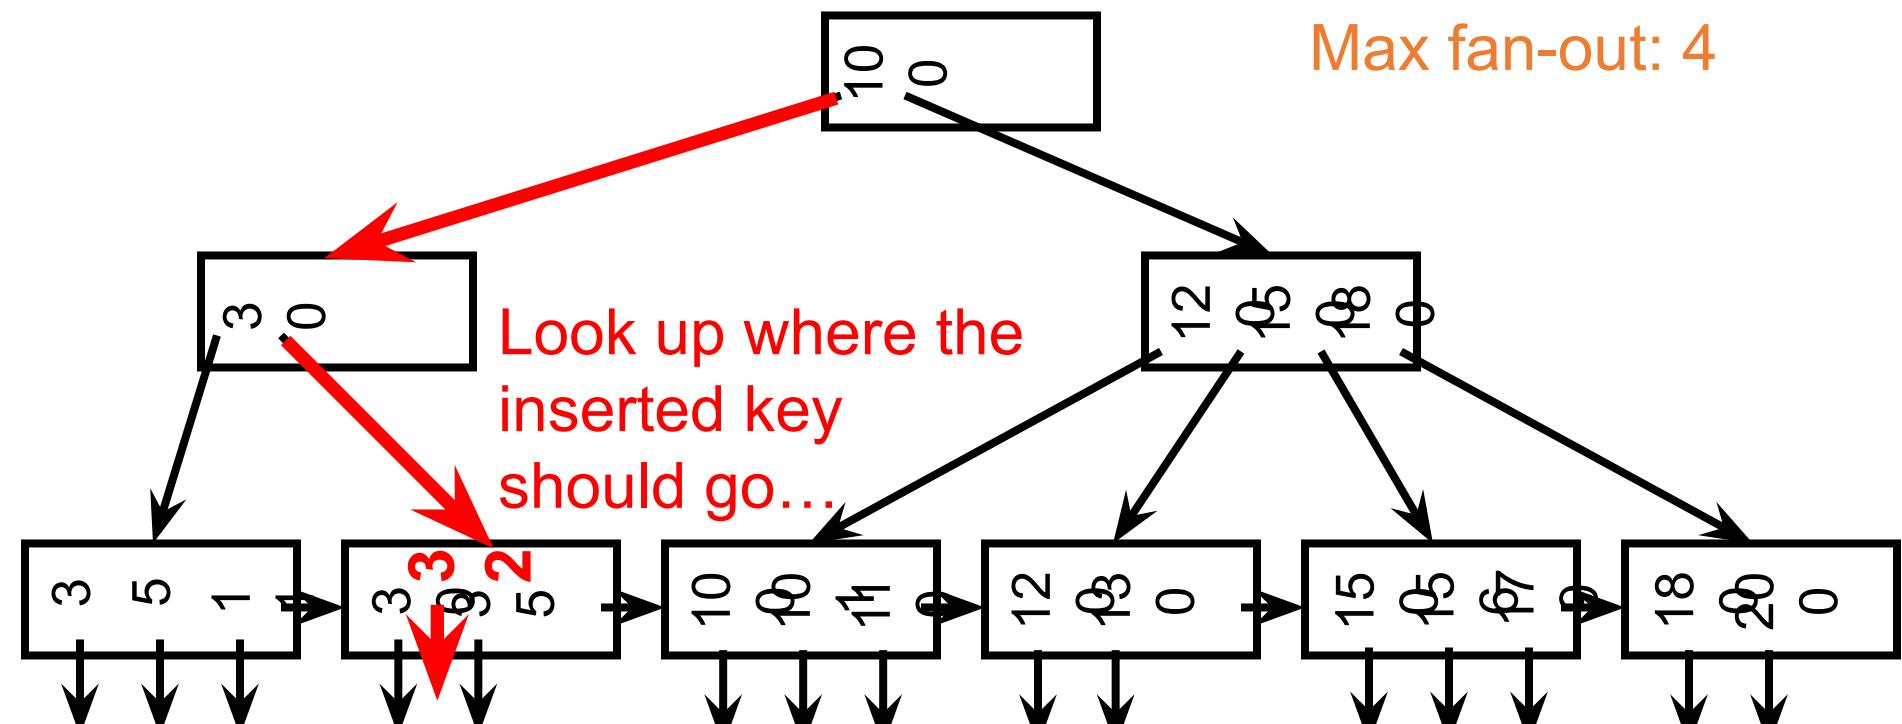
Range Query

- `SELECT * FROM R WHERE $k > 32$ AND $k < 179$;`



Insertion

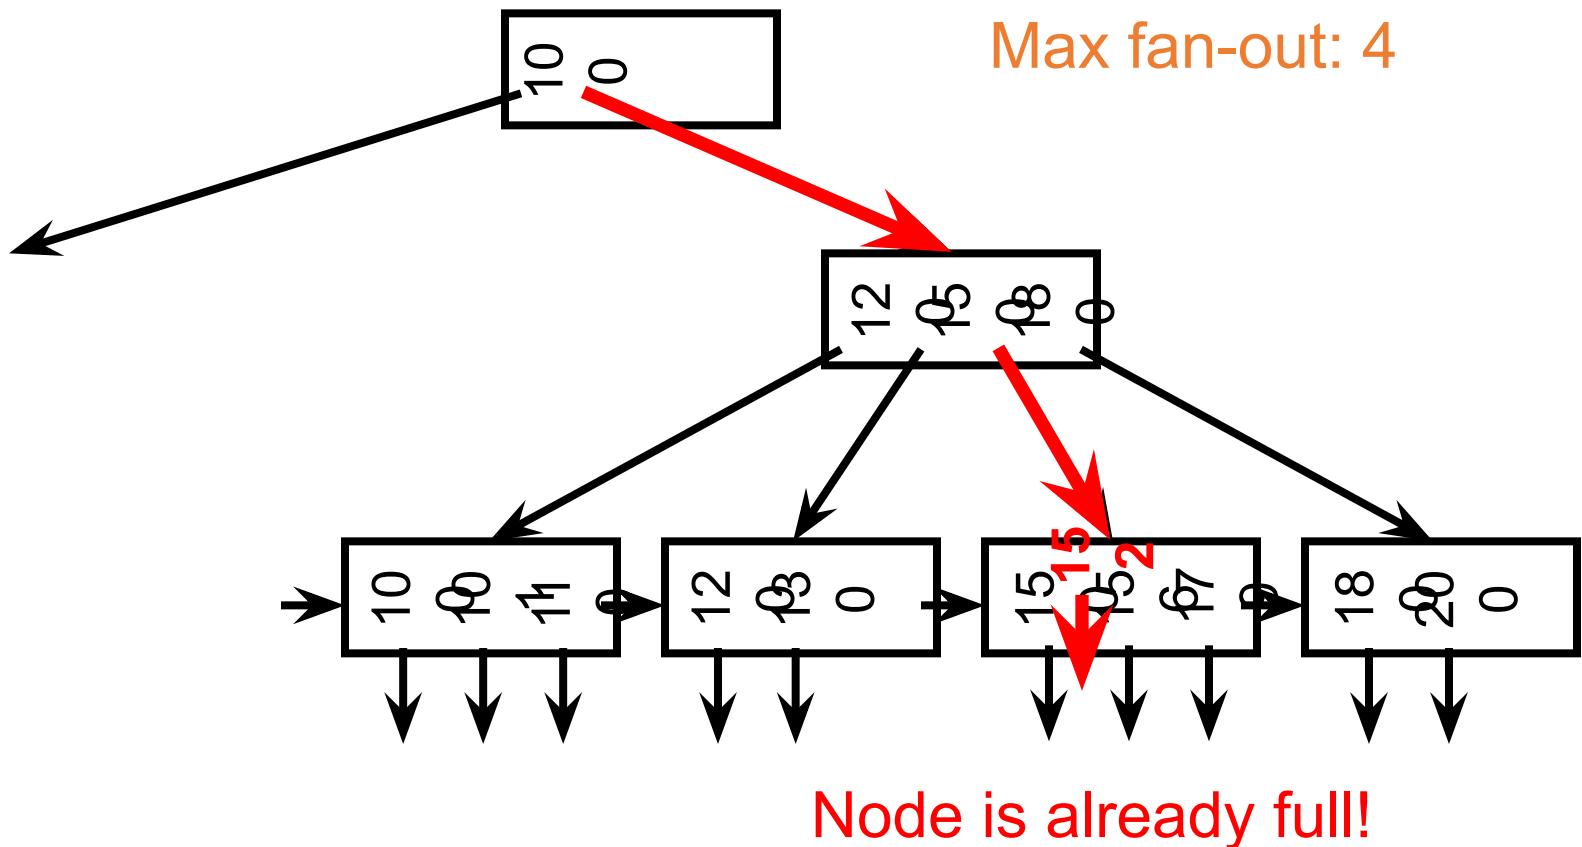
- Insert a record with search key value 32



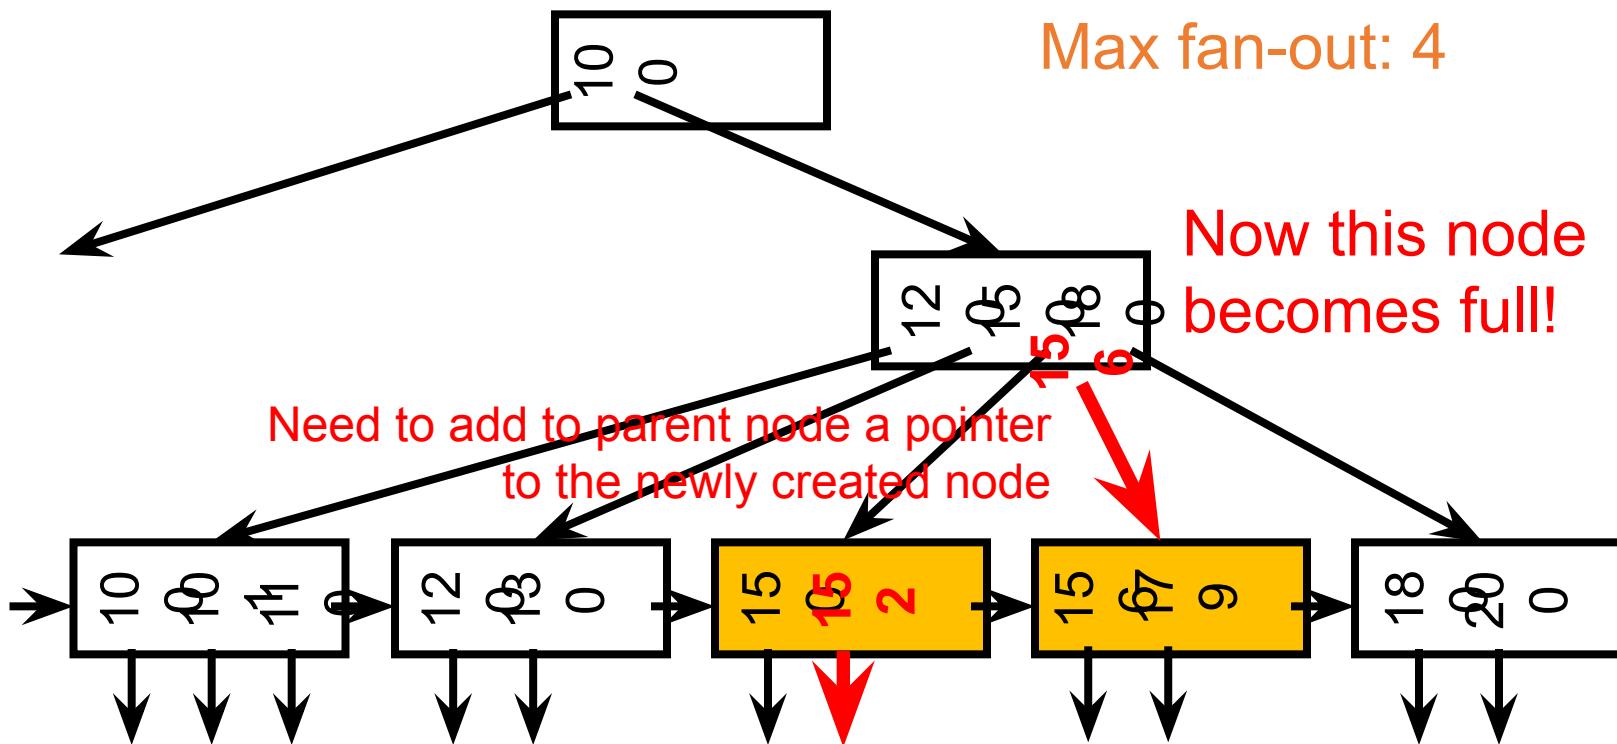
And insert it right there

Another Insertion Example

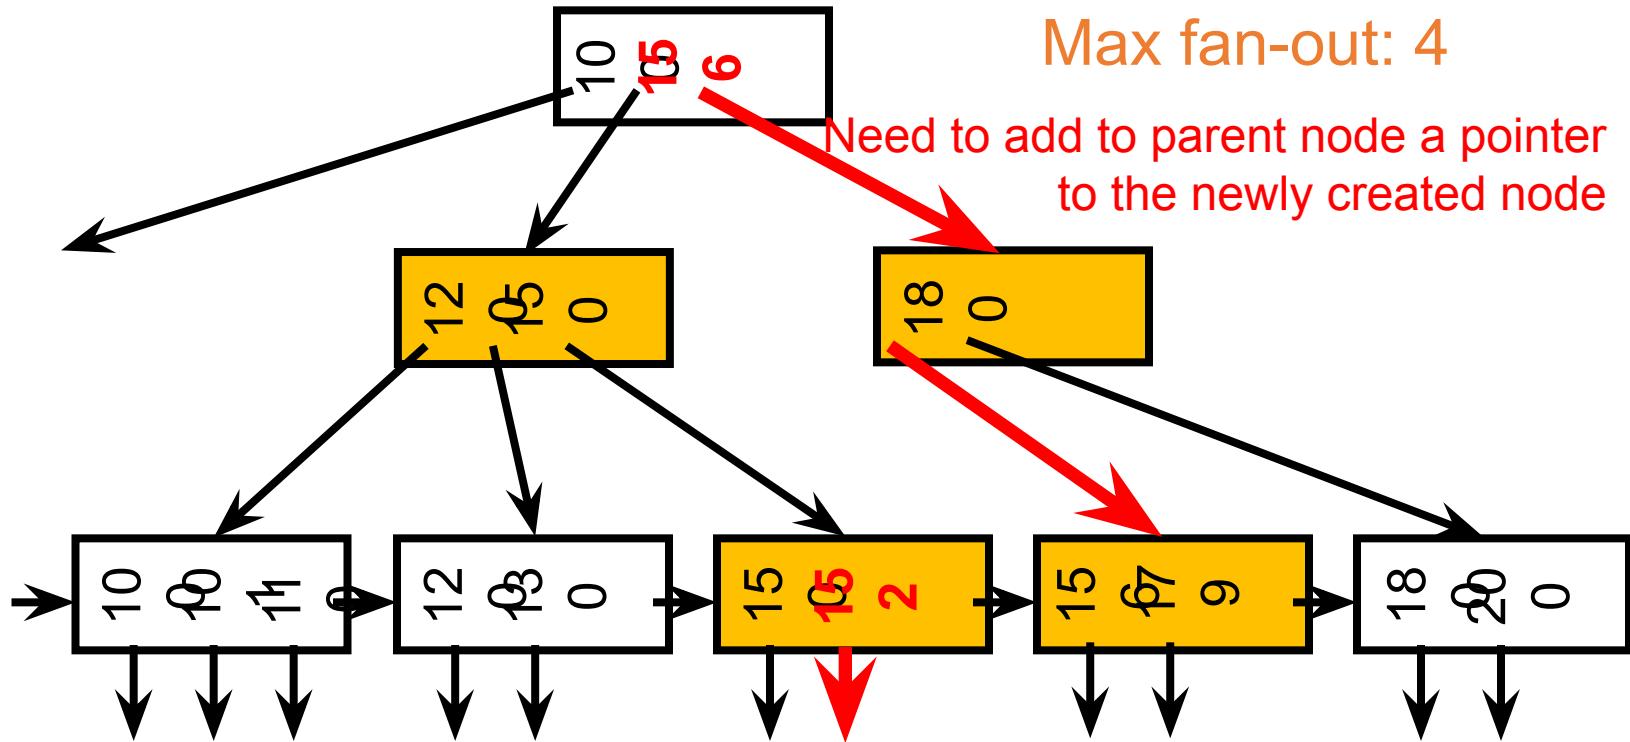
- Insert a record with search key value 152



Node Splitting



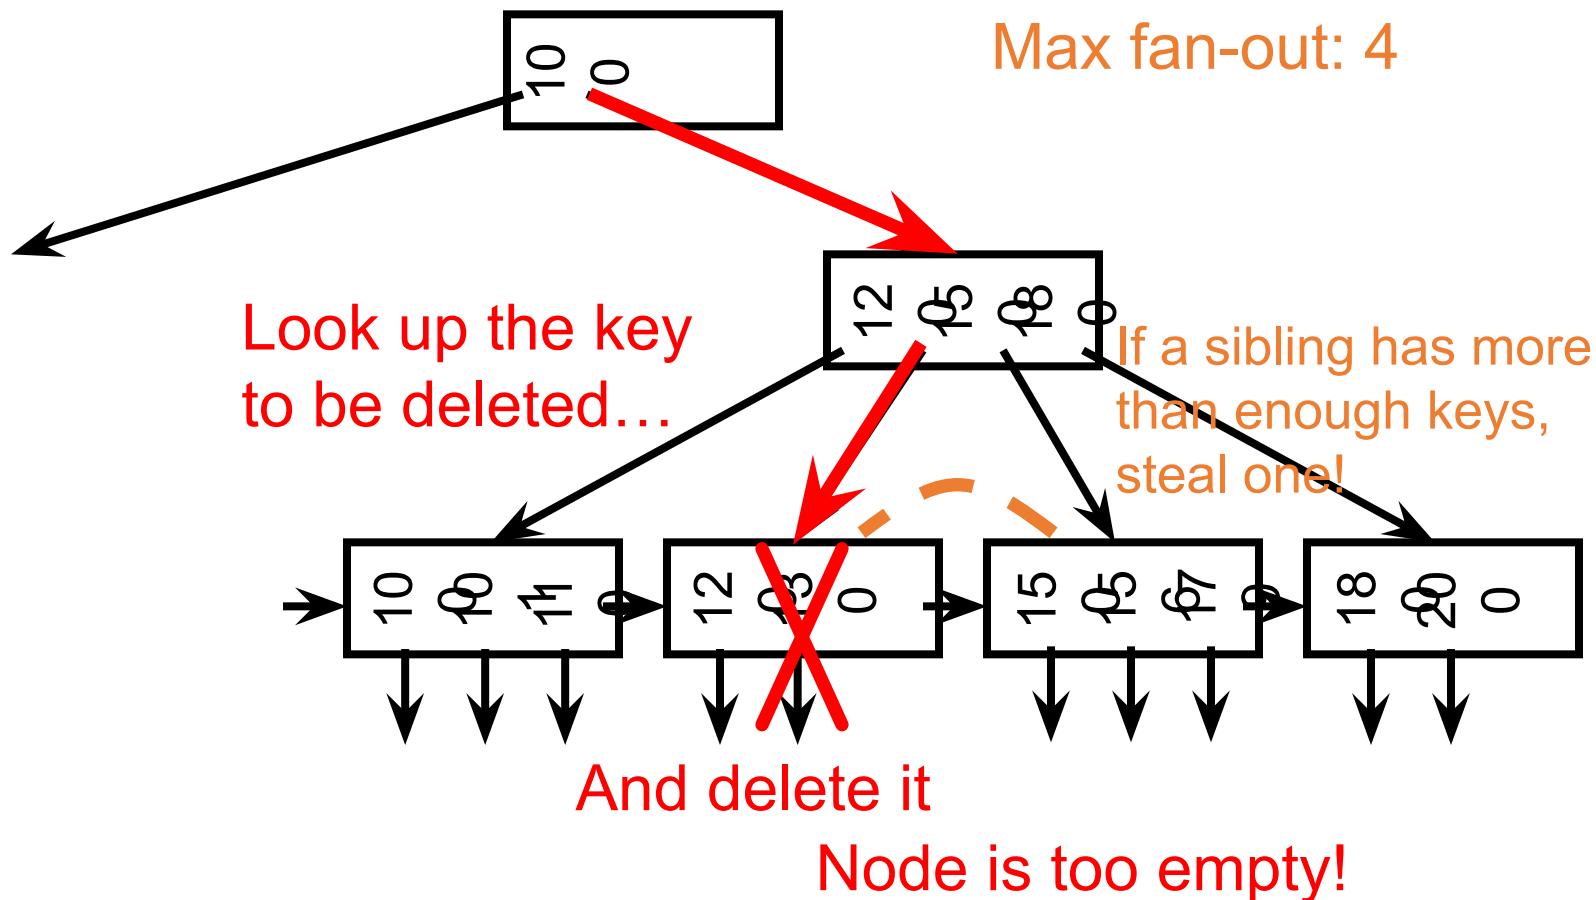
More Node Splitting



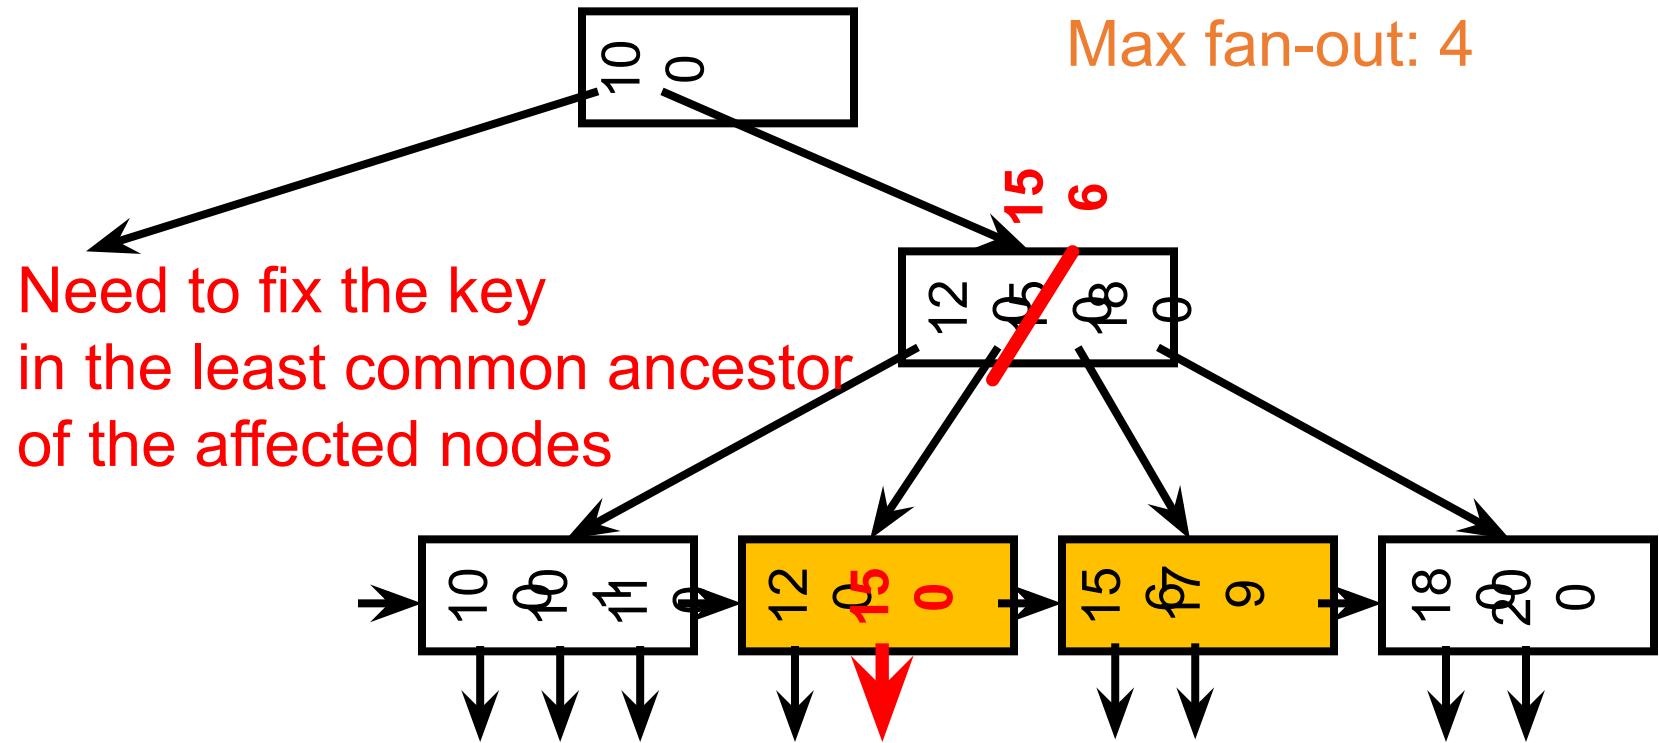
- In the worst case, node splitting can “propagate” all the way up to the root of the tree (not illustrated here)
- Splitting the root introduces a new root of fan-out 2 and causes the tree to grow “up” by one level (this is why roots can have $< f/2-1$ keys)

Deletions

- Delete a record with search key value 130



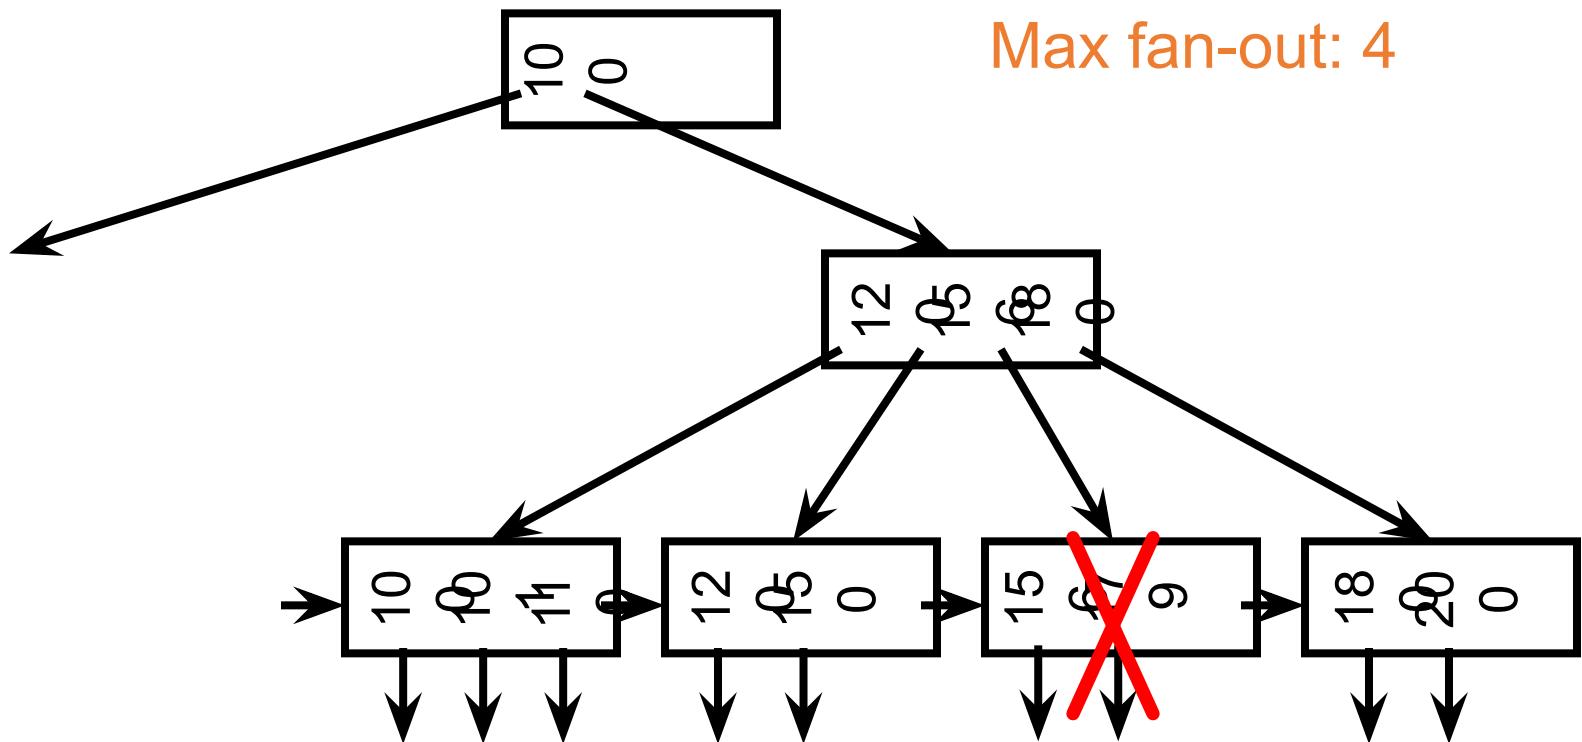
Stealing From a Sibling Node



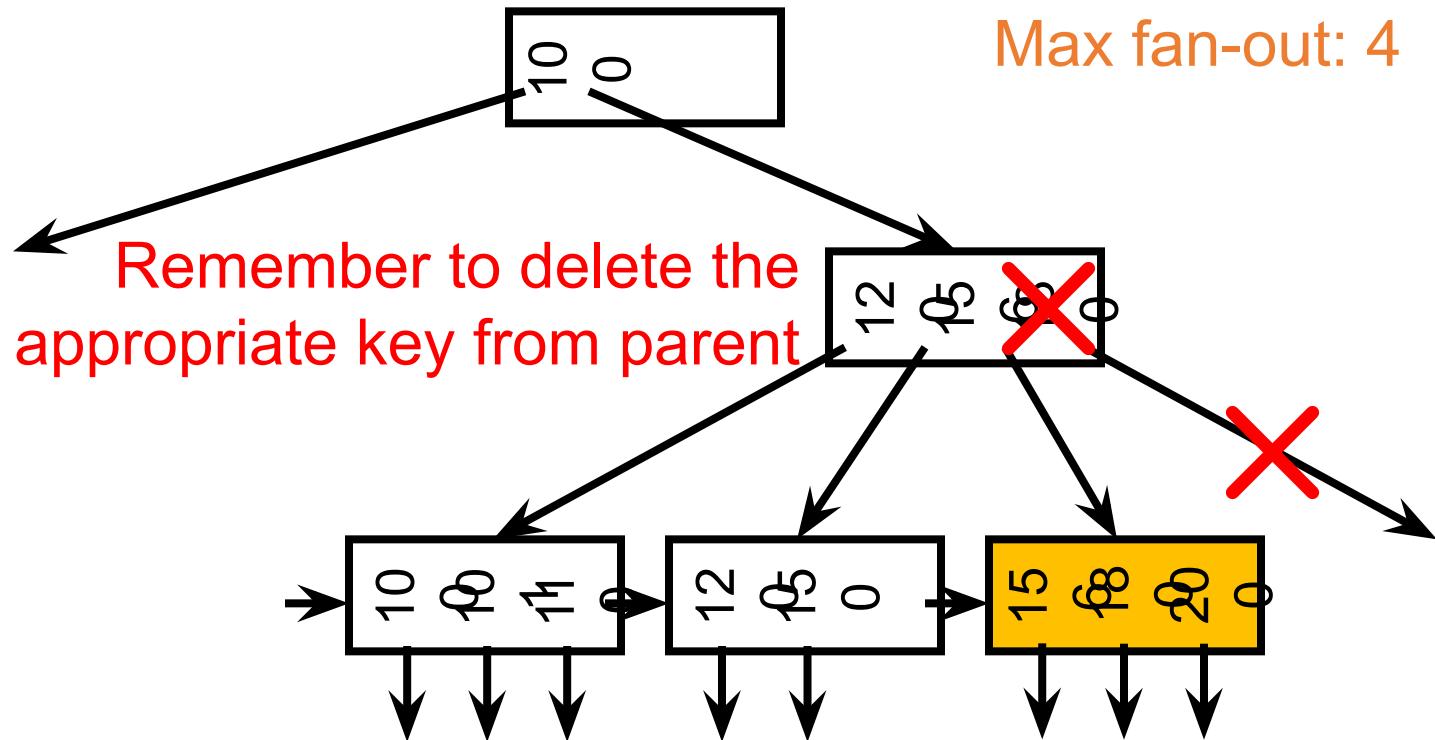
- If you are hacker, encourage you to implement the deletion subroutine of an external B+ tree. Quite challenging!

Another Deletion Example

- Delete a record with search key value 179



Merging



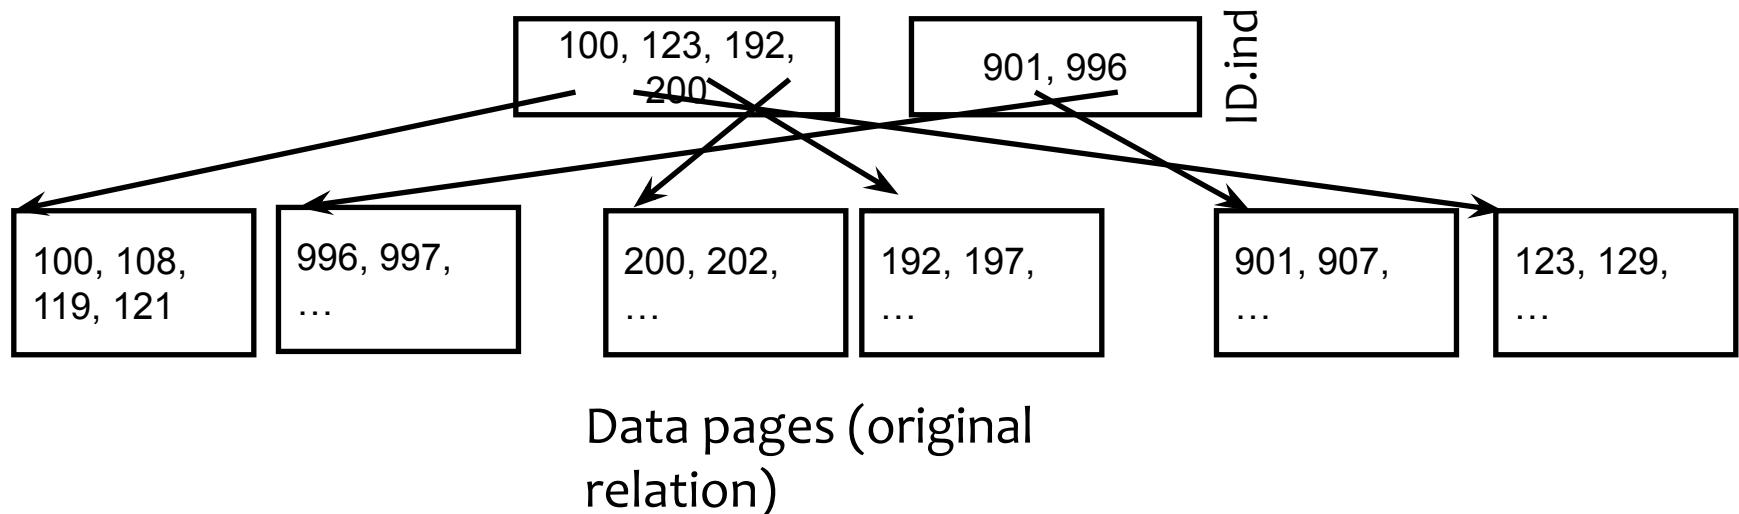
- Deletion can “propagate” all the way up to the root of the tree (not illustrated here)
- When the root becomes empty, the tree “shrinks” by one level

Performance analysis of B⁺-tree

- How many I/O's are required for each operation?
 - h , the **height of the tree**
 - Plus one or two to manipulate actual records
 - Plus $O(h)$ for reorganization (rare if f is large)
 - Minus one if we cache the root in memory
- How big is h ?
 - Roughly $\log_{\text{fanout}} N$, where N is the number of records
 - Fan-out is typically large (in hundreds)—many keys and pointers can fit into one block
 - A 4-level B⁺-tree is enough for many tables (e.g., if $f=200$, then you can accommodate 1.6B rows)

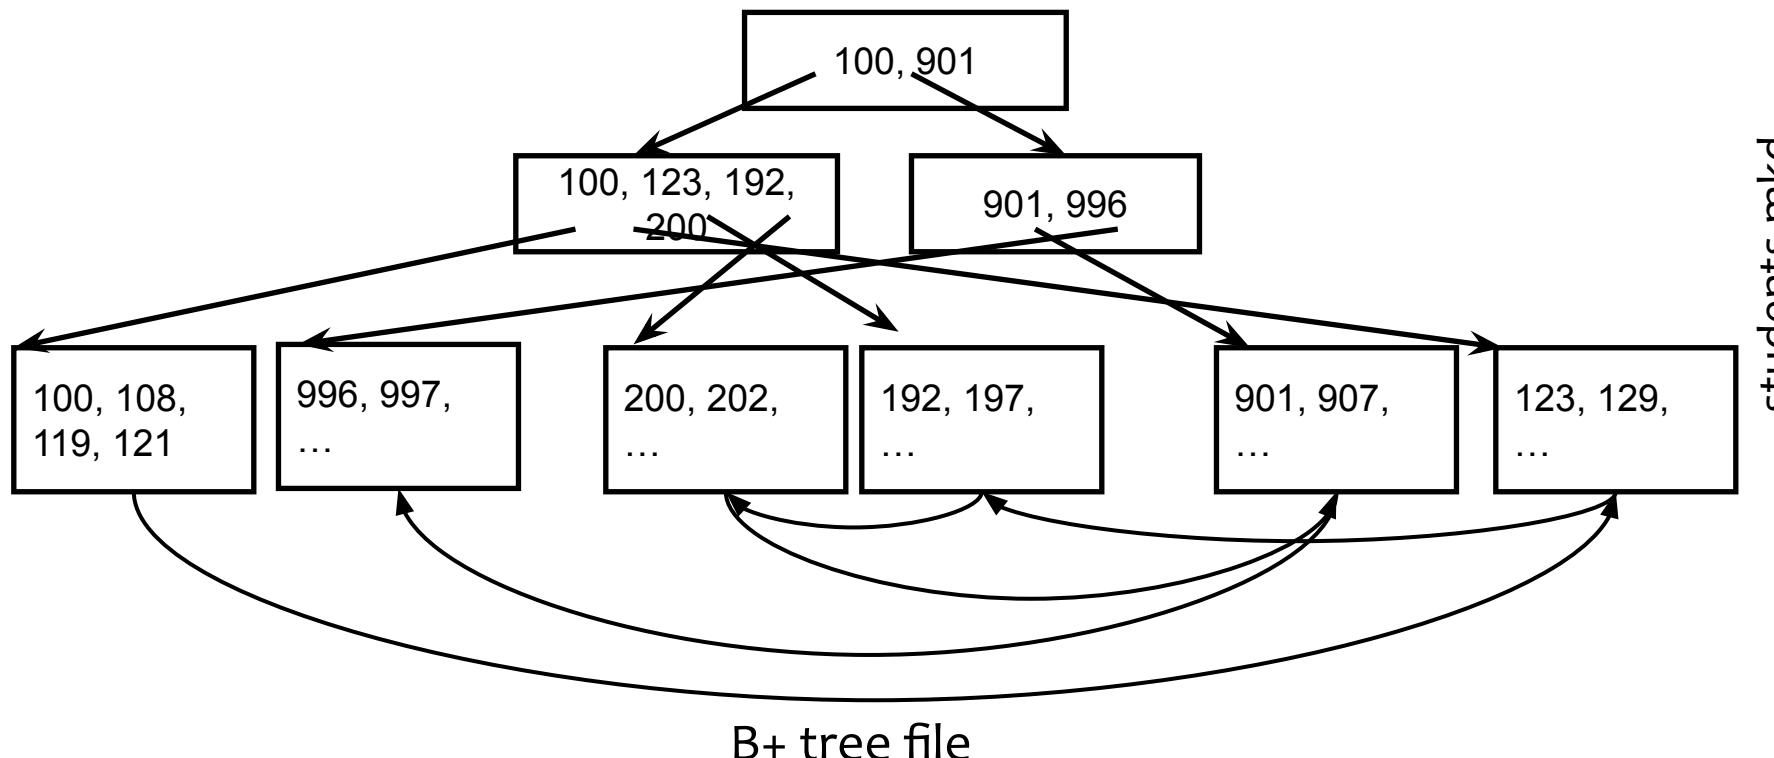
How to Keep A Table Sorted?

- Recall this key question
- Recall further note on clustered indices and page order.



How to Keep A Table Sorted?

- Recall this key question
- Recall further note on clustered indices and page order.
- Again assume leaf nodes are tuples
- Many RDBMSs use “B+ tree files” to store the tables, i.e., entire file is a B+ tree index, with leaf nodes storing tuples (instead of pointers to tuples)



Difference Between B and B+ Tree

- B-tree stores pointers to records or records in non-leaf nodes too
- **Pro:** These records can be accessed with fewer I/O's
- **Cons:**
 - Storing more data in nodes decrease fan-out and increases h
 - Records in leaves require more I/O's to access
 - Vast majority of the records live in leaves!

What Does B Stand For?

- No one really knows!
- But [Edward M. McCreight](#), co-inventor with [Rudolf Bayer](#), has a video that says:
 - They never resolved what B is but they had in mind:
 - Boeing, Bayer, and Balance

Other Common Indices

□ 2 Classes of Indices Overall

1. Tree-based: can do both lookups and range queries
 - B/B+ Trees, R Trees, Radix Tree

2. Hash-based

- Can only do look ups. Cannot do range queries.

- In practice: handle collisions

3. Many other indices: bitmap indices, probabilistic indices, suffix arrays, GiST or Inverted Index for different applications and data types.

Hash Table

		123	Milhouse	10
0	Windel	142	Bart	10
1		279	Jessica	10
2	Sherri	345	Martin	8
3	Martin	456	Ralph	8
4		512	Nelson	10
5	Ralph, Bart	679	Sherri	10
6		697	Terri	10
7	Nelson	857	Lisa	8
8	Milhouse	912	Windel	8
9	Terri	997	Jessica	8
10	Lisa			
11	Jessica			

students.mkd

Outline For Today

1. Recap & Finish Physical File Organization Designs
 - Row-oriented Physical Design (Recap)
 - Column-oriented Physical Design
 - Hybrid (PAX) Physical Design
 - Designs for Variable-length Fields
 - Designs for NULLs
2. Database Indices
 - 5 Index Designs in Increasing Level of Robustness
 - Using Indices In Practice

Using Indices In Practice (1)

- Indices can be defined on one or more attributes:
 - `CREATE INDEX NameIndex ON User(Lastname,Firstname);`
 - I.e., B+’s keys are (Lastname, Firstname) pairs and tuples are sorted first by LastName and then Firstname.
 - This index would be useful for these queries:

```
select * from User where Lastname = ‘Smith’
```

```
select * from User where Lastname = ‘Smith’ and Firsntname=‘John’
```

- But not this query:

```
select * from User where Firsntname=‘John’
```

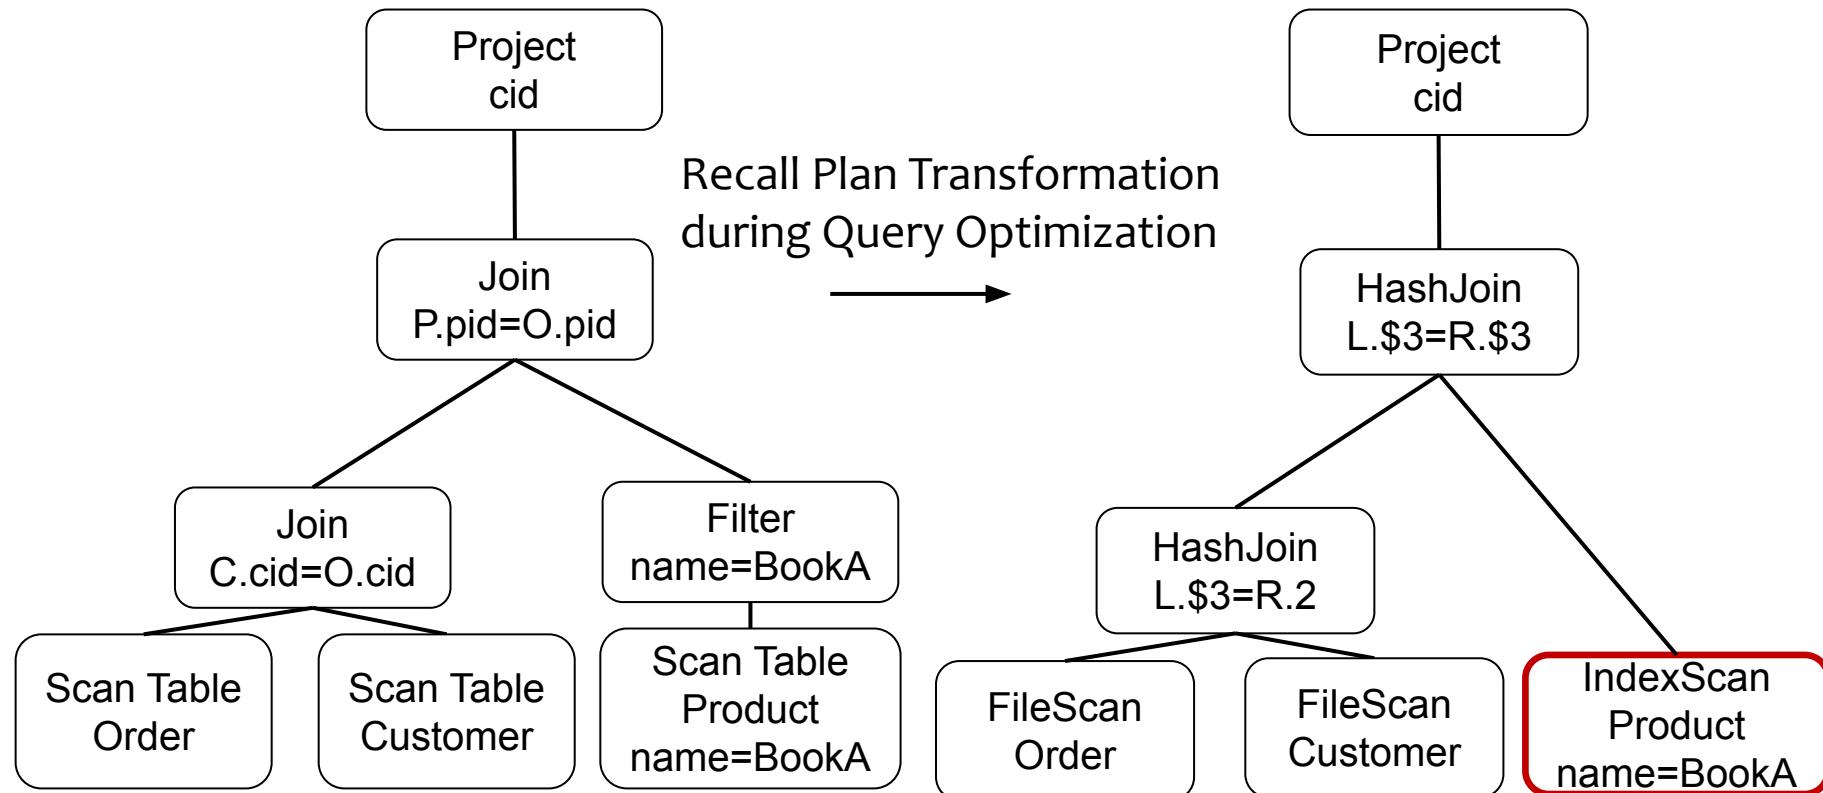
- Many systems use indices by default on the primary key
- Many systems use indices to implement UNIQUE constraints

```
CREATE TABLE Students(  
    studentID int,  
    sinNumber varchar(16) UNIQUE  
    PRIMARY KEY (studentID))
```

Will create 2 B+ indices:
1) on studentID;
2) 2) on sinNumber

Using Indices In Practice (2)

- Users only create indices. They do not refer to indices in queries.
- **Pro:** Some user queries will get much faster
 - B/c RDBMSs use indices during query evaluation
 - Ex: IndexScan operators, or IndexMergeJoin (in Oracle) or IndexNestedLoopJoin etc.



Using Indices In Practice (3)

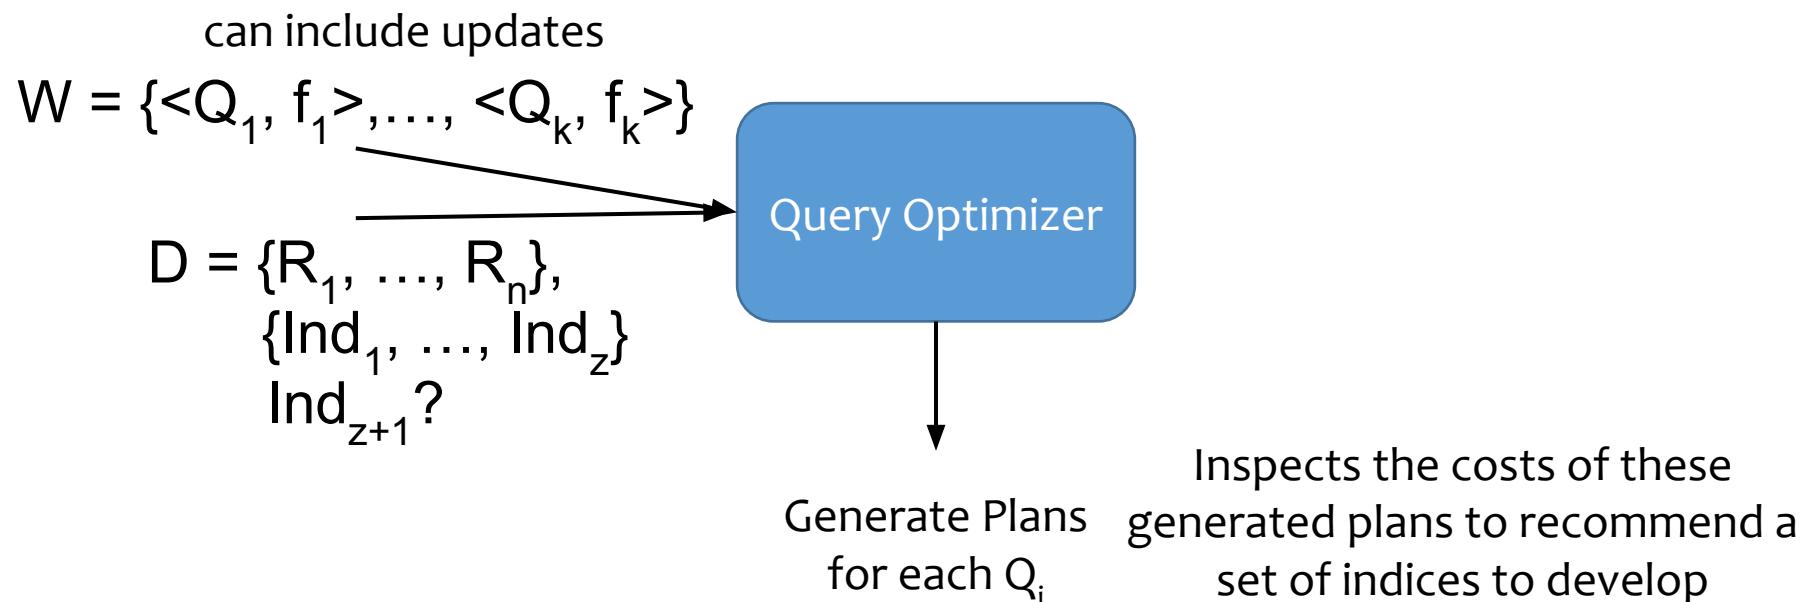
- Con: Updates will get slower because indices need to be maintained
- Q: How should users pick indices given a workload W (i.e., the set of queries an application asks and their frequencies)
- General Guideline:
 - Profile slow queries. Check if they have =, <, ≤, >, ≥ predicates

```
SELECT * FROM R WHERE A = value;  
SELECT * FROM R WHERE A = value AND B = 27;  
SELECT * FROM R, S WHERE R.A = S.C;  
SELECT * FROM S WHERE D > 50;
```

- E.g., above indices on R.A, R.A and R.B multicolumn, S.C, S.D are possible indices that can speed queries
- But one should weigh these benefits against slow downs due to updates

Using Indices In Practice (4)

- Many RDBMSs have “Physical Design Advisor” (PDA) tool
- Input: Database D (w/ existing indices), workload W
- Output: A set of recommended indices
- Internally PDA does a “what if” analysis:
 - Uses Query Optimizer & inspects the estimated runtimes/costs of plans the system would use for queries in W with & without additional indices



The Halloween Problem: An Interesting Note On Challenging Problems a DBMS has to solve

- Story from the early days of System R:

```
UPDATE Payroll  
SET salary = salary * 1.1  
WHERE salary >= 100000;
```

- There is a B⁺-tree index on *Payroll(salary)*
- The update never stopped (why?)

- Why?

- How could you try to solve this if you implemented System R.