

CS 348 Lectures 15-16

Query Optimization

Semih Salihoglu

Nov 9-11 2021



UNIVERSITY OF
WATERLOO | **DSg** Data
Systems
Group

Announcements

- Prof. Thomas Neumann DSG Seminar (Don't miss!)
- Nov 15th at 10:30am
- Register [here](#)
- Send me an email if you want to join the student meeting



Data Systems Seminar Series

Monday, November 15, 2021 • 10:30 a.m. EST

Adaptive Join Order Optimization using Search Space Linearization

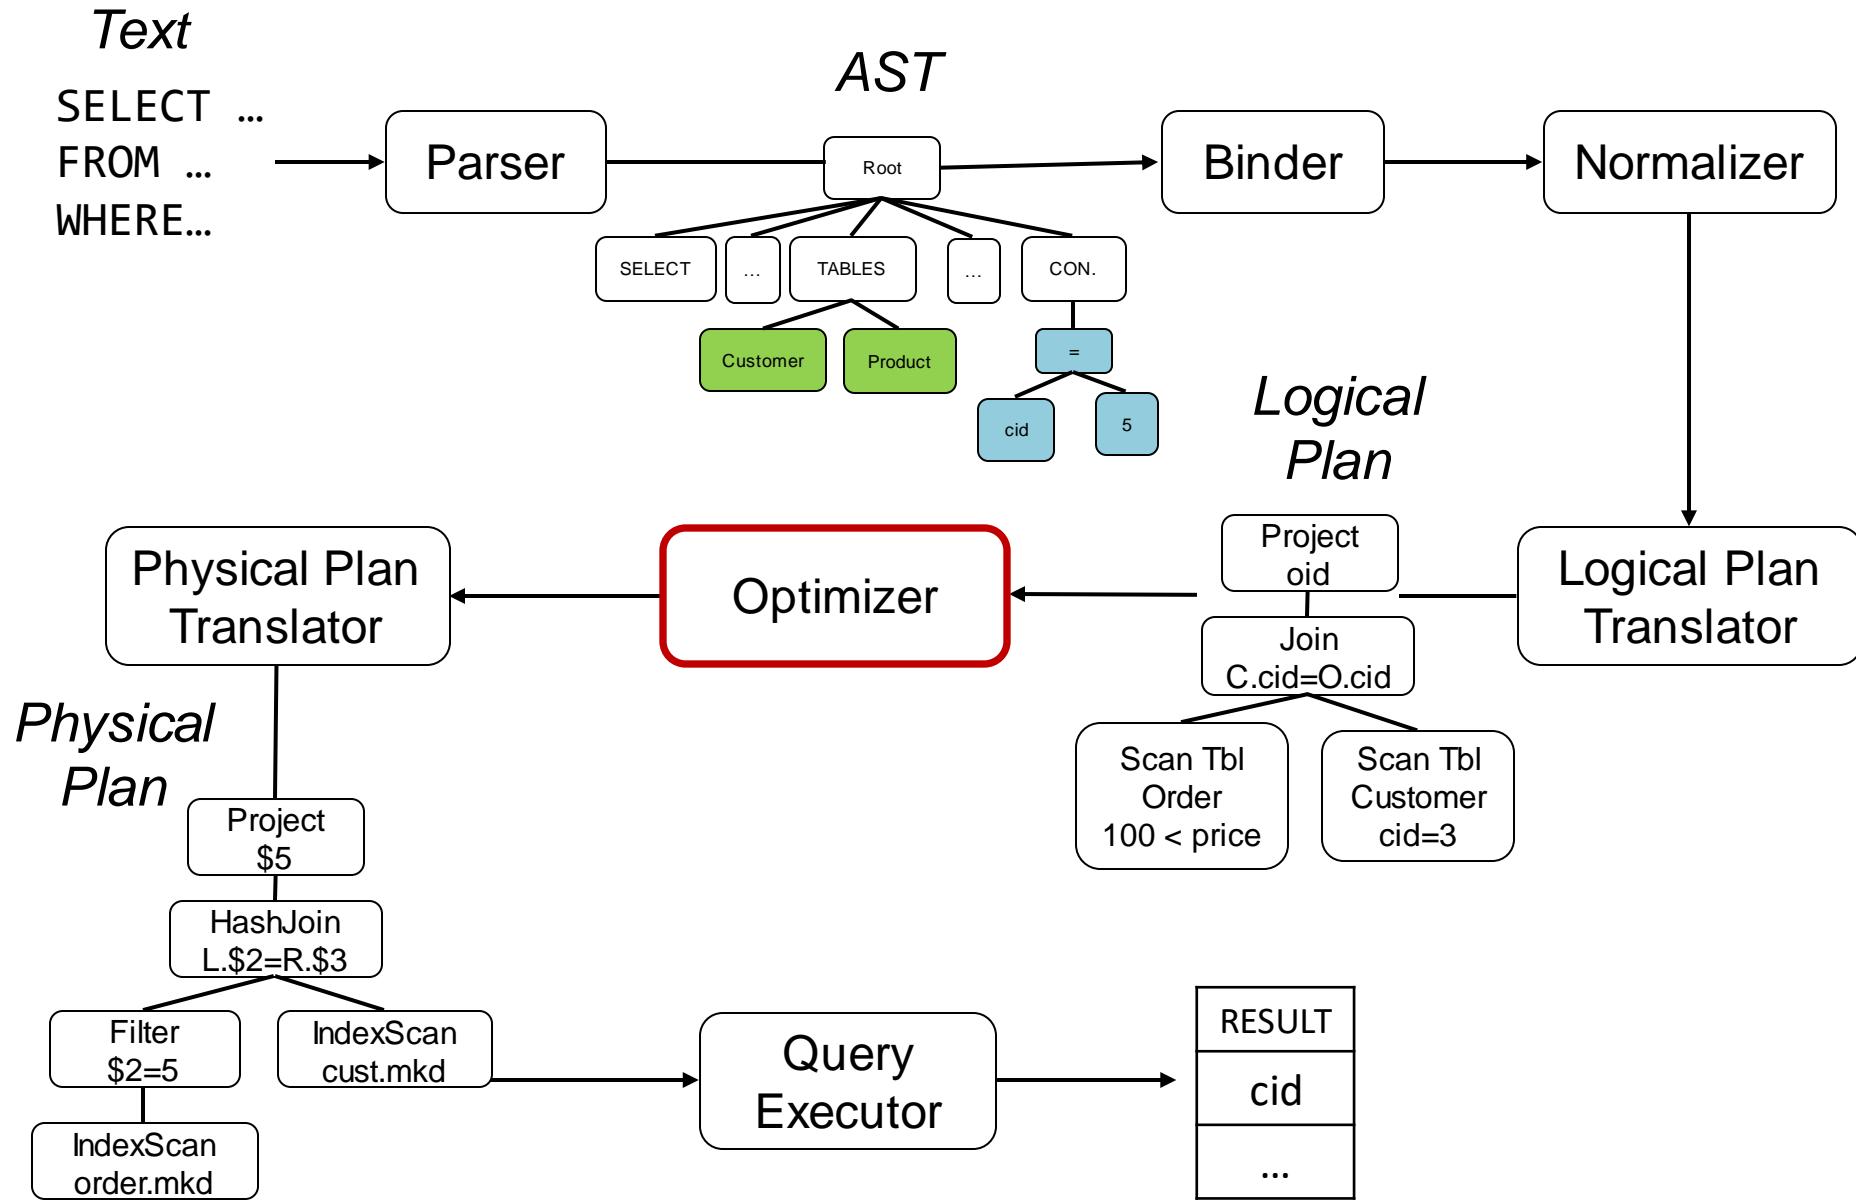
Join ordering is one of the core problems of query optimization, as differences in join order can affect the execution time of queries by orders of magnitudes. Unfortunately, the problem is NP hard in general, and real-world queries can join hundreds of relations, which makes exact solutions prohibitive expensive. In this talk we show how to tackle the join ordering problem by using a search space linearization technique. This adaptive optimization mechanism allows for a smooth transition from guaranteed optimality to a greedier approach, depending on the size of problem. In practice, a surprisingly large number of queries can be solved optimally or near optimally, with very low optimization times even for hundreds of relations.

Thomas Neumann is a full professor in the Department of Computer Science at the Technical University of Munich. His research interests are in the areas of database systems, query processing, and query optimization. In 2020, he received the Gottfried Wilhelm Leibniz Prize, which is considered the most important research award in Germany.

Please register • https://uwaterloo.zoom.us/meeting/register/tJAgdeCtqzwjGtCPTmlo5uMy_lOJFuRPkqjL

 UNIVERSITY OF WATERLOO | **DSg**
Data Systems Group

Recall: Overview of Compilation Steps



Outline For Today

1. Goal of Query Optimization and Overview of Techniques
2. Cost-based Optimization Principles & Cardinality Estimation
3. Cost-based DP Logical Join Plan Optimizer
4. Rule-based Optimizations/Transformations
5. Final Remarks on Query Optimization & Query Processing

Outline For Today

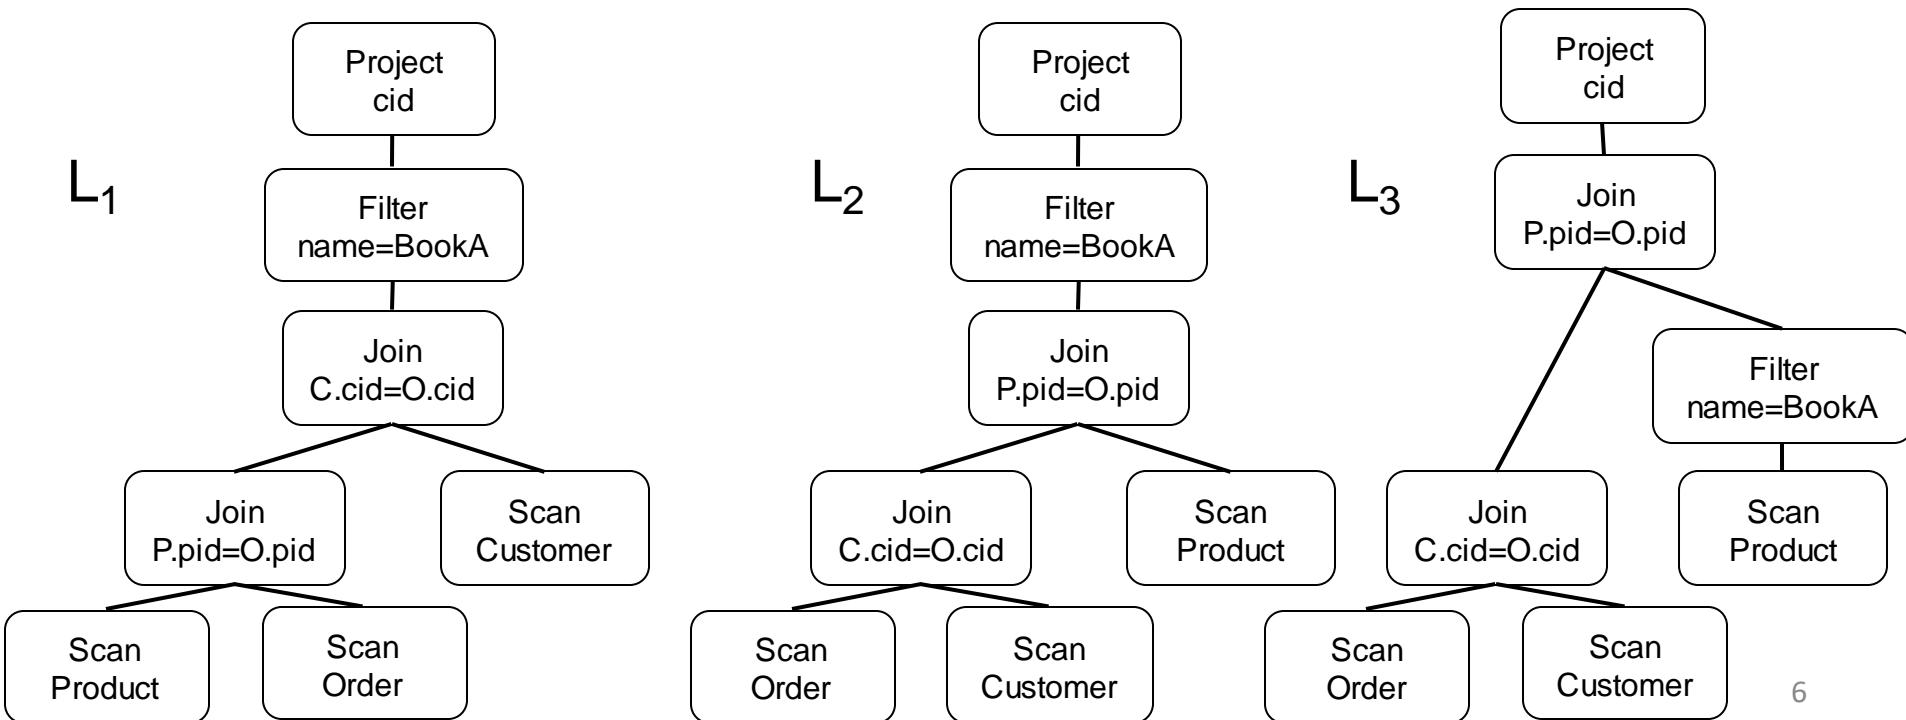
1. Goal of Query Optimization and Overview of Techniques
2. Cost-based Optimization Principles & Cardinality Estimation
3. Cost-based DP Logical Join Plan Optimizer
4. Rule-based Optimizations/Transformations
5. Final Remarks on Query Optimization & Query Processing

Goal of Query Optimization (1)

- Recall ultimately a *physical plan* executes to answer a query
- Given a query Q, many equivalent physical plans exist:
 1. Many equivalent logical plans exist

```
SELECT cid  
FROM Customer C, Order O, Product P  
WHERE C.cid = O.cid AND O.pid = P.pid  
      AND P.name = BookA
```

Logical Plans:

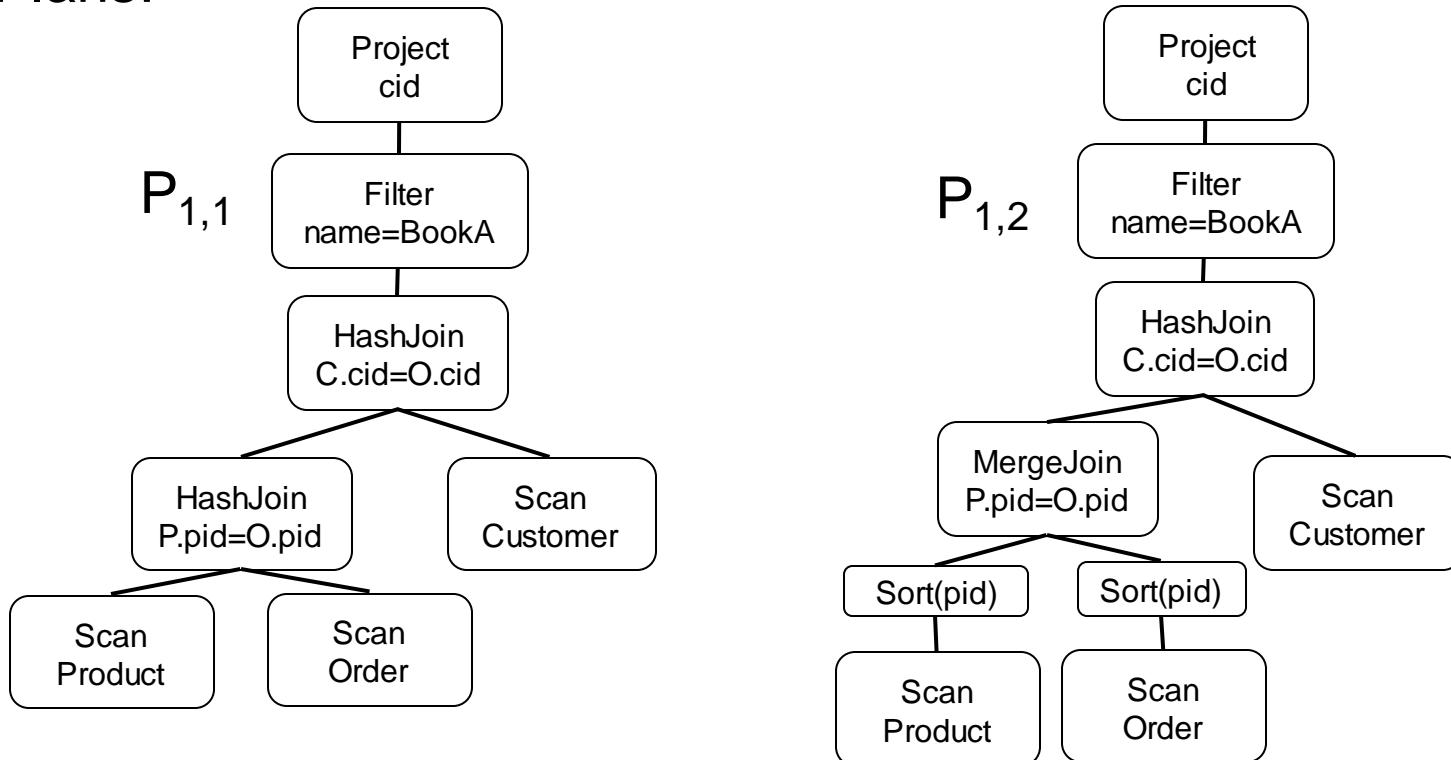


Goal of Query Optimization (1)

- Recall ultimately a *physical plan* executes to answer a query
- Given a query Q, many equivalent physical plans exist:
 1. Many equivalent logical plans exist
 2. Each logical plan can have many equivalent physical plans.

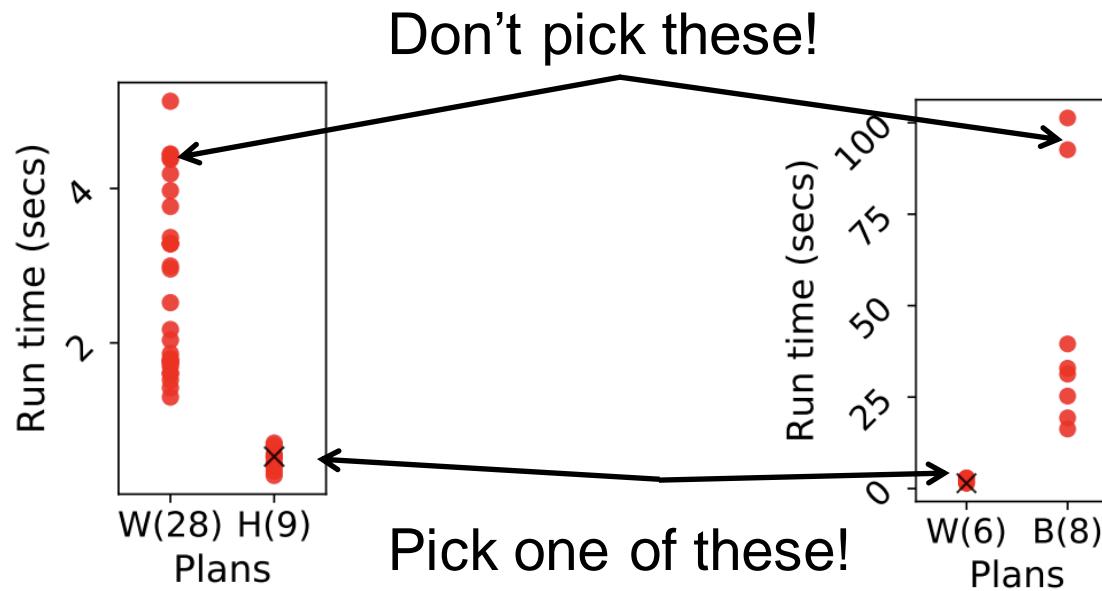
```
SELECT cid  
FROM Customer C, Order O, Product P  
WHERE C.cid = O.cid AND O.pid = P.pid  
      AND P.name = BookA
```

Physical Plans:



Goal of Query Optimization (2)

- Ultimately: Given Q , pick the “best” physical plan for Q :
 - Best: often means fastest, could mean “cheapest”
- DBMS developers are more humble:
 - Pick a reasonably good plan. Do not pick a very bad plan!
 - Example plan spectrum of join-heavy queries



Overview of Query Opt. Techniques

1. Enumerate a logical plan space (often

enumerates all join orders)

(extended) relational
algebraic expressions $\sqsubset L_1, L_2, \dots, L_k$

2. For one or more of L_i , (optionally)

enumerate a physical plan space:

$P_{i,1}, P_{i,2}, \dots, P_{i,t}$

3. Pick the best $P_{i,1}$

A common approach:

- Step 1 is cost-based or hybrid
- Step 2 is rule-based

Options for steps 1 & 2:

- i. Rule-based
- ii. Cost-based
- iii. Hybrid rule/cost-based

Outline For Today

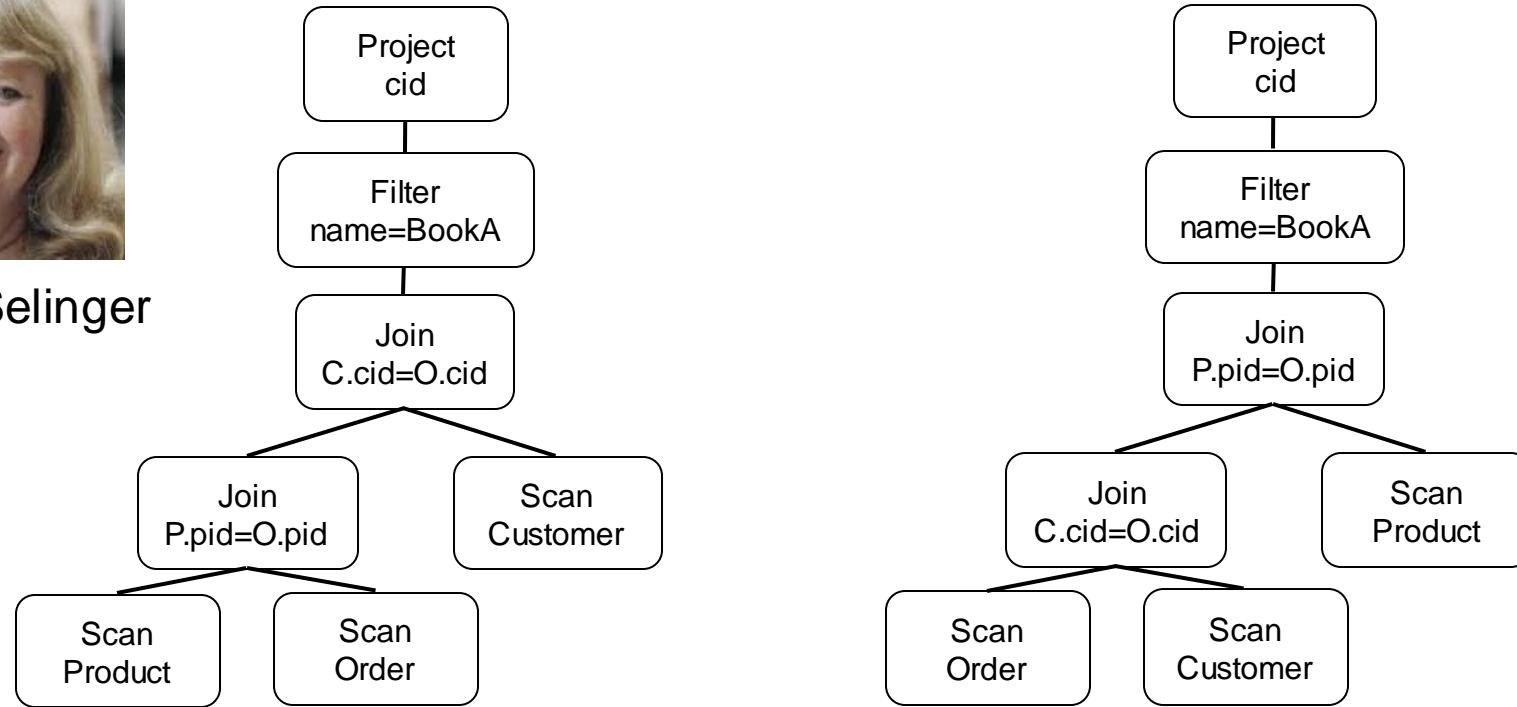
1. Goal of Query Optimization and Overview of Techniques
2. Cost-based Optimization Principles & Cardinality Estimation
3. Cost-based DP Logical Join Plan Optimizer
4. Rule-based Optimizations/Transformations
5. Final Remarks on Query Optimization & Query Processing

Cost-based Optimization Principles

- System R ('70s): First prototype relational DBMS from IBM



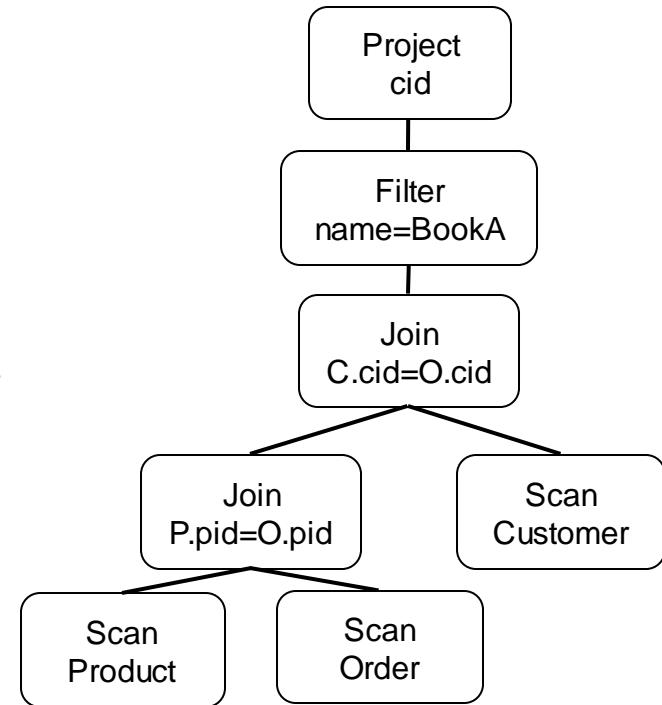
Patricia Selinger



- Give each enumerated log/phy plan, e.g., L_i , a $\text{Cost}(L_i) = c_i$
- Cost is the estimate of the system for how good/bad L_i is.
- Pick min cost plan

Cost-based Optimization Principles (1)

- Naturally: cost definition is broken into costs of operators.
 - i.e: $\text{Cost}(L_i) = \sum_j \text{cost}(o_j \in L_i)$
- Example cost metrics or components:
 - # I/Os a plan will make
 - # tuples that will pass through operators
 - # runtime of algorithm o_j is running
 - e.g., nested loop join of R, S: $|R|^*|S|$
 - Combination of above
- For any reasonable metric:
 - Need to estimate cardinality, i.e., size, of tuples o_j will process



The (notorious) cardinality estimation problem!

Cardinality Estimation

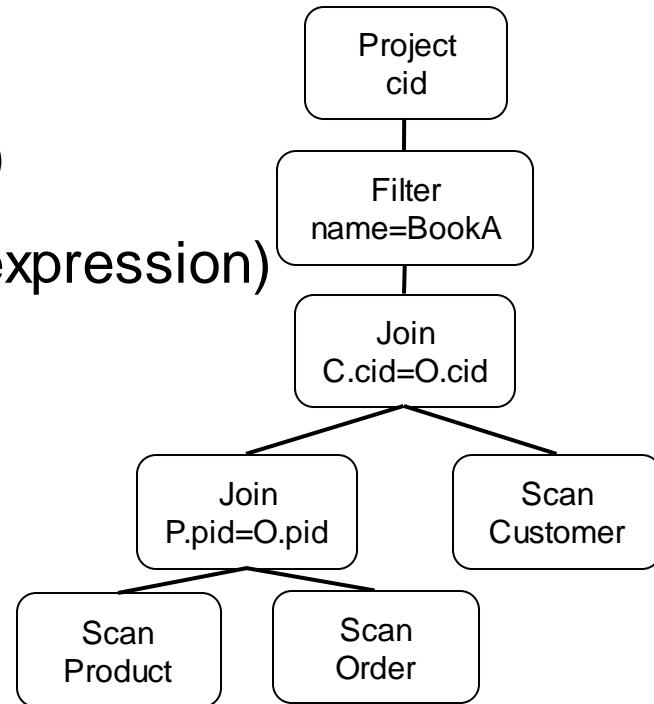
- Given a database

1. $D: R_1(A_{1,1}, \dots, A_{1,m_1}), \dots, R_n(A_{n,1}, \dots, A_{n,m_n})$
2. A (sub-) query Q (a relational algebra expression)

What is the $|Q|$?

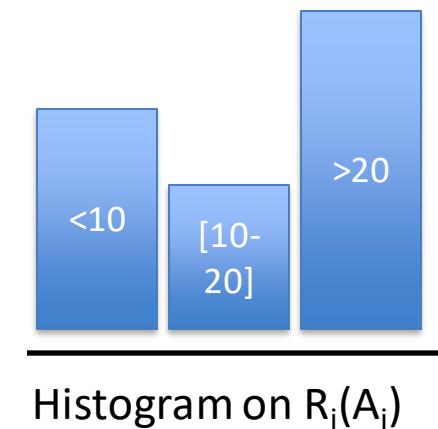
- E.g:

- $\sigma_{name=BookA}(Product)$?
- $Product \bowtie Order$?
- $\sigma_{name=BookA}(Product \bowtie Order \bowtie Customer)$?



2 High-level Card. Estimation Techniques

1. Sampling-based:
 - While optimizing Q, sample relations to make an estimate
2. Summary/statistics-based:
 - Use statistics about D to make estimates
 - Possible statistics:
 - $|R_i|$: size of each relation
 - $|\pi_{A_j}(R_i)|$ # distinct values in column A_j
 - Histograms: Distribution of values on A_j
 - Also use known constraints:
 - E.g: FK constraint from R to S: $|R \bowtie S| = |R|$
 - 2 common *simplification* assumptions (no other good reason):
 - (i) uniformity; (ii) independence



Example Statistics-based Estimation Techniques

Selections with Equality Predicates

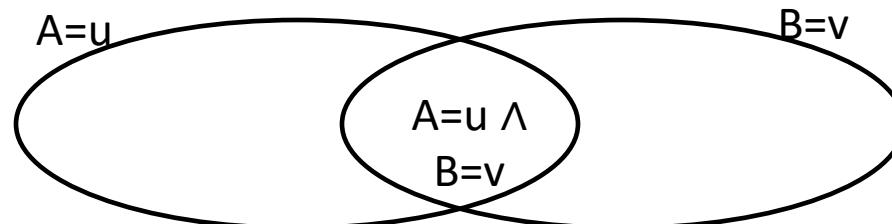
- $Q: \sigma_{A=v} R$
 - Suppose the following information is available:
 - Size of R : $|R|$
 - Number of distinct A values in R : $|\pi_A R|$
 - Assumptions:
 1. Values of A are *uniformly distributed* in R
 2. $v \in |\pi_A R|$
 - $|Q| \approx |R| / |\pi_A R|$
 - Selectivity factor of $(A = v)$ is $1 / |\pi_A R|$
 - Ex: $|Product| = 1000$, $|\pi_{name}(Product)| = 50$
 - $\sigma_{name=BookA} Product: 1000/50 = 20$
- wild assumption, often doesn't hold*
- fair assumption, often holds (b/c users search things they put in the db)*

Conjunctive Predicates

- $Q: \sigma_{A=u \wedge B=v} R$
- Additional assumption:
 3. $(A = u)$ and $(B = v)$ are *independent*
 - Counter example: age and salary
- $|Q| \approx |R| / |\pi_A R| \cdot |\pi_B R|$
 - Reduce total size by all selectivity factors
 - Directly derived from standard probability rules:
 - $\Pr(E_1) = p_1$, and $\Pr(E_2) = p_2$ and E_1 and E_2 are independent:
 - $\Pr(E_1 \wedge E_2) = p_1 * p_2$
 - Ex: $\Pr(\text{heads} \wedge \text{dice}=6) = 1/2 * 1/6 = 1/12$
- Ex: $|Prod| = 1000$, $|\pi_{name}(Prod)| = 50$, $|\pi_{merchant}(Prod)| = 4$
 - $\sigma_{name=BookA \wedge merchant=B\&N} Product: 1000/(50*4) = 5$

Negated and Disjunctive Predicates

- $Q: \sigma_{A \neq v} R$
 - $|Q| \approx |R| \cdot (1 - 1/|\pi_A R|)$
 - Selectivity factor of $\neg p$ is $(1 - \text{selectivity factor of } p)$
- $Q: \sigma_{A=u \vee B=v} R$
 - $|Q| \approx |R| \cdot (1/|\pi_A R| + 1/|\pi_B R|) ?$
 - No! Tuples satisfying $(A = u)$ and $(B = v)$ are counted twice
 - Use only for $\sigma_{A=u \vee A=v} R$ (b/c then $A=u$ and $A=v$ are disjoint)
 - $|Q| \approx |R| \cdot (1/|\pi_A R| + 1/|\pi_B R| - 1/|\pi_A R||\pi_B R|)$
 - Inclusion-exclusion principle from probability

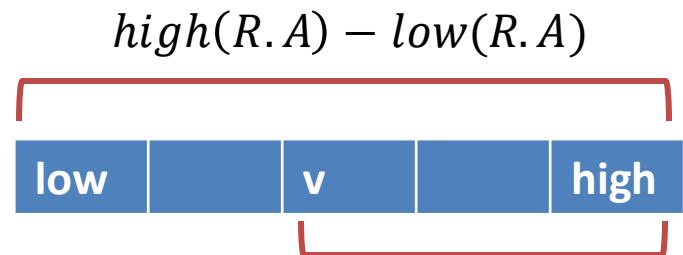


Range Predicates

- $\sigma_{A>u} R?$
- Case 1: Suppose the DBMS knew actual projection values:
 - Then range queries are a generalization of $\sigma_{A=u \vee A=\nu} R$
 - $\sigma_{A>u} R = |Q| \approx |R| \cdot \left(\frac{|\#vals>u|}{|\pi_A R|} \right) ?$
 - E.g: A was an int column and $|\pi_A R| = \{1, 2, 3, 4, 5\}$
 - $\sigma_{A>2} R = |R| * 3/5$

Case 2 of Range Predicates

- Case 2: We don't know actual values
- Not enough information!
 - Just pick a *magic constant*, e.g., $|Q| \approx |R| \cdot \frac{1}{3}$
- With more information
 - Largest $R.A$ value: $\text{high}(R.A)$
 - Smallest $R.A$ value: $\text{low}(R.A)$
 - $|Q| \approx |R| \cdot \frac{\text{high}(R.A) - v}{\text{high}(R.A) - \text{low}(R.A)}$
 - In practice: sometimes the second highest and lowest are used
 - The highest and the lowest are often used by inexperienced database designer to represent invalid values!



Equi-Join of Two Relations (1)

- $Q: R(A, B) \bowtie S(A, C)$
- Assumption: containment of value sets:
 - Every tuple in the “smaller” relation (one with fewer distinct values for the join attribute) joins with a tuple in the other relation
 - That is, if $|\pi_A R| \leq |\pi_A S|$ then $\pi_A R \subseteq \pi_A S$
 - Certainly not true in general
 - But holds in the common case of foreign key joins
- $|Q| \approx \frac{|R| \cdot |S|}{\max(|\pi_A R|, |\pi_A S|)}$
 - Selectivity factor of $R.A = S.A$ is $1 / \max(|\pi_A R|, |\pi_A S|)$

Equi-Join of Two Relations (2)

- Example:

R	
A	B
a ₁	b ₁
a ₁	b ₂
a ₁	b ₃
a ₁	b ₄

S	
A	C
a ₁	c ₁
a ₁	c ₂
a ₂	c ₃
a ₂	c ₄

$$\pi_A R = \{a_1\}$$

$$\pi_A S = \{a_1, a_2\}$$

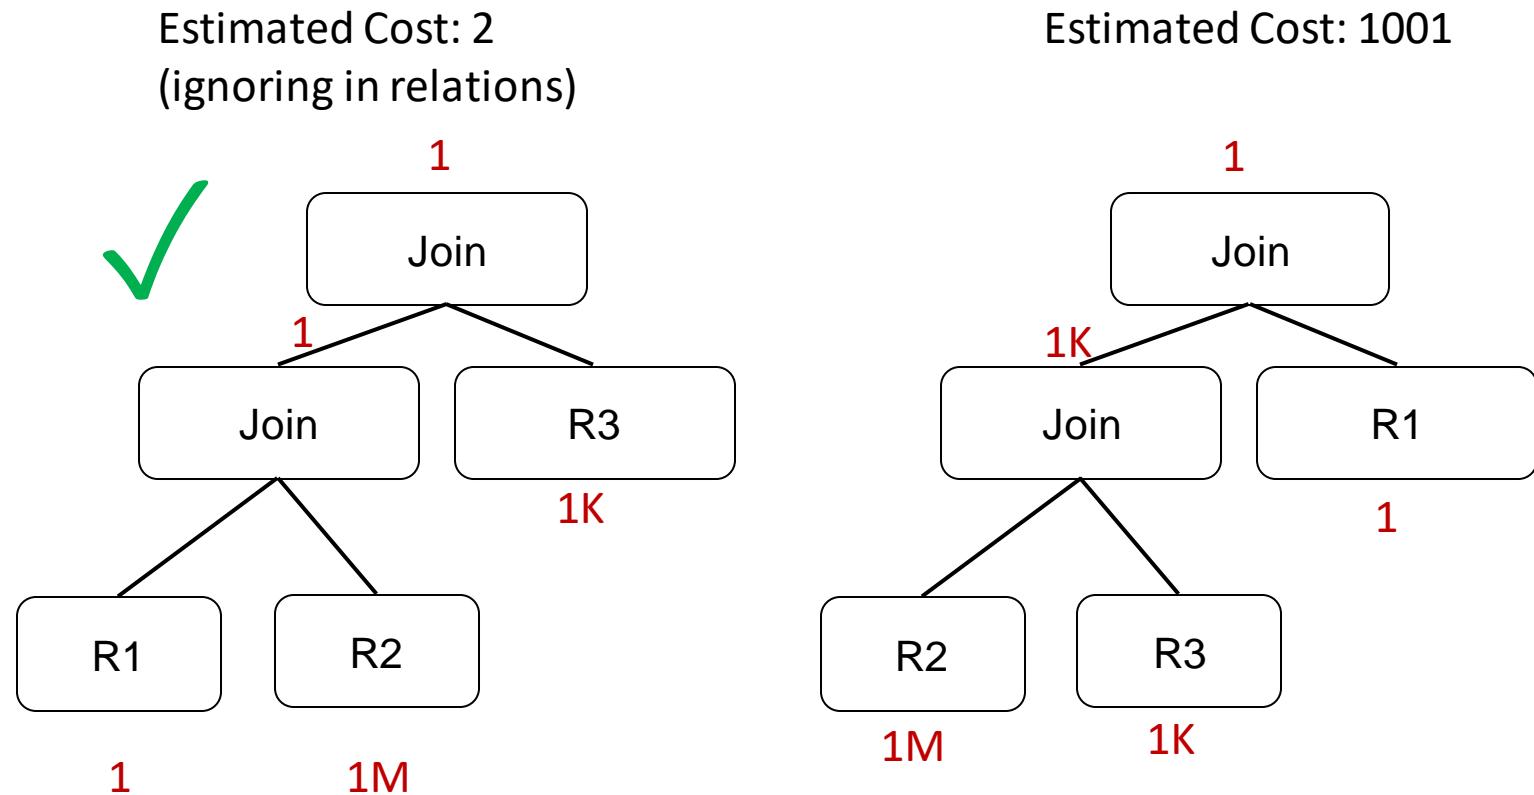
- $|Q| \approx \frac{|R| \cdot |S|}{\max(|\pi_A R|, |\pi_A S|)} = 4 \times 4 / 2 = 8$ (correct)
- If we had picked $\min(|\pi_A R|, |\pi_A S|)$, then we'd over-estimate
 - Intuitively a fraction of tuples from the larger-domain table will join with each tuple from smaller-domain table (not vice versa)

Other Estimations Techniques

- Using similar ideas, we can estimate the size of projection, duplicate elimination, union, difference, aggregation (with grouping)
- Lots of assumptions and very rough estimation
 - Accurate estimate is not needed
 - Maybe okay if we overestimate or underestimate consistently
- In practice: Very very difficult but very important for the optimizer.
 - B/c: ultimate goal is to help estimate costs of operators & plans
 - If we badly underestimate an expression => may lead to bad plans

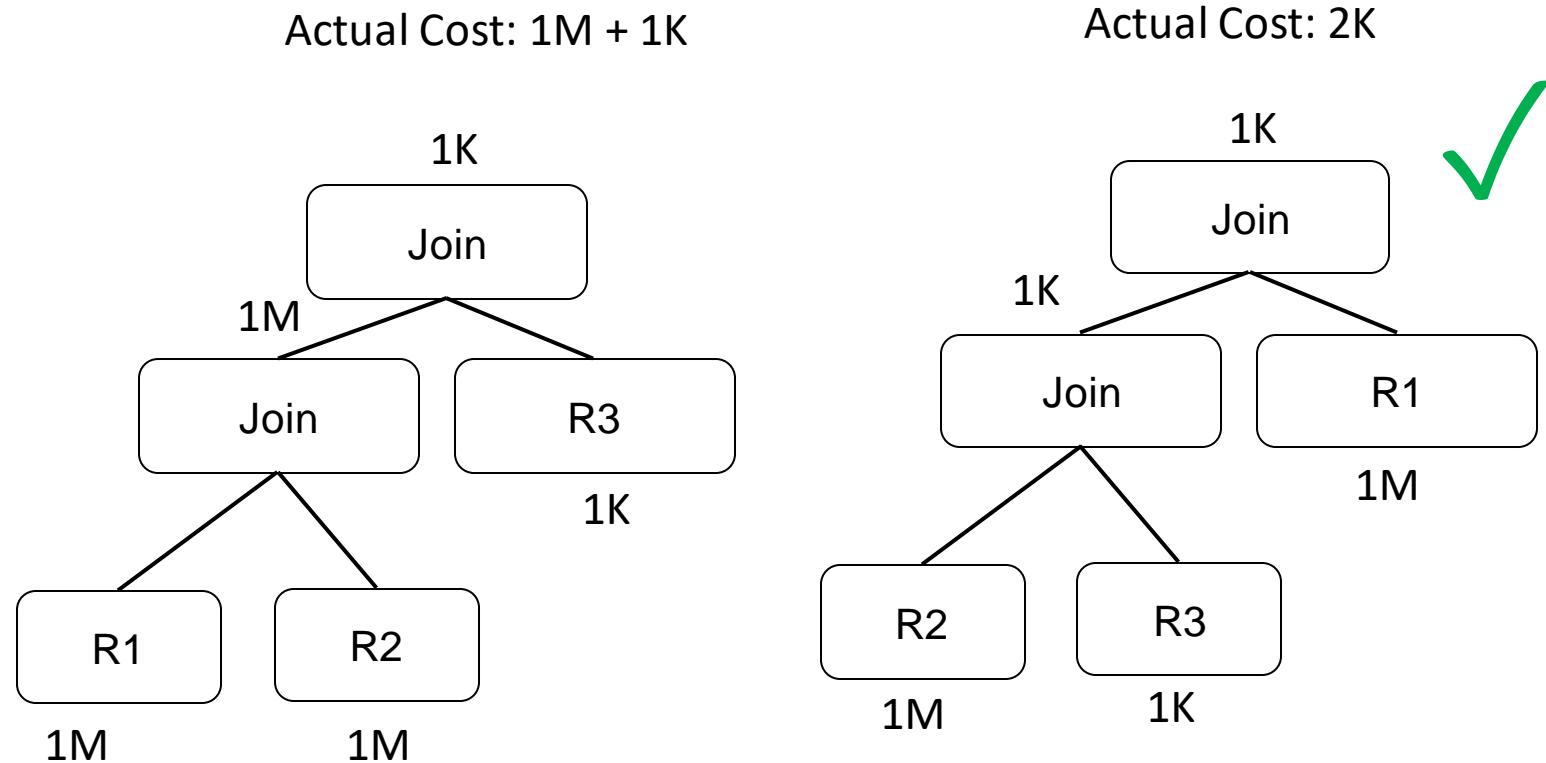
Example Poor Optimizer Choice

- Suppose $\text{cost}(o_j)$: # input tuples processed
- $\sigma_{p1}(R1) \bowtie R2 \bowtie R3$
- Suppose $\sigma_{p1}(R1) = 1M$ but DBMS underestimates as 1
- Suppose $|R2| = 1M$ and $|R3| = 1K$
- Suppose output of join has the size of the minimum input relation



Example Poor Optimizer Choice

- Suppose $\text{cost}(o_j)$: # input tuples processed
- $\sigma_{p_1}(R1) \bowtie R2 \bowtie R3$
- Suppose $\sigma_{p_1}(R1) = 1M$ but DBMS underestimates as 1
- Suppose $|R2| = 1M$ and $|R3| = 1K$
- Suppose output of join has the size of the minimum



Outline For Today

1. Goal of Query Optimization and Overview of Techniques
2. Cost-based Optimization Principles & Cardinality Estimation
3. **Cost-based DP Logical Join Plan Optimizer**
4. Rule-based Optimizations/Transformations
5. Final Remarks on Query Optimization & Query Processing

Widely Adopted Join Order Optimizer

Recall:

1. Enumerate a logical plan space (*often enumerates all join orders*)

$$L_1, L_2, \dots, L_k$$

A widely used optimization algorithm is to use dynamic programming:

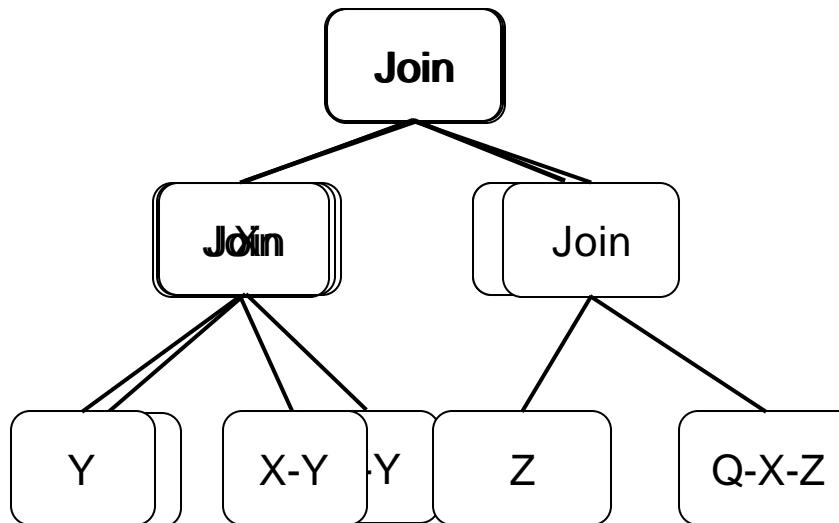
- Consider a join only query:

```
SELECT *
      FROM R1 NATURAL JOIN R2 NATURAL JOIN ... NATURAL JOIN Rn
```

- $Q = R1 \bowtie R2 \bowtie \dots \bowtie Rn$
- Note not-necessarily a ``chain'' query. It could be in any form, e.g:
 - $R1(A, B) \bowtie R2(B, C) \bowtie R3(C, A) \bowtie R4(A, B, C)$

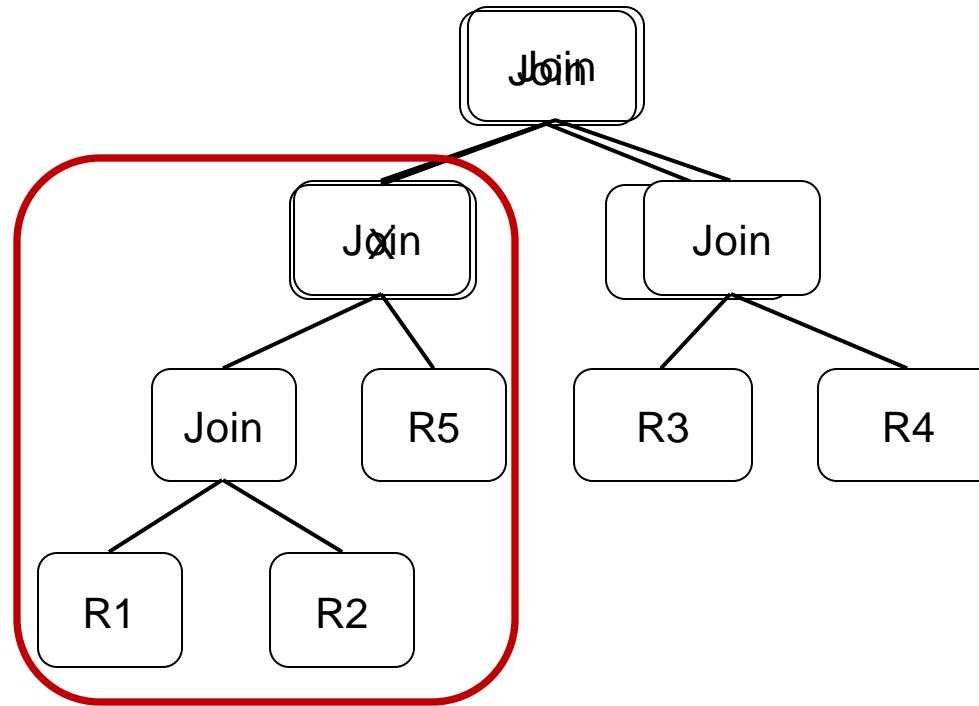
Plan Space

- In its most general form Plan Space=All possible join plan ``trees''
- In practice: If possible you'd avoid plans that do Cartesian Products
- Thought experiment: What does optimal tree L^* look like?



Optimal Sub-Join Tree Structure in L^*

- In L^* : What can we say about the sub-tree L^X starting from X ?
- Must be the best plan for the sub-query $Q^X = \bowtie_{\forall R_i \in X} R_i$
 - E.g: red-box must be the best plan for $R_1 \bowtie R_2 \bowtie R_5$ (o.w. just replace L^X with the best plan for Q^X : L^{X^*}).
- Therefore can use *dynamic programming algorithm to find join order.*



Cost-based DP Join Plan Optimizer

Input Q: $R_1 \bowtie R_2 \bowtie \dots \bowtie R_n$

Output Optimal Join Plan P:

OptPlans[]: a map that takes a sub-query Q_t and stores the already computed optimal plan:

```
for int t = 2 ... n // size of sub-queries
```

```
    for each  $Q_t \subseteq Q$  with t relations
```

```
         $P^*_{Q_t}$ : // best plan found so far
```

```
        for each ``split'' X,  $Q_t - X$ :
```

```
             $P^*_X = \text{OptPlans}[X]$ ;  $P^*_{Q_t-X} = \text{OptPlans}[Q_t-X]$ ;
```

```
             $P_{Q_t}: P^*_X \bowtie P^*_{Q_t-X}$ ; // Possible plan when split as X and  $Q_t-X$ 
```

```
             $P^*_{Q_t} = \min \text{ cost of } P^*_{Q_t}, P_{Q_t}$ 
```

```
        OptPlans[ $Q_t$ ] =  $P^*_{Q_t}$ 
```

where cardinality estimation of Q_t would happen

Optimization 1:

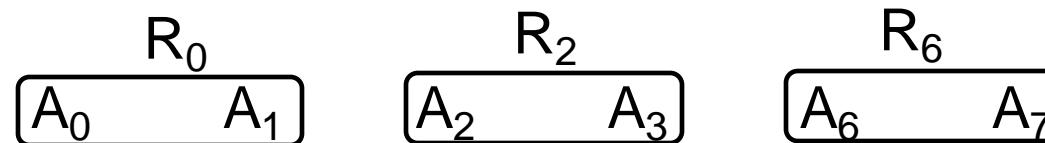
can enumerate over sub-queries that are ``connected'' to avoid Cartesian Products

Optimization 2:

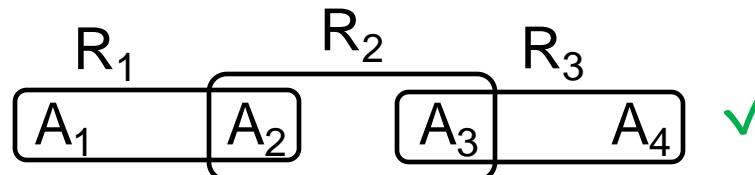
enumerate only if X and Q_t-X have common attributes; otherwise the possible plan would Cartesian product

Example Chain-based Join Optimizer (A4)

- A4: specialized version of DP Join Optimizer on ``chain queries``:
$$Q: R_0(A_0, A_1) \bowtie R_1(A_1, A_2) \bowtie \dots \bowtie R_{n-1}(A_{n-1}, A_n)$$
- Opt 1: Do not need to enumerate any dis-connected sub-query:
 - $Q_{t1}: R_0(A_0, A_1) \bowtie R_2(A_2, A_3) \bowtie R_6(A_6, A_7)$ **X No common attributes**



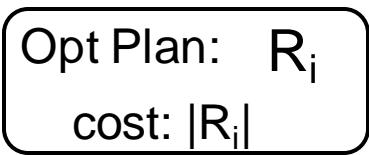
- $Q_{t2}: R_1(A_1, A_2) \bowtie R_2(A_2, A_3) \bowtie R_3(A_3, A_4)$



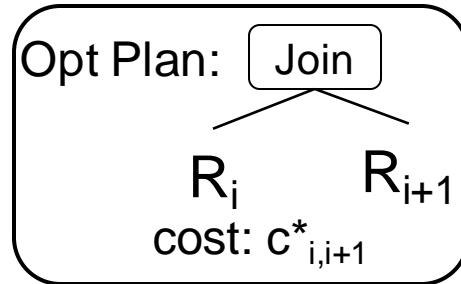
- Enumerate plans only for “consecutive”: $R_i \bowtie R_{i+1} \bowtie \dots \bowtie R_j$
- Enumerate only $j-i$ ``split points” for each $k: i \dots j-1$:
 - $R_i \bowtie R_{i+1} \bowtie \dots \bowtie R_k$ and $R_{k+1} \bowtie R_{k+2} \bowtie \dots \bowtie R_j$

Simulation

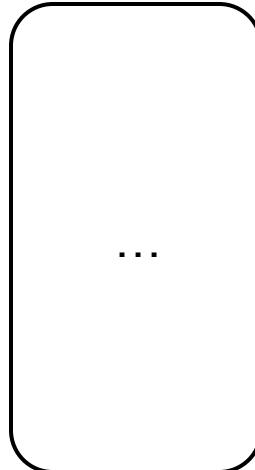
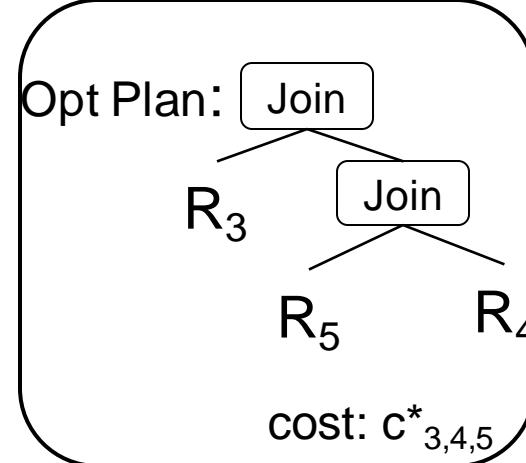
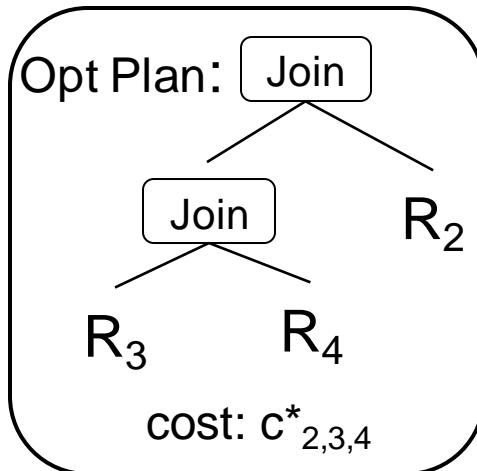
Opt Plans for 1-size
sub-queries R_i :



Opt Plans for 2-size
sub-queries $R_i \bowtie R_{i+1}$:



Opt Plans for 3-size
sub-queries
(using 1- and 2-size
opt. plans):



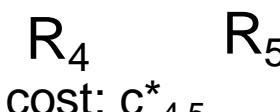
Simulation

When computing plans for a 4-size sub-query: e.g., $R_2 \bowtie R_3 \bowtie R_4 \bowtie R_5$:

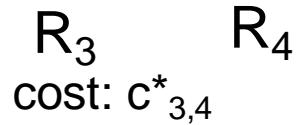
Opt Plan: 



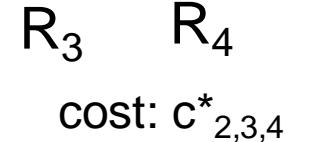
Opt Plan: 



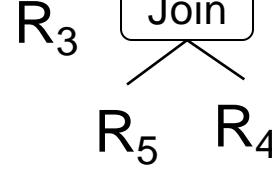
Opt Plan: 



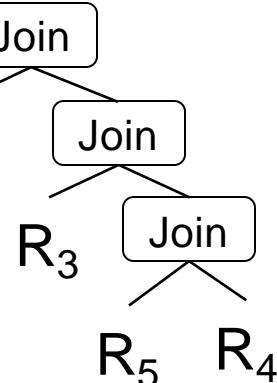
Opt Plan: 



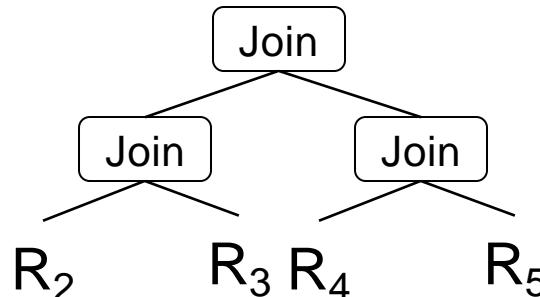
Opt Plan: 



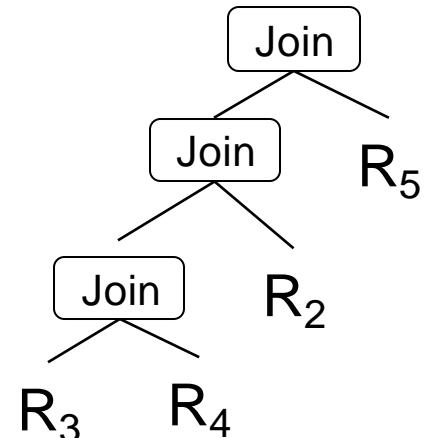
possible
plans:



R_2 as the split point



R_3 as the split point



R_4 as the split point

if left/right child matters compare 2x more plans

Outline For Today

1. Goal of Query Optimization and Overview of Techniques
2. Cost-based Optimization Principles & Cardinality Estimation
3. Cost-based DP Logical Join Plan Optimizer
4. Rule-based Optimizations/Transformations
5. Final Remarks on Query Optimization & Query Processing

Rule-based Transformations

- DP-based join optimizer algorithm only considered join-only queries
- What if there was a selection, projection, group-by aggregate etc?
- When possible we consider them as we enumerate plans but often in a *rule-based manner*

Example (1)

```
SELECT *  
FROM R1 NATURAL JOIN R2 NATURAL JOIN R3 NATURAL JOIN R4  
WHERE R1.A1 = "foo" AND R3.A3="bar"
```

- Intuitively instead of enumerating a plan for R1 we should enumerate a plan for relation: $\sigma_{A1=foo}(R1)$
- Similarly instead of R2, we should enumerate plans for $\sigma_{A3=bar}(R3)$
- Why?
- But not if the predicate was: R1.A1 = “foo” **OR** R3.A3=“bar”
- What to enumerate is governed by algebraic laws
 - *This is an important advantage of implementing a query language that's based on a formal algebra: i.e., relational algebra*

Example (2)

```
SELECT *  
FROM R1 NATURAL JOIN R2 NATURAL JOIN R3 NATURAL JOIN R4  
WHERE R1.A1 = "foo" AND R3.A3="bar"
```

- In relational algebra:

$$\sigma_{A1=foo \wedge A3=bar} (R1 \bowtie R2 \bowtie R3 \bowtie R4) = (\sigma_{A1=foo} (R1) \bowtie R2 \bowtie \sigma_{A3=bar} (R3) \bowtie R4)$$

- The expression effectively joins these smaller relations:

- i. $\sigma_{A1=foo} (R1)$

- ii. $R2$

- iii. $\sigma_{A3=bar} (R3)$

- iv. $R4$

- What if WHERE clause was $R1.A1 = "foo" \text{ OR } R3.A3="bar"$?
 - Apply the predicate only for sub-queries with both R1 and R3.
- The above algebraic law is called: *pushing down selections*

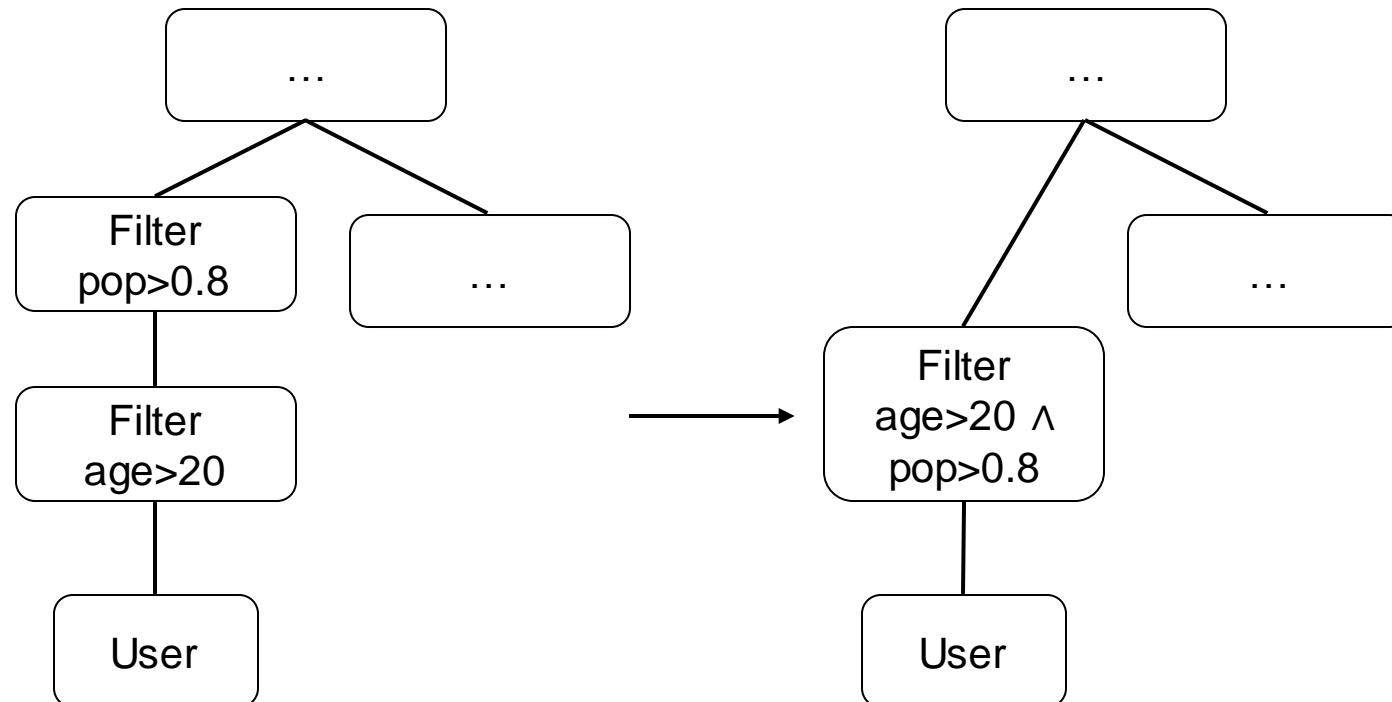
Ex Algebraic Transformation Rules (1)

Will use pure rel. algebra notation but can use our logical plan notation

- Convert σ_p - \times to/from \bowtie_p : $\sigma_p(R \times S) = R \bowtie_p S$
 - Example: $\sigma_{User.uid=Member.uid}(User \times Member) = User \bowtie Member$
- Merge/split σ 's: $\sigma_{p_1}(\sigma_{p_2}R) = \sigma_{p_1 \wedge p_2}R$
 - Example: $\sigma_{age > 20}(\sigma_{pop = 0.8}User) = \sigma_{age > 20 \wedge pop = 0.8}User$
- Merge/split π 's: $\pi_{L_1}(\pi_{L_2}R) = \pi_{L_1}R$, where $L_1 \subseteq L_2$
 - Example: $\pi_{age}(\pi_{age,pop}User) = \pi_{age}User$

Example In Logical Plan Notation

- Merge/split σ 's: $\sigma_{p_1}(\sigma_{p_2}R) = \sigma_{p_1 \wedge p_2}R$
- Example: $\sigma_{age > 20}(\sigma_{pop = 0.8}User) = \sigma_{age > 20 \wedge pop = 0.8}User$



Ex Algebraic Transformation Rules (2)

- Push down/pull up σ (not predicate is a conjunction):

$$\sigma_{p \wedge p_r \wedge p_s}(R \bowtie_{p'} S) = (\sigma_{p_r} R) \bowtie_{p \wedge p'} (\sigma_{p_s} S), \text{ where}$$

- p_r is a predicate involving only R columns
- p_s is a predicate involving only S columns
- p and p' are predicates involving both R and S columns
 - i.e., p an additional join predicate

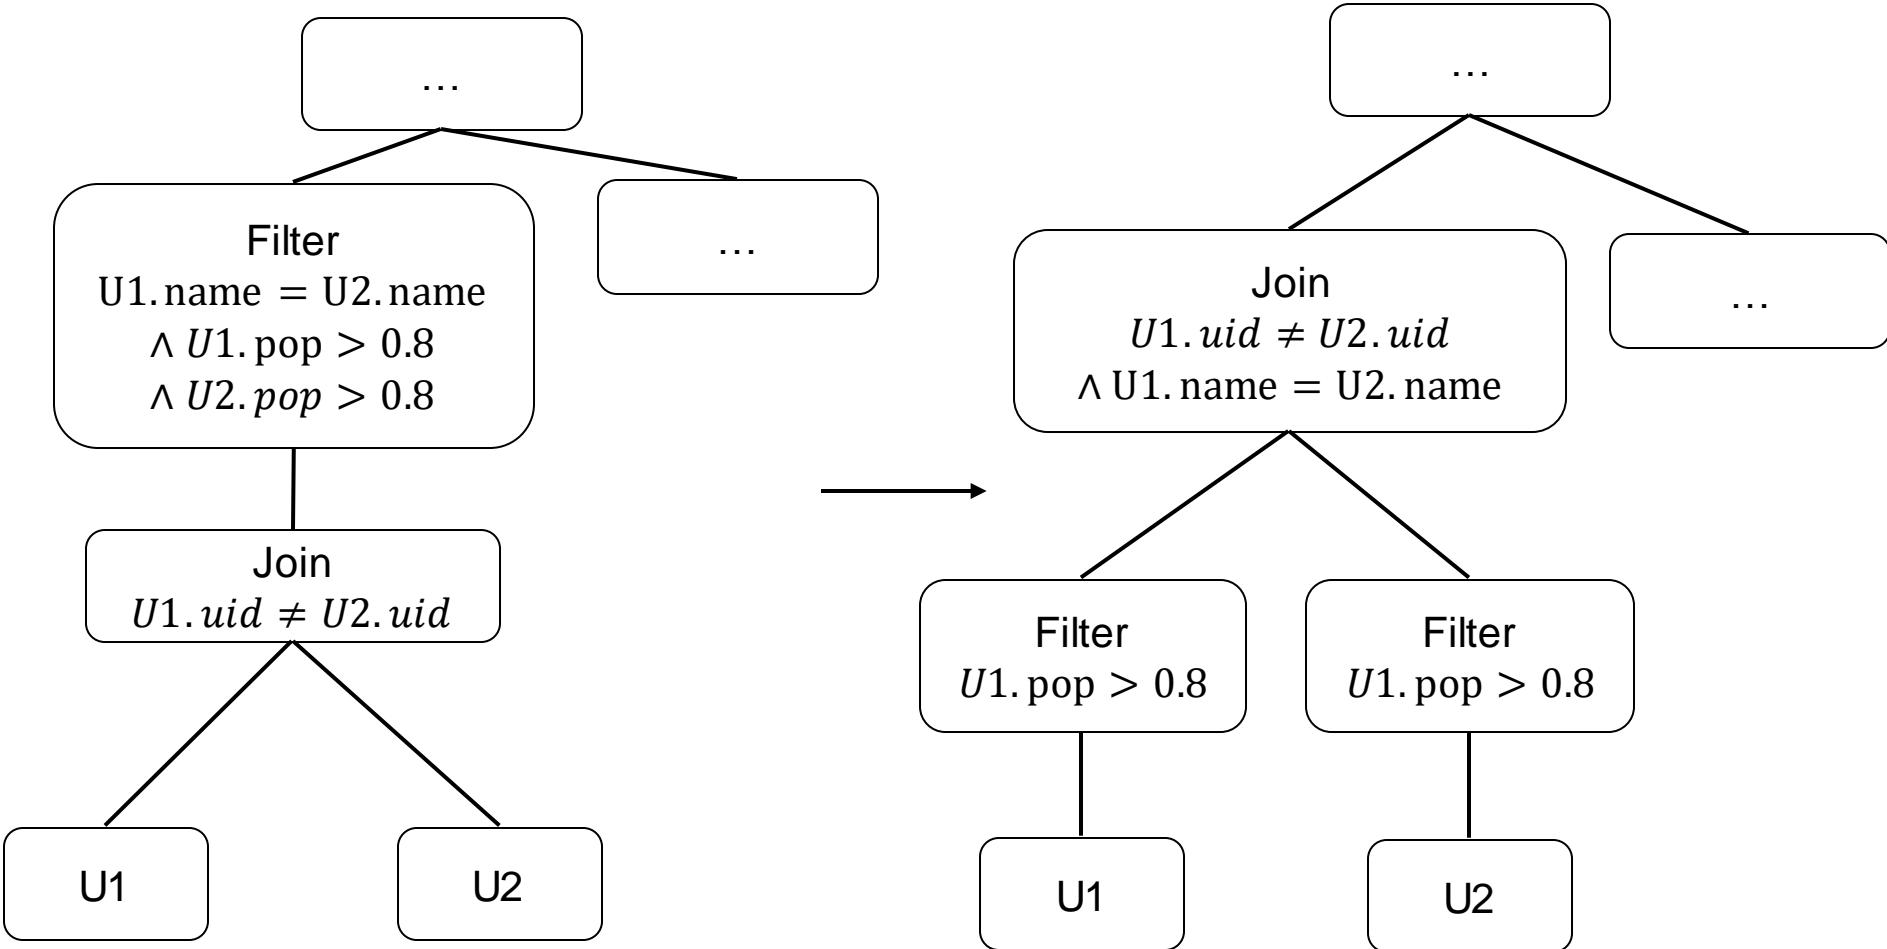
- Example:

$$\begin{aligned} & \sigma_{U1.name=U2.name \wedge U1.pop>0.8 \wedge U2.pop>0.8}(\rho_{U1}User \bowtie_{U1.uid \neq U2.uid} \rho_{U2}User) \\ &= \sigma_{pop>0.8}(\rho_{U1}User) \bowtie_{U1.uid \neq U2.uid, U1.name=U2.name} (\sigma_{pop>0.8}(\rho_{U2}User)) \end{aligned}$$

- Why should you always do this optimization?
 - Selections are relatively cheap (e.g., compared to joins or group-by and aggregates) and can only reduce the number tuples processed.

Example In Logical Plan Notation

$$\begin{aligned} & \sigma_{U1.name=U2.name \wedge U1.pop>0.8 \wedge U2.pop>0.8} (\rho_{U1} User \bowtie_{U1.uid \neq U2.uid} \rho_{U2} User) \\ & = \sigma_{pop>0.8} (\rho_{U1} User) \bowtie_{U1.uid \neq U2.uid, U1.name=U2.name} (\sigma_{pop>0.8} (\rho_{U2} User)) \end{aligned}$$



Ex Algebraic Transformation Rules (3)

- Push down π : $\pi_L(\sigma_p R) = \pi_L(\sigma_p(\pi_{L L'} R))$, where
 - L' is the set of columns referenced by p that are not in L
 - Example:

$$\pi_{age}(\sigma_{pop > 0.8} User) = \pi_{age}(\sigma_{pop > 0.8}(\pi_{age, pop} User))$$

- Not as important and effective as pushing σ
- Many more (seemingly trivial) equivalences...
- Can be systematically used to transform plans

Outline For Today

1. Goal of Query Optimization and Overview of Techniques
2. Cost-based Optimization Principles & Cardinality Estimation
3. Cost-based DP Logical Join Plan Optimizer
4. Rule-based Optimizations/Transformations
5. Final Remarks on Query Optimization & Query Processing

Final Remarks (1)

- Query Optimizer and Cardinality Estimator: Brain of the DBMS
 - *Ultimate Goal: Pick a reasonable plan (i.e., one processing few tuples)*
- Query Processor and Storage: Skeleton
 - They do actual data searching and computation
- Several insights have emerged over the years in DBMS literature:
 - Cost model is not very critical: keep a simple model (e.g., # tuples)
 - Cardinality estimation: matters a lot
 - But! Extremely difficult to integrate a good estimator. Always a hack with wild unrealistic assumptions here and there to make it implementable: magic constants, uniformity assumptions, independence assumptions etc.
- My advice: Optimizer is important but keep it simple.
 - Do not be complacent on the query processor and storage! Work very hard on these and optimize relentlessly!

Final Remarks (2)

- CS 448: Database Systems Implementation
 - Gets into many more details about the internals of query processing and optimization and other DBMS components!
 - A4's programming question is meant to give you a glimpse of CS 448 assignments.