

CS 348 Lectures 15-16

Query Optimization

Semih Salihoglu

March 8-10 2022




UNIVERSITY OF
WATERLOO | **DSg** Data
Systems
Group


Recall: Overview of Compilation Steps

Text


SELECT ...
FROM ...
WHERE...



Logical Plan



Physical Plan



Outline For Today

1. Goal of Query Optimization and Overview of Techniques
2. Cost-based Optimization Principles
3. Cost-based DP Logical Join Plan Optimizer
4. Cardinality Estimation Techniques
5. Rule-based Optimizations/Transformations
6. Final Remarks on Query Optimization & Query Processing

Outline For Today


1. Goal of Query Optimization and Overview of Techniques
2. Cost-based Optimization Principles
3. Cost-based DP Logical Join Plan Optimizer
4. Cardinality Estimation Techniques
5. Rule-based Optimizations/Transformations
6. Final Remarks on Query Optimization & Query Processing

Goal of Query Optimization (1)

- Recall ultimately a *physical plan* executes to answer a query
- Given a query Q, many equivalent physical plans exist:
 1. Many equivalent logical plans exist

```
SELECT cid  
FROM Customer C, Order O, Product P  
WHERE C.cid = O.cid AND O.pid = P.pid  
      AND P.name = BookA
```

Logical Plans:




Goal of Query Optimization (1)

- Recall ultimately a *physical plan* executes to answer a query
- Given a query Q, many equivalent physical plans exist:
 1. Many equivalent logical plans exist
 2. Each logical plan can have many equivalent physical plans.


```
SELECT cid  
FROM Customer C, Order O, Product P  
WHERE C.cid = O.cid AND O.pid = P.pid  
      AND P.name = BookA
```

Physical Plans:



Goal of Query Optimization (2)

- Ultimately: Given Q , pick the “best” physical plan for Q :
 - Best: often means fastest, could mean “cheapest”
- DBMS developers are more humble:
 - Pick a reasonably good plan. Do not pick a very bad plan!
 - Example plan spectrum of join-heavy queries



Overview of Query Opt. Techniques

1. Enumerate a logical plan space (often

enumerates all join orders)

(extended) relational
algebraic expressions $\sqsubset L_1, L_2, \dots, L_k$

2. For one or more of L_i , (optionally)

enumerate a physical plan space:

$$P_{i,1}, P_{i,2}, \dots, P_{i,t}$$

3. Pick the best $P_{i,1}$

A common approach:

- Step 1 is cost-based or hybrid
- Step 2 is rule-based

Options for steps 1 & 2:

- i. Rule-based
- ii. Cost-based
- iii. Hybrid rule/cost-based

Outline For Today


1. Goal of Query Optimization and Overview of Techniques
2. Cost-based Optimization Principles
3. Cost-based DP Logical Join Plan Optimizer
4. Cardinality Estimation Techniques
5. Rule-based Optimizations/Transformations
6. Final Remarks on Query Optimization & Query Processing

Cost-based Optimization Principles

- System R ('70s): First prototype relational DBMS from IBM




Patricia Selinger



- Give each enumerated log/phy plan, e.g., L_i , a $\text{Cost}(L_i) = c_i$
- Cost is the estimate of the system for how good/bad L_i is.
- Pick min cost plan

Cost-based Optimization Principles (1)


- Naturally: cost definition is broken into costs of operators.
i.e: $\text{Cost}(L_i) = \sum_j \text{cost}(o_j \in L_i)$
- Example cost metrics or components:
 - # I/Os a plan will make
 - # tuples that will pass through operators
 - # runtime of algorithm o_j is running
 - e.g., nested loop join of R, S: $|R|^*|S|$
 - Combination of above
- For any reasonable metric:
 - Need to estimate cardinality, i.e., size, of tuples o_j will process



The (notorious) cardinality estimation problem!

2 Components of Cost-based Optimization

1. What is the cost metric?
 - Can be complicated, e.g., different ops could have different costs
 - But inevitably depends on cardinality, i.e., number, of tuples processed by each operator
2. How do we estimate cardinality of tables processed by each op?
 - Need a “cardinality estimation technique to estimate cardinality



Cardinality Estimation


- Given a database

1. $D: R_1(A_{1,1}, \dots, A_{1,m_1}), \dots, R_n(A_{n,1}, \dots, A_{n,m_n})$
2. A (sub-) query Q (a relational algebra expression)

What is the $|Q|$?


- E.g:

- $\sigma_{name=BookA}(Product)$?
- $Product \bowtie Order$?
- $\sigma_{name=BookA}(Product \bowtie Order \bowtie Customer)$?




Example Poor Optimizer Choice

- Suppose $\text{cost}(o_j)$: # input tuples processed
- $\sigma_{p1}(R1) \bowtie R2 \bowtie R3$
- Suppose $\sigma_{p1}(R1) = 1M$ but DBMS underestimates as 1
- Suppose $|R2| = 1M$ and $|R3| = 1K$
- Suppose output of join has the size of the minimum input relation



Example Poor Optimizer Choice

- Suppose $\text{cost}(o_j)$: # input tuples processed
- $\sigma_{p1}(R1) \bowtie R2 \bowtie R3$
- Suppose $\sigma_{p1}(R1) = 1M$ but DBMS underestimates as 1
- Suppose $|R2| = 1M$ and $|R3| = 1K$
- Suppose output of join has the size of the minimum



Outline For Today

1. Goal of Query Optimization and Overview of Techniques
2. Cost-based Optimization Principles
3. **Cost-based DP Logical Join Plan Optimizer**
4. Cardinality Estimation Techniques
5. Rule-based Optimizations/Transformations
6. Final Remarks on Query Optimization & Query Processing

Widely Adopted Join Order Optimizer

Recall:

1. Enumerate a logical plan space (*often enumerates all join orders*)

$$L_1, L_2, \dots, L_k$$

A widely used optimization algorithm is to use dynamic programming:

- Consider a join only query:

```
SELECT *
  FROM R1 NATURAL JOIN R2 NATURAL JOIN ... NATURAL JOIN Rn
```


- $Q = R1 \bowtie R2 \bowtie \dots \bowtie Rn$

- Note not-necessarily a ``chain'' query. It could be in any form, e.g:

- $R1(A, B) \bowtie R2(B, C) \bowtie R3(C, A) \bowtie R4(A, B, C)$


Plan Space

- In its most general form Plan Space=All possible join plan ``trees''
- In practice: If possible you'd avoid plans that do Cartesian Products
- Thought experiment: What does optimal tree L^* look like?



Optimal Sub-Join Tree Structure in L^*

- In L^* : What can we say about the sub-tree L^X starting from X ?
- Must be the best plan for the sub-query $Q^X = \bowtie_{\forall R_i \in X} R_i$
 - E.g: red-box must be the best plan for $R_1 \bowtie R_2 \bowtie R_5$ (o.w. just replace L^X with the best plan for Q^X : L^{X^*} .)
- Therefore can use *dynamic programming algorithm to find join order.*



Cost-based DP Join Plan Optimizer

Input Q: $R_1 \bowtie R_2 \bowtie \dots \bowtie R_n$

Output Optimal Join Plan P:

OptPlans[]: a map that takes a sub-query Q_t and stores the already computed optimal plan:

```
for int t = 2 ... n // size of sub-queries
```

```
    for each  $Q_t \subseteq Q$  with t relations
```

```
         $P^*_{Q_t}$ : // best plan found so far
```

```
        for each ``split'' X,  $Q_t - X$ :
```

```
             $P^*_X = \text{OptPlans}[X]$ ;  $P^*_{Q_t-X} = \text{OptPlans}[Q_t-X]$ ;
```

```
             $P_{Q_t}: P^*_X \bowtie P^*_{Q_t-X}$ ; // Possible plan when split as X and  $Q_t-X$ 
```

```
             $P^*_{Q_t} = \min \text{ cost of } P^*_{Q_t}, P_{Q_t}$ 
```

```
        OptPlans[ $Q_t$ ] =  $P^*_{Q_t}$ 
```

where cardinality estimation of Q_t would happen

Optimization 1:


can enumerate over sub-queries that are ``connected'' to avoid Cartesian Products

Optimization 2:
enumerate only if


X and Q_t-X have common attributes; otherwise the possible plan would Cartesian product

Example Chain-based Join Optimizer (A4)

- A4: specialized version of DP Join Optimizer on ``chain queries``:
$$Q: R_0(A_0, A_1) \bowtie R_1(A_1, A_2) \bowtie \dots \bowtie R_{n-1}(A_{n-1}, A_n)$$
- Opt 1: Do not need to enumerate any dis-connected sub-query:
 - $Q_{t1}: R_0(A_0, A_1) \bowtie R_2(A_2, A_3) \bowtie R_6(A_6, A_7)$ **X No common attributes**




- $Q_{t2}: R_1(A_1, A_2) \bowtie R_2(A_2, A_3) \bowtie R_3(A_3, A_4)$




- Enumerate plans only for “consecutive”: $R_i \bowtie R_{i+1} \bowtie \dots \bowtie R_j$
- Enumerate only $j-i$ ``split points'' for each $k: i \dots j-1$:
 - $R_i \bowtie R_{i+1} \bowtie \dots \bowtie R_k$ and $R_{k+1} \bowtie R_{k+2} \bowtie \dots \bowtie R_j$

Simulation




Opt Plans for 1-size
sub-queries R_i :



Opt Plans for 2-size
sub-queries $R_i \bowtie R_{i+1}$:



Opt Plans for 3-size
sub-queries
(using 1- and 2-size
opt. plans):







Simulation

When computing plans for a 4-size sub-query: e.g., $R_2 \bowtie R_3 \bowtie R_4 \bowtie R_5$:


Opt Plan: Join




Opt Plan: Join




possible plans:



R_2 as the split point



R_3 as the split point



R_4 as the split point


if left/right child matters compare 2x more plans

Outline For Today

1. Goal of Query Optimization and Overview of Techniques
2. Cost-based Optimization Principles
3. Cost-based DP Logical Join Plan Optimizer
- 4. Cardinality Estimation Techniques**
5. Rule-based Optimizations/Transformations
6. Final Remarks on Query Optimization & Query Processing


Recall Example Poor Optimizer Choice

- Suppose $\text{cost}(o_j)$: # input tuples processed
- $\sigma_{p1}(R1) \bowtie R2 \bowtie R3$
- Suppose $\sigma_{p1}(R1) = 1M$ but DBMS underestimates as 1
- Suppose $|R2| = 1M$ and $|R3| = 1K$
- Suppose output of join has the size of the minimum input relation




Recall Example Poor Optimizer Choice

- Suppose $\text{cost}(o_j)$: # input tuples processed
- $\sigma_{p1}(R1) \bowtie R2 \bowtie R3$
- Suppose $\sigma_{p1}(R1) = 1M$ but DBMS underestimates as 1
- Suppose $|R2| = 1M$ and $|R3| = 1K$
- Suppose output of join has the size of the minimum



2 High-level Card. Estimation Techniques

1. Sampling-based:
 - While optimizing Q, sample relations to make an estimate
2. Summary/statistics-based:
 - Use statistics about D to make estimates
 - Possible statistics:
 - $|R_i|$: size of each relation
 - $|\pi_{A_j}(R_i)|$ # distinct values in column A_j
 - Histograms: Distribution of values on A_j
 - Also use known constraints:
 - E.g: FK constraint from R to S: $|R \bowtie S| = |R|$
 - 2 common *simplification* assumptions (no other good reason):
 - (i) uniformity; (ii) independence



Example Statistics-based Estimation Techniques

Selections with Equality Predicates


- $Q: \sigma_{A=v} R$
 - Suppose the following information is available:
 - Size of R : $|R|$
 - Number of distinct A values in R : $|\pi_A R|$
 - Assumptions:
 1. Values of A are *uniformly distributed* in R
 2. $v \in |\pi_A R|$
 - $|Q| \approx |R| / |\pi_A R|$
 - Selectivity factor of $(A = v)$ is $1 / |\pi_A R|$
 - Ex: $|Product| = 1000$, $|\pi_{name}(Product)| = 50$
 - $\sigma_{name=BookA} Product: 1000/50 = 20$
- wild assumption, often doesn't hold*
- fair assumption, often holds (b/c users search things they put in the db)*

Conjunctive Predicates

- $Q: \sigma_{A=u \wedge B=v} R$
- Additional assumption:
 3. $(A = u)$ and $(B = v)$ are *independent*
 - Counter example: age and salary
- $|Q| \approx |R| / |\pi_A R| \cdot |\pi_B R|$
 - Reduce total size by all selectivity factors
 - Directly derived from standard probability rules:
 - $\Pr(E_1) = p_1$, and $\Pr(E_2) = p_2$ and E_1 and E_2 are independent:
 - $\Pr(E_1 \wedge E_2) = p_1 * p_2$
 - Ex: $\Pr(\text{heads} \wedge \text{dice}=6) = 1/2 * 1/6 = 1/12$
- Ex: $|Prod| = 1000$, $|\pi_{name}(Prod)| = 50$, $|\pi_{merchant}(Prod)| = 4$
 - $\sigma_{name=BookA \wedge merchant=B\&N} Product: 1000/(50*4) = 5$

Negated and Disjunctive Predicates

- $Q: \sigma_{A \neq v} R$
 - $|Q| \approx |R| \cdot (1 - 1/|\pi_A R|)$
 - Selectivity factor of $\neg p$ is $(1 - \text{selectivity factor of } p)$
- $Q: \sigma_{A=u \vee B=v} R$
 - $|Q| \approx |R| \cdot (1/|\pi_A R| + 1/|\pi_B R|) ?$
 - No! Tuples satisfying $(A = u)$ and $(B = v)$ are counted twice
 - Use only for $\sigma_{A=u \vee A=v} R$ (b/c then $A=u$ and $A=v$ are disjoint)
 - $|Q| \approx |R| \cdot (1/|\pi_A R| + 1/|\pi_B R| - 1/|\pi_A R||\pi_B R|)$
 - Inclusion-exclusion principle from probability




Range Predicates

- $\sigma_{A>u} R?$
- Case 1: Suppose the DBMS knew actual projection values:
 - Then range queries are a generalization of $\sigma_{A=u \vee A=v} R$
 - $\sigma_{A>u} R = |Q| \approx |R| \cdot \left(\frac{|\#vals>u|}{|\pi_A R|} \right) ?$
 - E.g: A was an int column and $|\pi_A R| = \{1, 2, 3, 4, 5\}$
 - $\sigma_{A>2} R = |R| * 3/5$

Case 2 of Range Predicates

- Case 2: We don't know actual values
- Not enough information!
 - Just pick a *magic constant*, e.g., $|Q| \approx |R| \cdot \frac{1}{3}$
- With more information
 - Largest $R.A$ value: $\text{high}(R.A)$
 - Smallest $R.A$ value: $\text{low}(R.A)$
 - $|Q| \approx |R| \cdot \frac{\text{high}(R.A) - v}{\text{high}(R.A) - \text{low}(R.A)}$
 - In practice: sometimes the second highest and lowest are used
 - The highest and the lowest are often used by inexperienced database designer to represent invalid values!



Equi-Join of Two Relations (1)

- $Q: R(A, B) \bowtie S(A, C)$
- Assumption: containment of value sets:
 - Every tuple in the “smaller” relation (one with fewer distinct values for the join attribute) joins with a tuple in the other relation
 - That is, if $|\pi_A R| \leq |\pi_A S|$ then $\pi_A R \subseteq \pi_A S$
 - Certainly not true in general
 - But holds in the common case of foreign key joins
- $|Q| \approx \frac{|R| \cdot |S|}{\max(|\pi_A R|, |\pi_A S|)}$
 - Selectivity factor of $R.A = S.A$ is $1 / \max(|\pi_A R|, |\pi_A S|)$

Equi-Join of Two Relations (2)

- Example:

R	
A	B
a ₁	b ₁
a ₁	b ₂
a ₁	b ₃
a ₁	b ₄

$$\pi_A R = \{a_1\}$$

S	
A	C
a ₁	c ₁
a ₁	c ₂
a ₂	c ₃
a ₂	c ₄

$$\pi_A S = \{a_1, a_2\}$$

- $|Q| \approx \frac{|R| \cdot |S|}{\max(|\pi_A R|, |\pi_A S|)} = 4 \times 4 / 2 = 8$ (correct)
- If we had picked $\min(|\pi_A R|, |\pi_A S|)$, then we'd over-estimate
 - Intuitively a fraction of tuples from the larger-domain table will join with each tuple from smaller-domain table (not vice versa)

Other Estimations Techniques

- Using similar ideas, we can estimate the size of projection, duplicate elimination, union, difference, aggregation (with grouping)
- Lots of assumptions and very rough estimation
 - Accurate estimate is not needed
 - Maybe okay if we overestimate or underestimate consistently
- In practice: Very very difficult but very important for the optimizer.
 - B/c: ultimate goal is to help estimate costs of operators & plans
 - If we badly underestimate an expression => may lead to bad plans

Outline For Today

1. Goal of Query Optimization and Overview of Techniques
2. Cost-based Optimization Principles
3. Cost-based DP Logical Join Plan Optimizer
4. Cardinality Estimation Techniques
5. Rule-based Optimizations/Transformations
6. Final Remarks on Query Optimization & Query Processing

Rule-based Transformations

- DP-based join optimizer algorithm only considered join-only queries
- What if there was a selection, projection, group-by aggregate etc?
- When possible we consider them as we enumerate plans but often in a *rule-based manner*

Example (1)

```
SELECT *  
FROM R1 NATURAL JOIN R2 NATURAL JOIN R3 NATURAL JOIN R4  
WHERE R1.A1 = "foo" AND R3.A3="bar"
```

- Intuitively instead of enumerating a plan for R1 we should enumerate a plan for relation: $\sigma_{A1=foo}(R1)$
- Similarly instead of R2, we should enumerate plans for $\sigma_{A3=bar}(R3)$
- Why?
- But not if the predicate was: $R1.A1 = "foo"$ **OR** $R3.A3="bar"$
- What to enumerate is governed by algebraic laws
 - *This is an important advantage of implementing a query language that's based on a formal algebra: i.e., relational algebra*

Example (2)

```
SELECT *  
FROM R1 NATURAL JOIN R2 NATURAL JOIN R3 NATURAL JOIN R4  
WHERE R1.A1 = "foo" AND R3.A3="bar"
```

- In relational algebra:

$$\sigma_{A1=foo \wedge A3=bar} (R1 \bowtie R2 \bowtie R3 \bowtie R4) = (\sigma_{A1=foo} (R1) \bowtie R2 \bowtie \sigma_{A3=bar} (R3) \bowtie R4)$$

- The expression effectively joins these smaller relations:

- i. $\sigma_{A1=foo} (R1)$

- ii. $R2$

- iii. $\sigma_{A3=bar} (R3)$

- iv. $R4$

- What if WHERE clause was $R1.A1 = "foo" \text{ OR } R3.A3="bar"$?
 - Apply the predicate only for sub-queries with both R1 and R3.
- The above algebraic law is called: *pushing down selections*


Ex Algebraic Transformation Rules (1)

Will use pure rel. algebra notation but can use our logical plan notation

- Convert σ_p - \times to/from \bowtie_p : $\sigma_p(R \times S) = R \bowtie_p S$
 - Example: $\sigma_{User.uid=Member.uid}(User \times Member) = User \bowtie Member$
- Merge/split σ 's: $\sigma_{p_1}(\sigma_{p_2}R) = \sigma_{p_1 \wedge p_2}R$
 - Example: $\sigma_{age > 20}(\sigma_{pop = 0.8}User) = \sigma_{age > 20 \wedge pop = 0.8}User$
- Merge/split π 's: $\pi_{L_1}(\pi_{L_2}R) = \pi_{L_1}R$, where $L_1 \subseteq L_2$
 - Example: $\pi_{age}(\pi_{age,pop}User) = \pi_{age}User$

Example In Logical Plan Notation

- Merge/split σ 's: $\sigma_{p_1}(\sigma_{p_2}R) = \sigma_{p_1 \wedge p_2}R$
- Example: $\sigma_{age > 20}(\sigma_{pop = 0.8}User) = \sigma_{age > 20 \wedge pop = 0.8}User$



Ex Algebraic Transformation Rules (2)

- Push down/pull up σ (not predicate is a conjunction):

$$\sigma_{p \wedge p_r \wedge p_s}(R \bowtie_{p'} S) = (\sigma_{p_r} R) \bowtie_{p \wedge p'} (\sigma_{p_s} S), \text{ where}$$

- p_r is a predicate involving only R columns
- p_s is a predicate involving only S columns
- p and p' are predicates involving both R and S columns
 - i.e., p an additional join predicate

- Example:


$$\begin{aligned} & \sigma_{U1.name=U2.name \wedge U1.pop>0.8 \wedge U2.pop>0.8}(\rho_{U1}User \bowtie_{U1.uid \neq U2.uid} \rho_{U2}User) \\ &= \sigma_{pop>0.8}(\rho_{U1}User) \bowtie_{U1.uid \neq U2.uid, U1.name=U2.name} (\sigma_{pop>0.8}(\rho_{U2}User)) \end{aligned}$$

- Why should you always do this optimization?

- Selections are relatively cheap (e.g., compared to joins or group-by and aggregates) and can only reduce the number tuples processed.

Example In Logical Plan Notation

$$\begin{aligned} & \sigma_{U1.name=U2.name \wedge U1.pop>0.8 \wedge U2.pop>0.8} (\rho_{U1} User \bowtie_{U1.uid \neq U2.uid} \rho_{U2} User) \\ & = \sigma_{pop>0.8} (\rho_{U1} User) \bowtie_{U1.uid \neq U2.uid, U1.name=U2.name} (\sigma_{pop>0.8} (\rho_{U2} User)) \end{aligned}$$



Ex Algebraic Transformation Rules (3)

- Push down π : $\pi_L(\sigma_p R) = \pi_L(\sigma_p(\pi_{L'L'}R))$, where
 - L' is the set of columns referenced by p that are not in L
 - Example:

$$\pi_{age}(\sigma_{pop>0.8} User) = \pi_{age}(\sigma_{pop>0.8}(\pi_{age,pop} User))$$

- Not as important and effective as pushing σ
- Many more (seemingly trivial) equivalences...
- Can be systematically used to transform plans

Outline For Today

1. Goal of Query Optimization and Overview of Techniques
2. Cost-based Optimization Principles
3. Cost-based DP Logical Join Plan Optimizer
4. Cardinality Estimation Techniques
5. Rule-based Optimizations/Transformations
6. Final Remarks on Query Optimization & Query Processing

Final Remarks (1)

- Query Optimizer and Cardinality Estimator: Brain of the DBMS
 - *Ultimate Goal: Pick a reasonable plan (i.e., one processing few tuples)*
- Query Processor and Storage: Skeleton
 - They do actual data searching and computation
- Several insights have emerged over the years in DBMS literature:
 - Cost model is not very critical: keep a simple model (e.g., # tuples)
 - Cardinality estimation: matters a lot
 - But! Extremely difficult to integrate a good estimator. Always a hack with wild unrealistic assumptions here and there to make it implementable: magic constants, uniformity assumptions, independence assumptions etc.
- My advice: Optimizer is important but keep it simple.
 - Do not be complacent on the query processor and storage! Work very hard on these and optimize relentlessly!

Final Remarks (2)

- CS 448: Database Systems Implementation
 - Gets into many more details about the internals of query processing and optimization and other DBMS components!
 - A4's programming question is meant to give you a glimpse of CS 448 assignments.