

Concurrency control & recovery system

Transaction Processing

Introduction to Database Management

CS348 Spring 2021

Review


- ACID
 - **Atomicity**: TX's are either completely done or not done at all
 - **Consistency**: TX's should leave the database in a consistent state
 - **Isolation**: TX's must behave as if they are executed in isolation
 - **Durability**: Effects of committed TX's are resilient against failures
- SQL transactions
 - Begins implicitly
 - SELECT ...;
 - UPDATE ...;
 - ROLLBACK | COMMIT;

Outline

- Concurrency control -- isolation
 - Review serializable execution histories
 - Locking-based concurrency control
- Recovery – atomicity and durability
 - Naïve approaches
 - Logging for undo and redo

Concurrency control

- Goal: ensure the “I” (isolation) in ACID





Good versus bad execution histories

Serial					
Good!		Bad!		Good! Why?	
T_1	T_2	T_1	T_2	T_1	T_2
r1(x)		r1(x)		r1(x)	
w1(x)		Read 400	r2(x)	w1(x)	r2(x)
r1(y)		Write 400 – 100	w1(x)	Read 400	w2(x)
w1(y)			w2(x)	Write 400 – 50	r1(y)
	r2(x)	r1(y)	r2(z)		r2(C)
	w2(x)		w1(y)		w1(y)
	r2(z)		w2(z)		
H_a	w2(z)	H_b	w2(z)	H_c	w2(C)

Good versus bad execution histories

Serialization graph
(Lecture 15)





Good versus bad execution histories

How to avoid
this?

Not serializable

Bad!



Concurrency control

Possible classification

- Pessimistic – assume that conflicts will happen and take preventive action
 - Two-phase locking (2PL)
 - Timestamp ordering
- Optimistic – assume that conflicts are rare and run transactions and fix if there is a problem
- We will only review 2PL

Locking


- Rules
 - If a transaction wants to **read** an object, it must first request a **shared lock (S mode)** on that object
 - If a transaction wants to **modify** an object, it must first request an **exclusive lock (X mode)** on that object
 - Allow one exclusive lock, or multiple shared locks

Mode of the lock requested		
	S	X
Mode of lock(s) currently held by other transactions		
S	Yes	No
X	No	No


Compatibility matrix

Grant the lock?

Basic locking is not enough




Basic locking is not enough



Two-phase locking (2PL)

- All lock requests precede all unlock requests
 - Phase 1: obtain locks, phase 2: release locks



Remaining problems of 2PL

T_1	T_2
$r_1(x)$	
$w_1(x)$	
	$r_2(x)$
	$w_2(x)$
$r_1(y)$	
$w_1(y)$	
	$r_2(y)$
	$w_2(y)$
Abort!	

- T_2 has read uncommitted data written by T_1
- If T_1 aborts, then T_2 must abort as well
- **Cascading aborts** possible if other transactions have read data written by T_2

- Even worse, what if T_2 commits before T_1 ?
 - Schedule is **not recoverable** if the system crashes right after T_2 commits


Deadlocks

- A transaction is deadlocked if it is blocked and will remain blocked until there is an intervention.
- Locking-based concurrency control algorithms may cause deadlocks requiring abort of one of the transactions

- Consider the partial history

- Neither T_1 nor T_2 can make progress

Cannot obtain
the lock on y
until T_2 unlocks



Strict 2PL

- Only release X-locks at commit/abort time
 - A writer will block all other readers until the writer commits or aborts
- Used in many commercial DBMS
 - Oracle is a notable exception
- Why do we use strict 2PL? (assignment question)




Outline

- Concurrency control -- isolation
 - Review serializable execution histories
 - Locking-based concurrency control
- Recovery – atomicity and durability
 - Naïve approaches
 - Logging for undo and redo

Execution model


To read/write X

- The disk block containing X must be first brought into memory
- X is read/written in memory
- The memory block containing X, if modified, must be written back (flushed) to disk eventually



Failures

- System crashes right after a transaction T_1 commits;
but not all effects of T_1 were written to disk
 - How do we complete/redo T_1 (**durability**)?
- System crashes in the middle of a transaction T_2 ;
partial effects of T_2 were written to disk
 - How do we undo T_2 (**atomicity**)?



Naïve approach: Force -- durability

T_1 (balance transfer of \$100 from A to B)

```
read(A, a); a = a - 100;  
write(A, a);  
read(B, b); b = b + 100;  
write(B, b);  
commit;
```

Memory buffer

~~A = 800 700~~
~~B = 400 500~~

Disk

~~A = 800 700~~
~~B = 400 500~~

Force: all writes must be reflected on disk
when a transaction commits



Naïve approach: Force -- durability

T1 (balance transfer of \$100 from A to B)

```
read(A, a); a = a - 100;  
write(A, a);  
read(B, b); b = b + 100;  
write(B, b);  
commit;
```



Force: all writes must be reflected on disk
when a transaction commits



Without force: not all writes are on disk when T1 commits

If system crashes right after T1 commits, effects of T1 will be lost

Naïve approach: No steal -- atomicity

T1 (balance transfer of \$100 from A to B)

```
read(A, a); a = a - 100;
```

```
write(A, a);
```

```
read(B, b); b = b + 100;
```

```
-----  
write(B, b);
```

```
commit;
```



Memory buffer

A = 800	700
B = 400	500

Disk

A = 800	700
B = 400	

Bad!

With steal: some writes are on disk before T commits


If system crashes before T1 commits, there is no way to undo the changes

Naïve approach

- **Force:** When a transaction commits, all writes of this transaction must be reflected on disk
 - Ensures durability
 - ☞ Problem of force: Lots of **random writes** hurt performance
- **No steal:** Writes of a transaction can only be flushed to disk at commit time
 - Ensures atomicity
 - ☞ Problem of no steal: Holding on to all dirty blocks requires lots of memory

Logging


- **Database log:** sequence of **log records**, recording all changes made to the database, written to stable storage (e.g., disk) during normal operation



- Hey, one change turns into two—bad for performance?
 - But writes are **sequential** (append to the end of log)

Log format

- When a transaction T_i starts
 - $\langle T_i, \text{start} \rangle$
- Record values before and after each modification:
 - $\langle T_i, X, \text{old_value_of_}X, \text{new_value_of_}X \rangle$
 - T_i is transaction id
 - X identifies the data item
- A transaction T_i is committed when its commit log record is written to disk
 - $\langle T_i, \text{commit} \rangle$



When to write log records into stable store?

- Write-ahead logging (WAL): Before X is modified on disk, the log record pertaining to X must be flushed
- Without WAL, system might crash after X is modified on disk but before its log record is written to disk—no way to undo

Undo/redo logging example

T_1 (balance transfer of \$100 from A to B)



read(A, a); $a = a - 100$;

write(A, a);

read(B, b); $b = b + 100$;

write(B, b);

Memory buffer	
A =	800 700
B =	400 500



WAL: Before A,B are modified on disk, their log info must be flushed

Undo/redo logging example cont.


T_1 (balance transfer of \$100 from A to B)

read(A, a); $a = a - 100$;



write(A, a);

read(B, b); $b = b + 100$;

write(B, b);



Steal: can flush
before commit




If system crashes before T_1 commits, we have
the old value of A stored on the log to **undo T_1**

Undo/redo logging example cont.

T_1 (balance transfer of \$100 from A to B)


```
read(A, a); a = a - 100;  
write(A, a);  
read(B, b); b = b + 100;  
write(B, b);  
commit;
```



If system crashes before we flush the changes of A, B to the disk, we have their new committed values on the log to redo T_1

Log example

- Redo phase:



List of active transactions at crash:


T1 T2 T3



Start of log


redo
redo
redo
redo
redo
redo
redo

End of log




Log example

- Redo phase:



Start of log

End of log




List of active transactions at crash:

T_1 ~~T_2~~ T_3



Log example

- Redo phase:



List of active transactions at crash:


T1 T2 T3 T4



Start of log


redo
redo
redo
redo
redo
redo
redo
redo
redo

End of log




Log example

- Redo phase:




List of active transactions at crash:

T1 T2 ~~T3~~ T4



Log example

- Redo phase:



List of active transactions at crash:

T1 ~~T2 T3~~ T4



Start of log

→
redo
redo


End of log →

Log

T_1 , start
 T_1 , x, 99, 100
 T_2 , start
 T_2 , y, 199, 200
 T_3 , start
 T_3 , z, 51, 50
 T_2 , w, 1000, 10
 T_2 , commit
 T_4 , start
 T_3 , abort
 T_4 , y, 200, 50


Log example

- Undo phase: T1, T4



List of active transactions at crash:

T1 ~~T2 T3~~ T4



Undo/redo logging

- U : used to track the set of active transactions at crash
- Redo phase: scan **forward** to end of the log
 - For a log record $\langle T, \text{start} \rangle$, add T to U
 - For a log record $\langle T, \text{commit} | \text{abort} \rangle$, remove T from U
 - For a log record $\langle T, X, \text{old}, \text{new} \rangle$, issue $\text{write}(X, \text{new})$
- Basically repeats history!
- Undo phase: scan log **backward**
 - Undo the effects of transactions in U
 - That is, for each log record $\langle T, X, \text{old}, \text{new} \rangle$ where T is in U , issue $\text{write}(X, \text{old})$, and log this operation too (part of the “repeating-history” paradigm)
 - Log $\langle T, \text{abort} \rangle$ when all effects of T have been undone

Checkpointing

- Shortens the amount of log that need to be undone or redone when a failure occurs
- A checkpoint record contains a list of active transactions
- Steps:
 1. Write a **begin_checkpoint** record into the log
 2. Collect the checkpoint data into the stable storage
 3. Write an **end_checkpoint** record into the log

Summary

- Concurrency control
 - 2PL: guarantees a conflict-serializable schedule
 - Deadlock problem
- Recovery: undo/redo logging
 - Normal operation: write-ahead logging, no force, steal
 - Recovery: first redo (forward), and then undo (backward)

